

Data Parallel Programming III

Accelerate

high performance computing in Haskell

Gabriele Keller, Utrecht University



acceleratehs.org



AccelerateHS

Overview

- Programming in Accelerate
- Implementation aspects of Accelerate
- Implementing an irregular parallel algorithm in Accelerate



Accelerate

- a spin-off from Data-Parallel Haskell
- similar to Futhark in many respects
 - restricts the computations to (mostly) regular nested computations
 - occupies a different place in the design space:
 - somewhat more abstract
 - is **deeply embedded** in Haskell
 - not a stand-alone language, but a Haskell library with the look and feel of one
 - our research is both about high-performance computing and the implementation of embedded languages

shallow embedding

fixed set of data types, operations

each operation can be highly efficient

no domain specific inter-function
optimisations

host-language
compiler in
charge of
optimisations

deep embedding

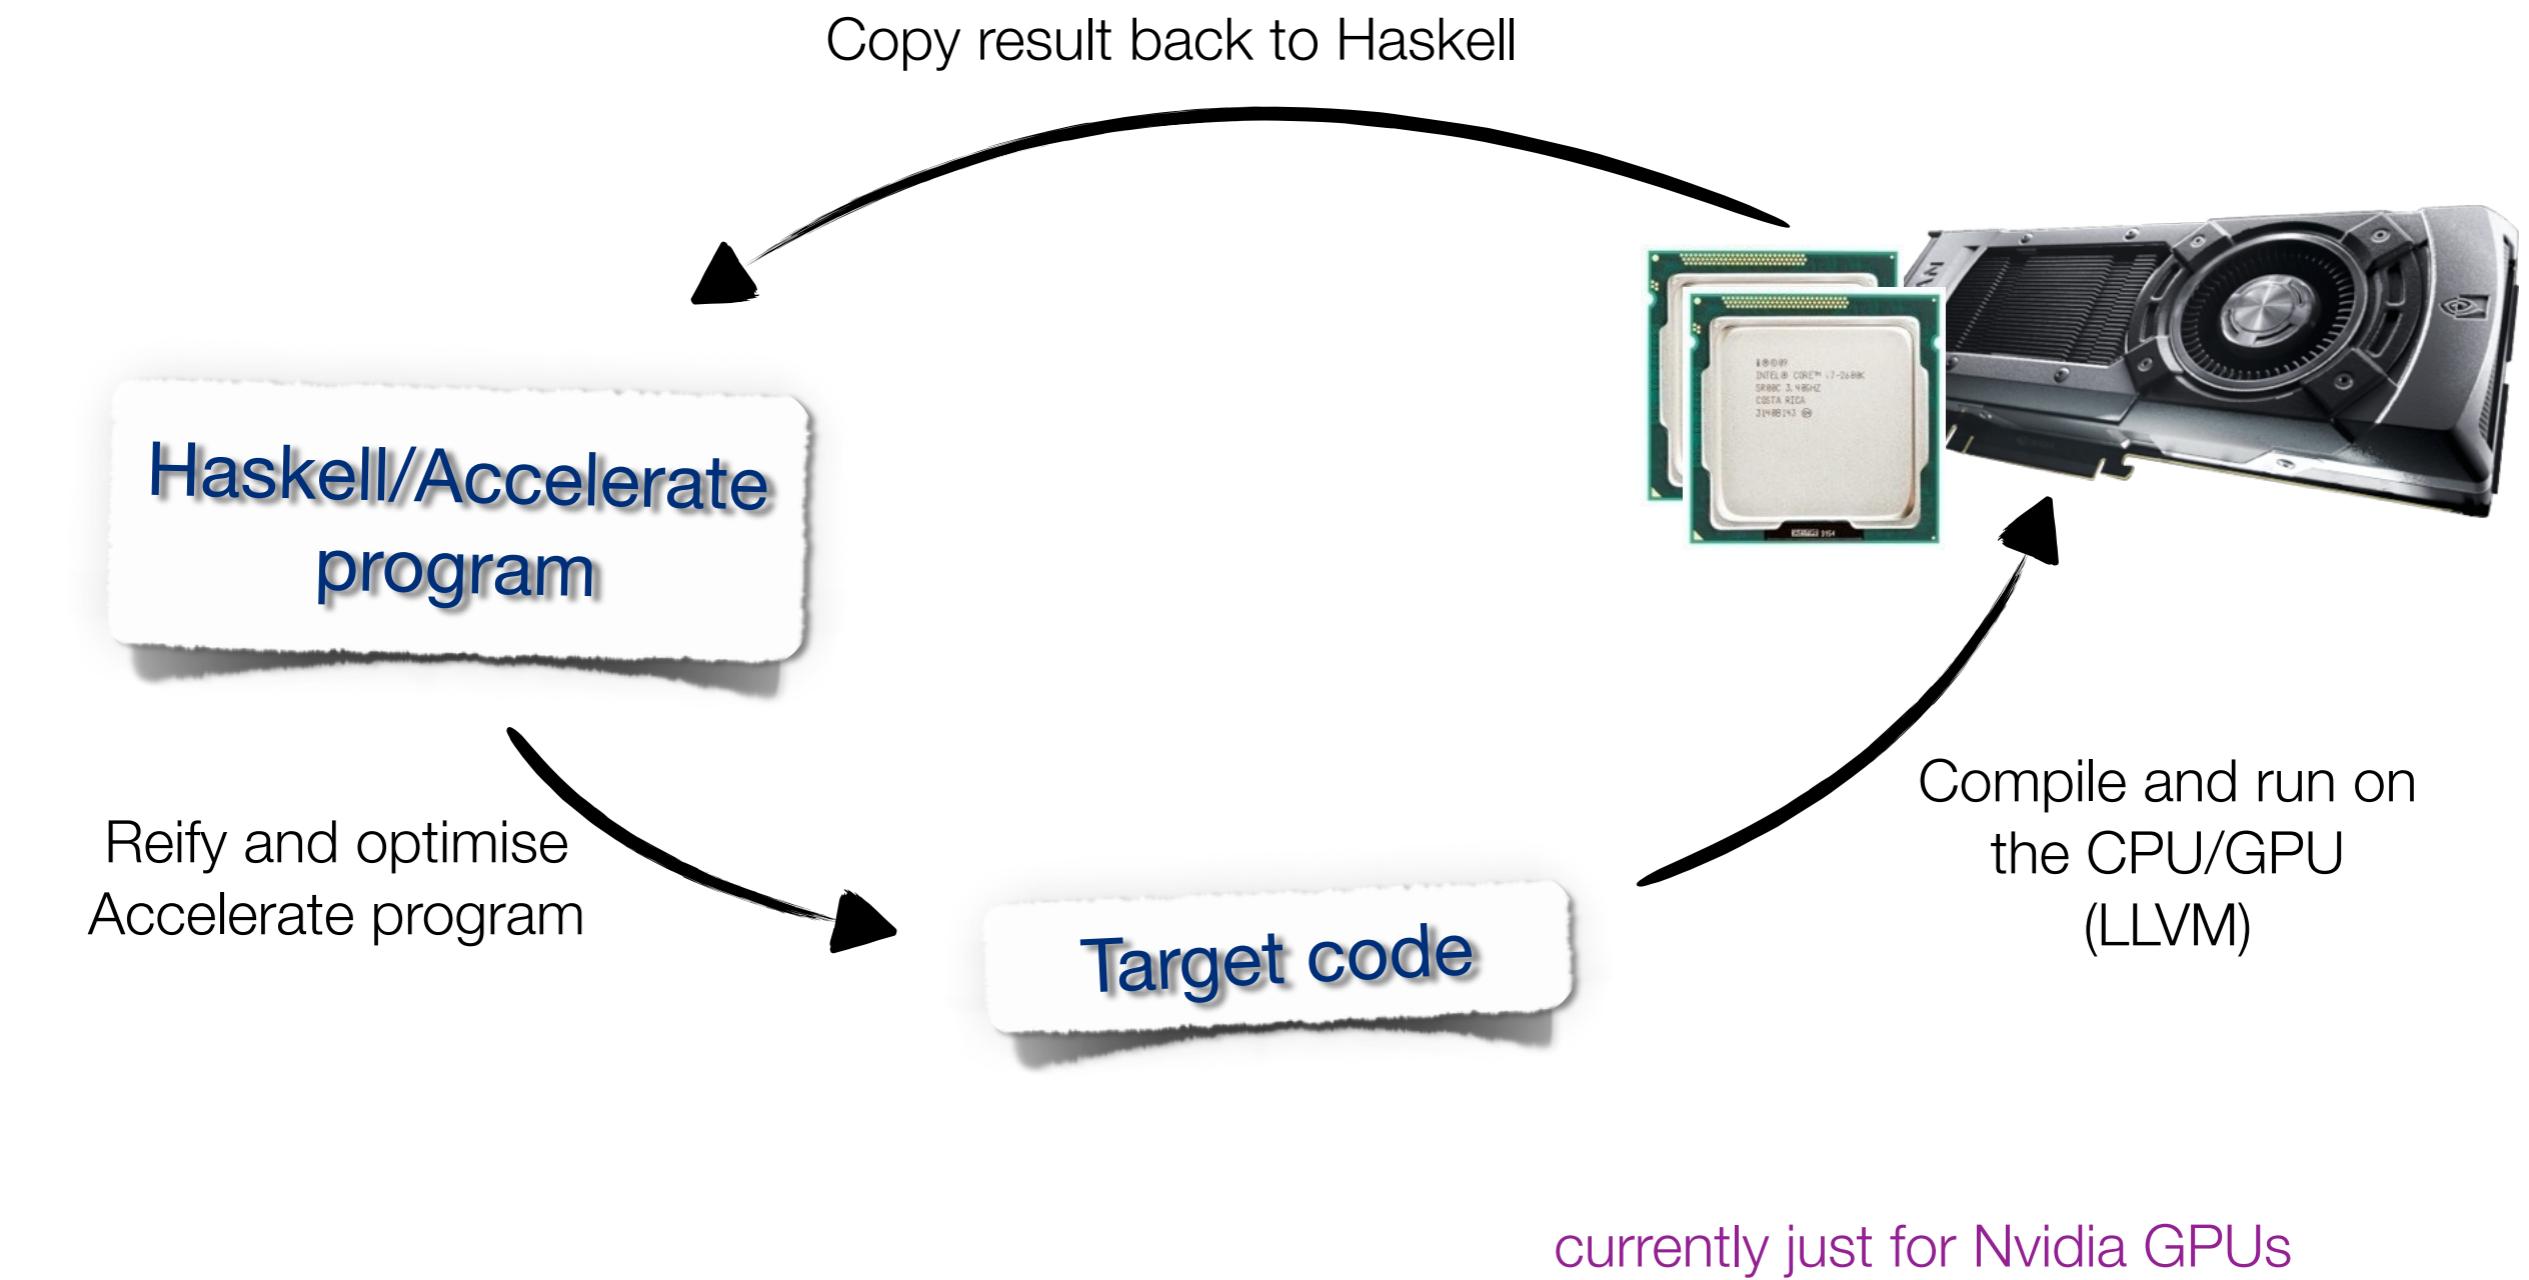
‘language in a language’

abstract syntax tree of the computation available
at run time

EDSL implementation in charge of code generation,
optimisation

Accelerate

An deeply embedded language for data-parallel arrays



Example: dot product

Futhark:

```
let dotp (x: []i32) (y: []i32): i32 =  
    reduce (+) 0 (map2 (*) x y)
```

Example: dot product

```
import Prelude

dotp :: Num a
      => [a] -> [a] -> a
dotp xs ys = foldl (+) 0 (zipWith (*) xs ys)
```

Example: dot product

```
import Data.Vector.Unboxed

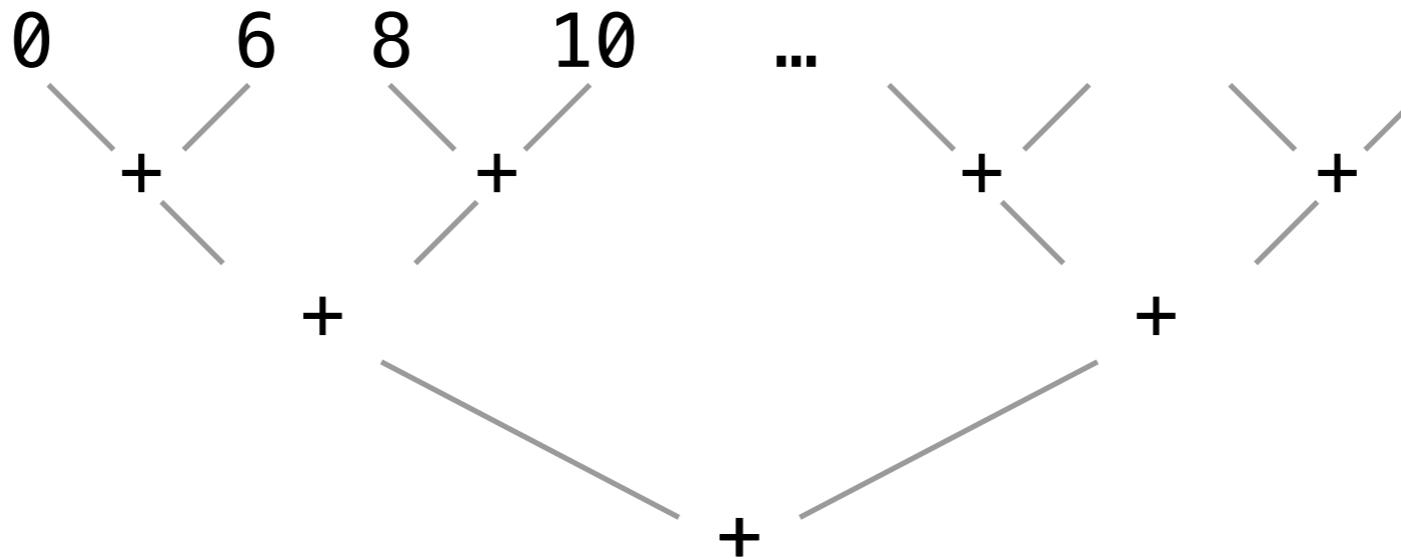
dotp :: (Unbox a, Num a)
      => Vector a
      -> Vector a
      -> a
dotp xs ys = foldl (+) 0 (zipWith (*) xs ys)
```

Accelerate

Collective array operations compile to parallel code

```
import Data.Array.Accelerate
dotp :: (Elt a, Num a)
      => Acc (Vector a)
      => Acc (Vector a)
      -> Acc (Scalar a)

dotp xs ys = fold (+) 0 (zipWith (*) xs ys)
```



Accelerate

Collective operations which compile to parallel code

```
import Data.Array.Accelerate
dotp :: (Elt a, Num a)
      => Acc (Vector a)
      -> Acc (Vector a)
      -> Acc (Scalar a)

dotp xs ys = fold (+) 0 (zipWith (*) xs ys)
```

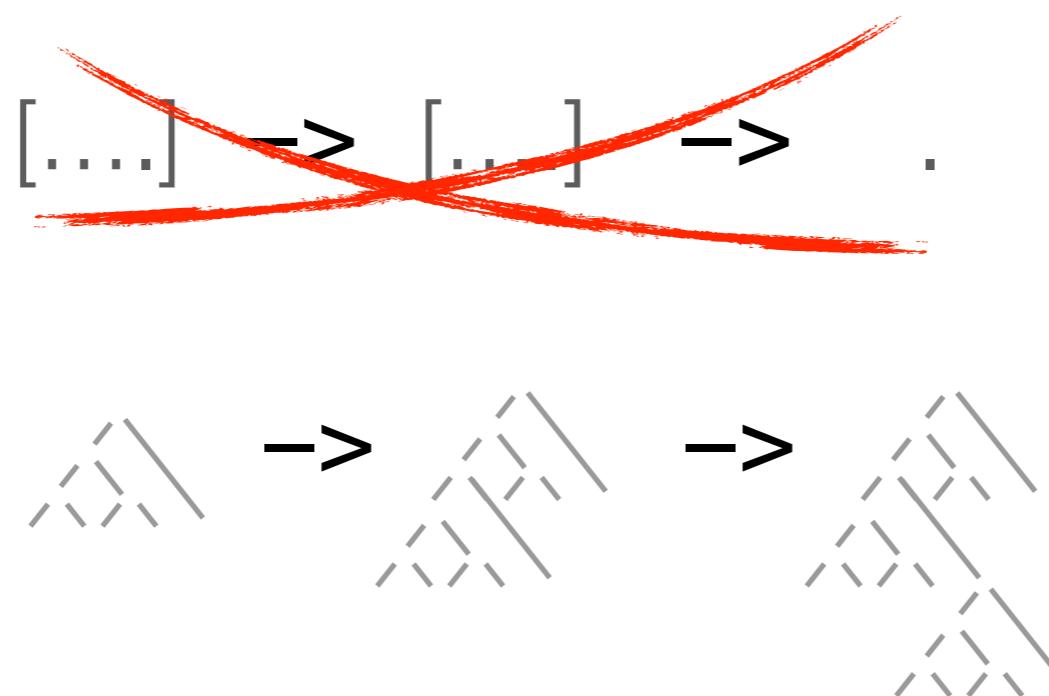
Accelerate

Collective operations which compile to parallel code

```
import Data.Array.Accelerate
dotp :: (Elt a, Num a)
      => Acc (Vector a)
      => Acc (Vector a)
      -> Acc (Scalar a)

dotp xs ys = fold (+) 0 (zipWith (*) xs ys)
```

language of collective,
parallel operations



Accelerate

Collective operations which compile to parallel code

```
import Data.Array.Accelerate
dotp :: (Elt a, Num a)
      => Acc (Vector a)
      -> Acc (Vector a)
      -> Acc (Scalar a)

dotp xs ys = fold (+) 0 (zipWith (*) xs ys)
```

Accelerate

Collective operations which compile to parallel code

language of sequential,
scalar expressions

fold (+) 0

```
fold :: (Shape sh, Elt e)
      => (Exp e -> Exp e -> Exp e)
      -> Exp e
      -> Acc (Array (sh :. Int) e)
      -> Acc (Array sh e)
```

language of collective,
parallel operations

rank-polymorphic

To enforce hardware restrictions,
general irregular nested parallel computation can't be expressed

```
import Data.Array.Accelerate
dotp :: (Elt a, Num a, Shape sh)
      => Acc (Array (sh :. Int) a)
      -> Acc (Array (sh :. Int) a)
      -> Acc (Array sh a)
dotp xs ys = fold (+) 0 (zipWith (*) xs ys)
```

\approx

```
dotp  $\cup$  (map dotp)  $\cup$  (map (map dotp))  $\cup$  ...
```

Accelerate

Collective operations which compile to parallel code

```
fold :: (Shape sh, Elt e)
      => (Exp e -> Exp e -> Exp e)
      -> Exp e
      -> Acc (Array (sh :. Int) e)
      -> Acc (Array sh e)
```

shape sh of the form Z :. Int :. Int :. ...

```
type DIM0      = Z
```

```
type Scalar a = Array DIM0 a
```

```
type DIM1      = DIM0 :. Int
```

```
type Vector a = Array DIM1 a
```

Accelerate

$$a = \begin{bmatrix} 0 & 1 & 2 & 3 & 4 & 5 & 6 & 7 & 8 & 9 \\ 10 & 11 & 12 & 13 & 14 & 15 & 16 & 17 & 18 & 19 \\ 20 & 21 & 22 & 23 & 24 & 25 & 26 & 27 & 28 & 29 \\ 30 & 31 & 32 & 33 & 34 & 35 & 36 & 37 & 38 & 39 \\ 40 & 41 & 42 & 43 & 44 & 45 & 46 & 47 & 48 & 49 \end{bmatrix}$$

```
shape a = (Z :. 5 :. 10)
```

```
arr ! (Z :. 2 :. 3) = 23
```

```
slice arr (Z :. 2 :. All) = [20 21 22 23 24 25 26 27 28 29]
```

- Accelerate supports the usual collection oriented operations:

- maps/zipWith, various flavours of scan and folds,
- indexing, slicing, permutations
- segmented operations
- parallel ‘forall’:

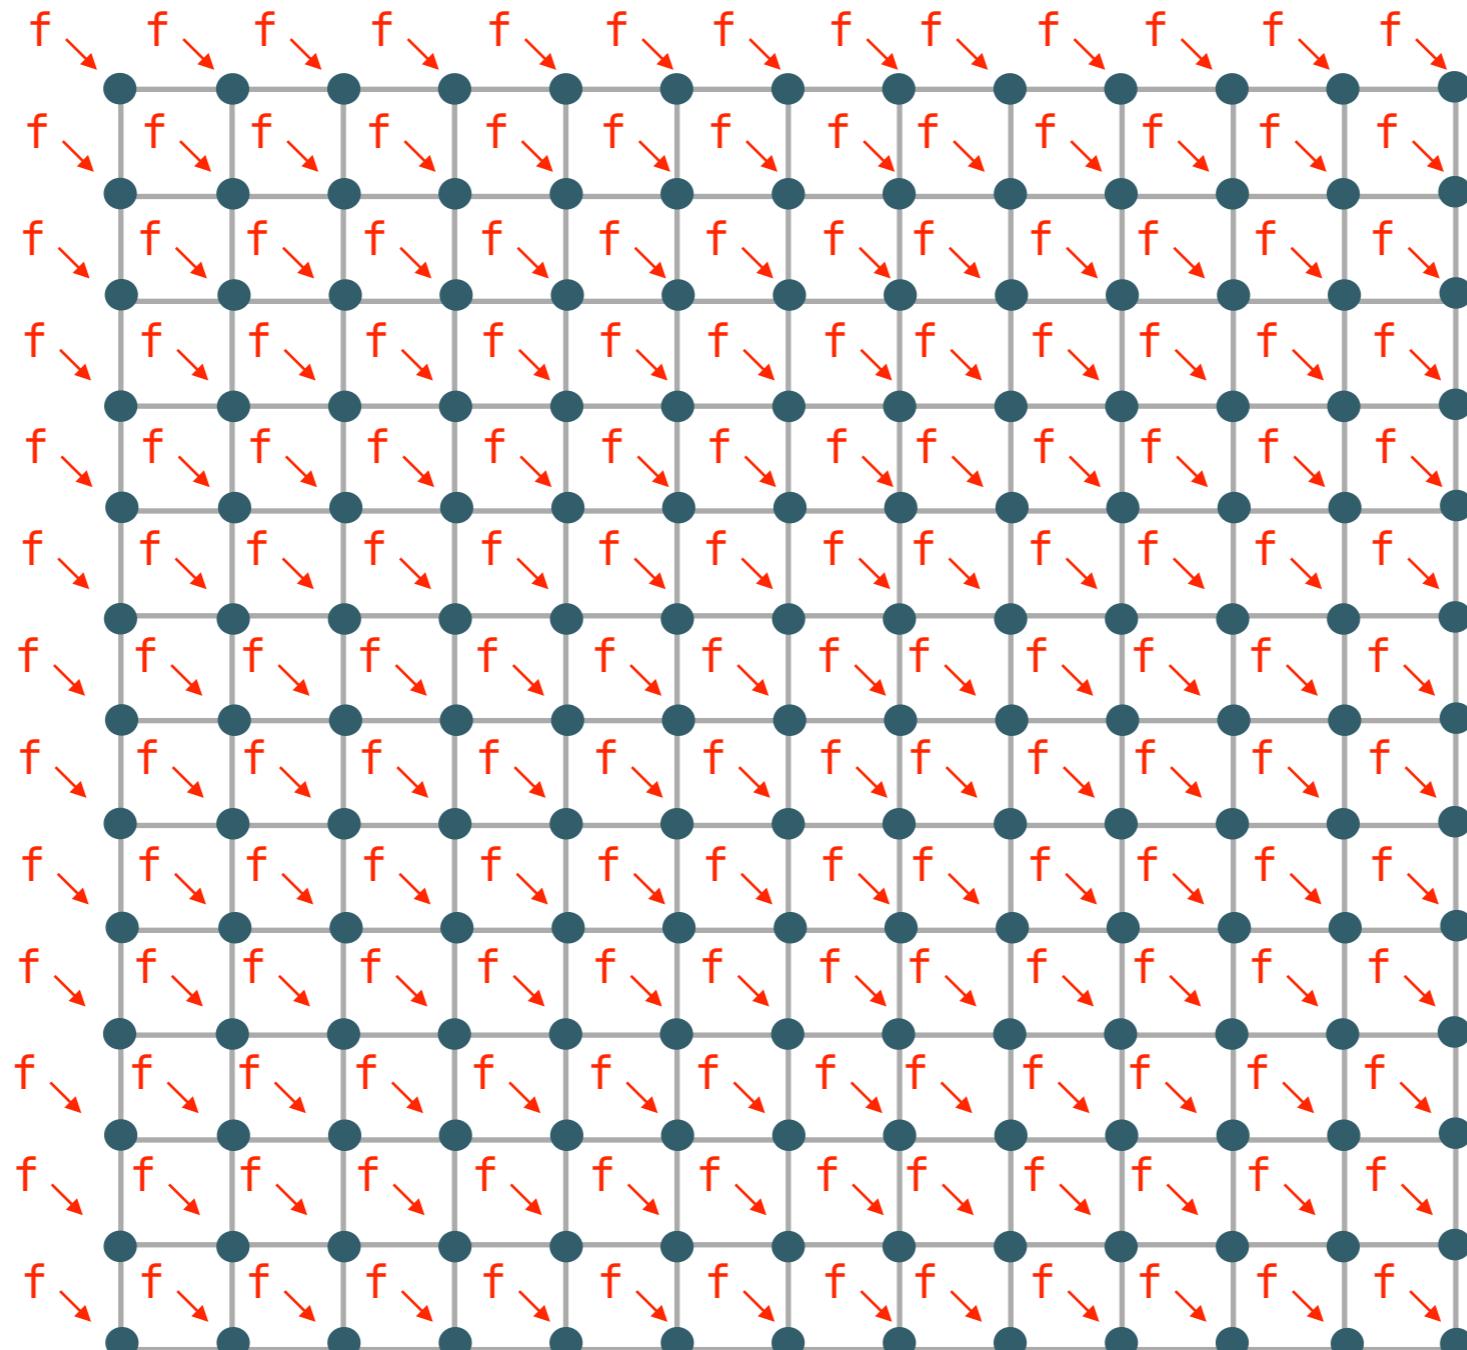
```
generate :: (Shape sh, Elt a)
           => Exp sh -> (Exp sh -> Exp a)
           -> Acc (Array sh a)
```

- while loops and iteration on element and array level

```
while :: Elt a
        => (Exp a -> Exp Bool)
        -> (Exp a -> Exp a)
        -> Exp a
```

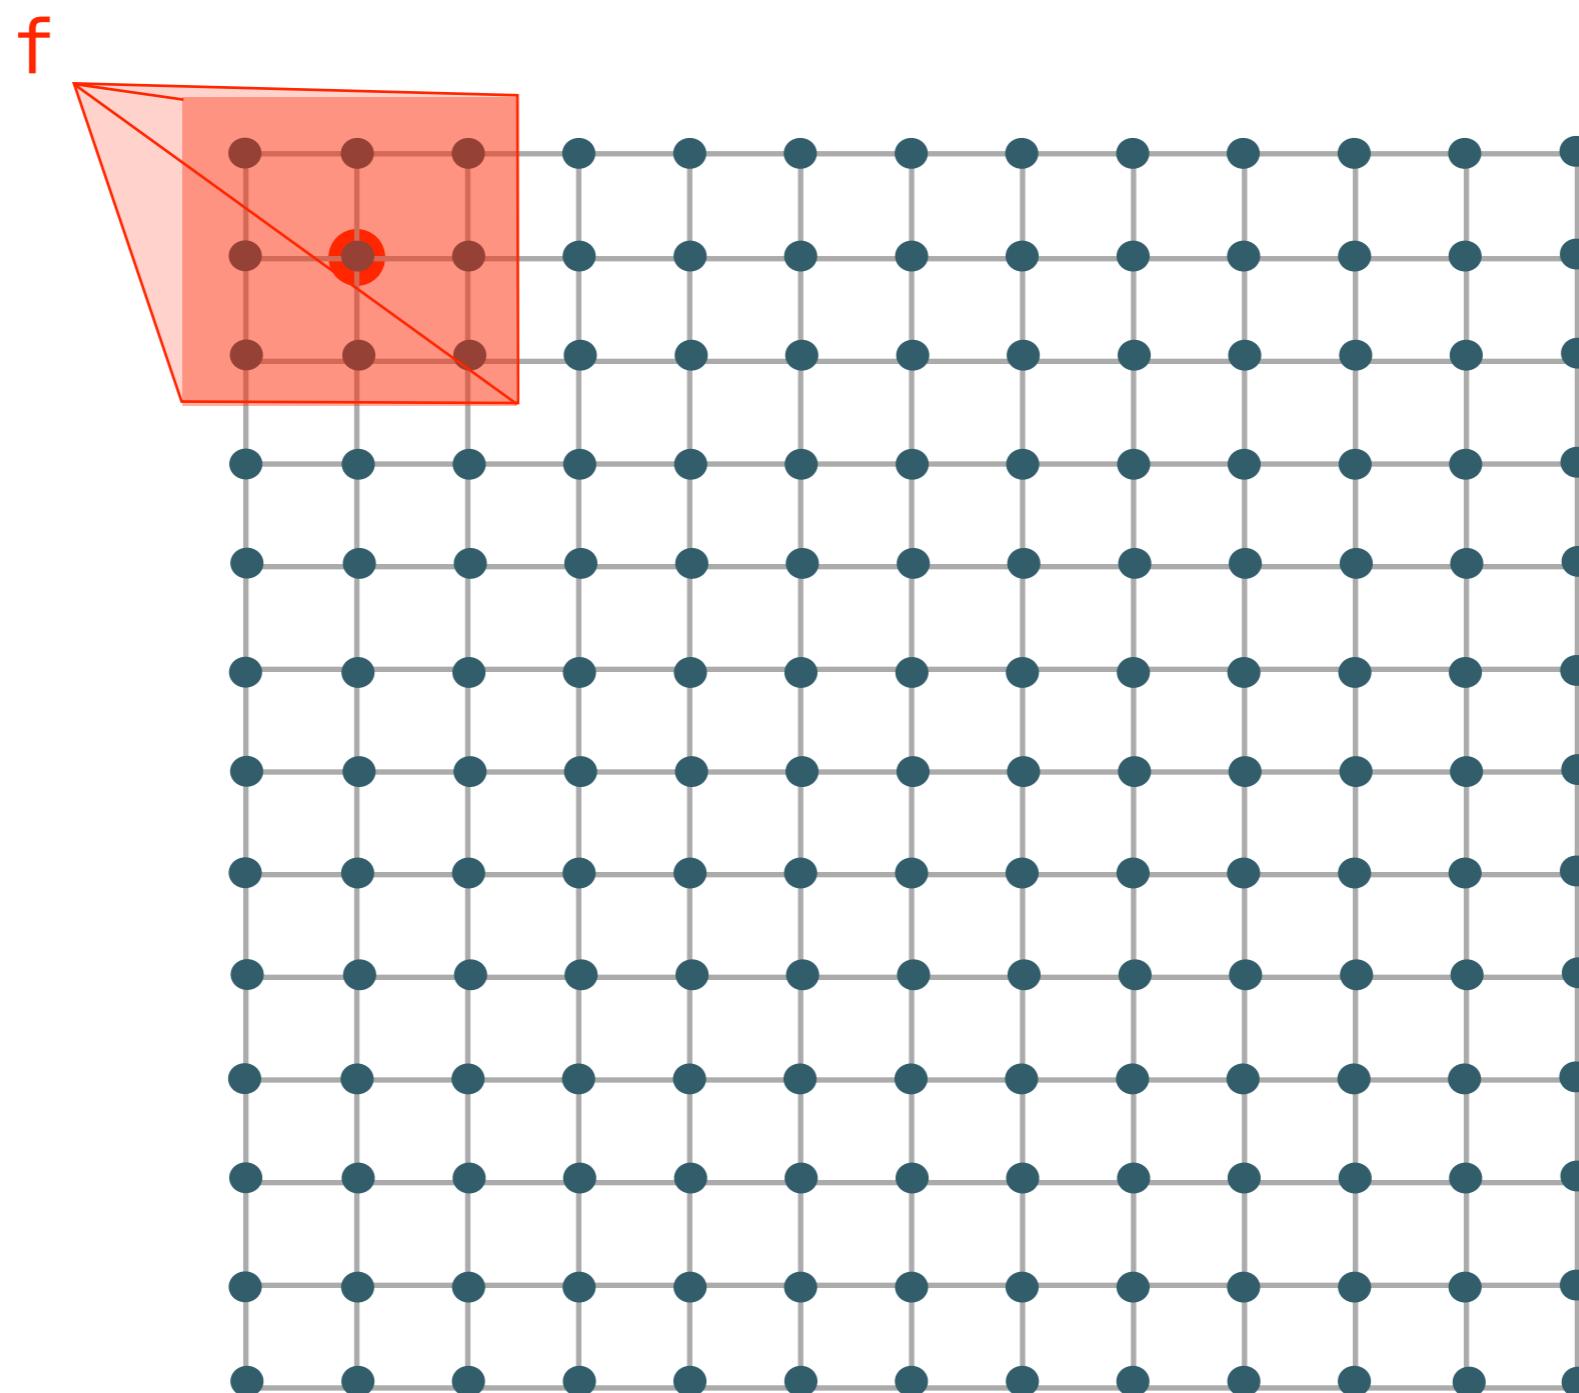
Stencil (convolution matrix) computations

```
map :: (Shape sh, Elt a, Elt b)  
=> (Exp a -> Exp b) -> Acc (Array sh a)  
-> Acc (Array sh a)
```

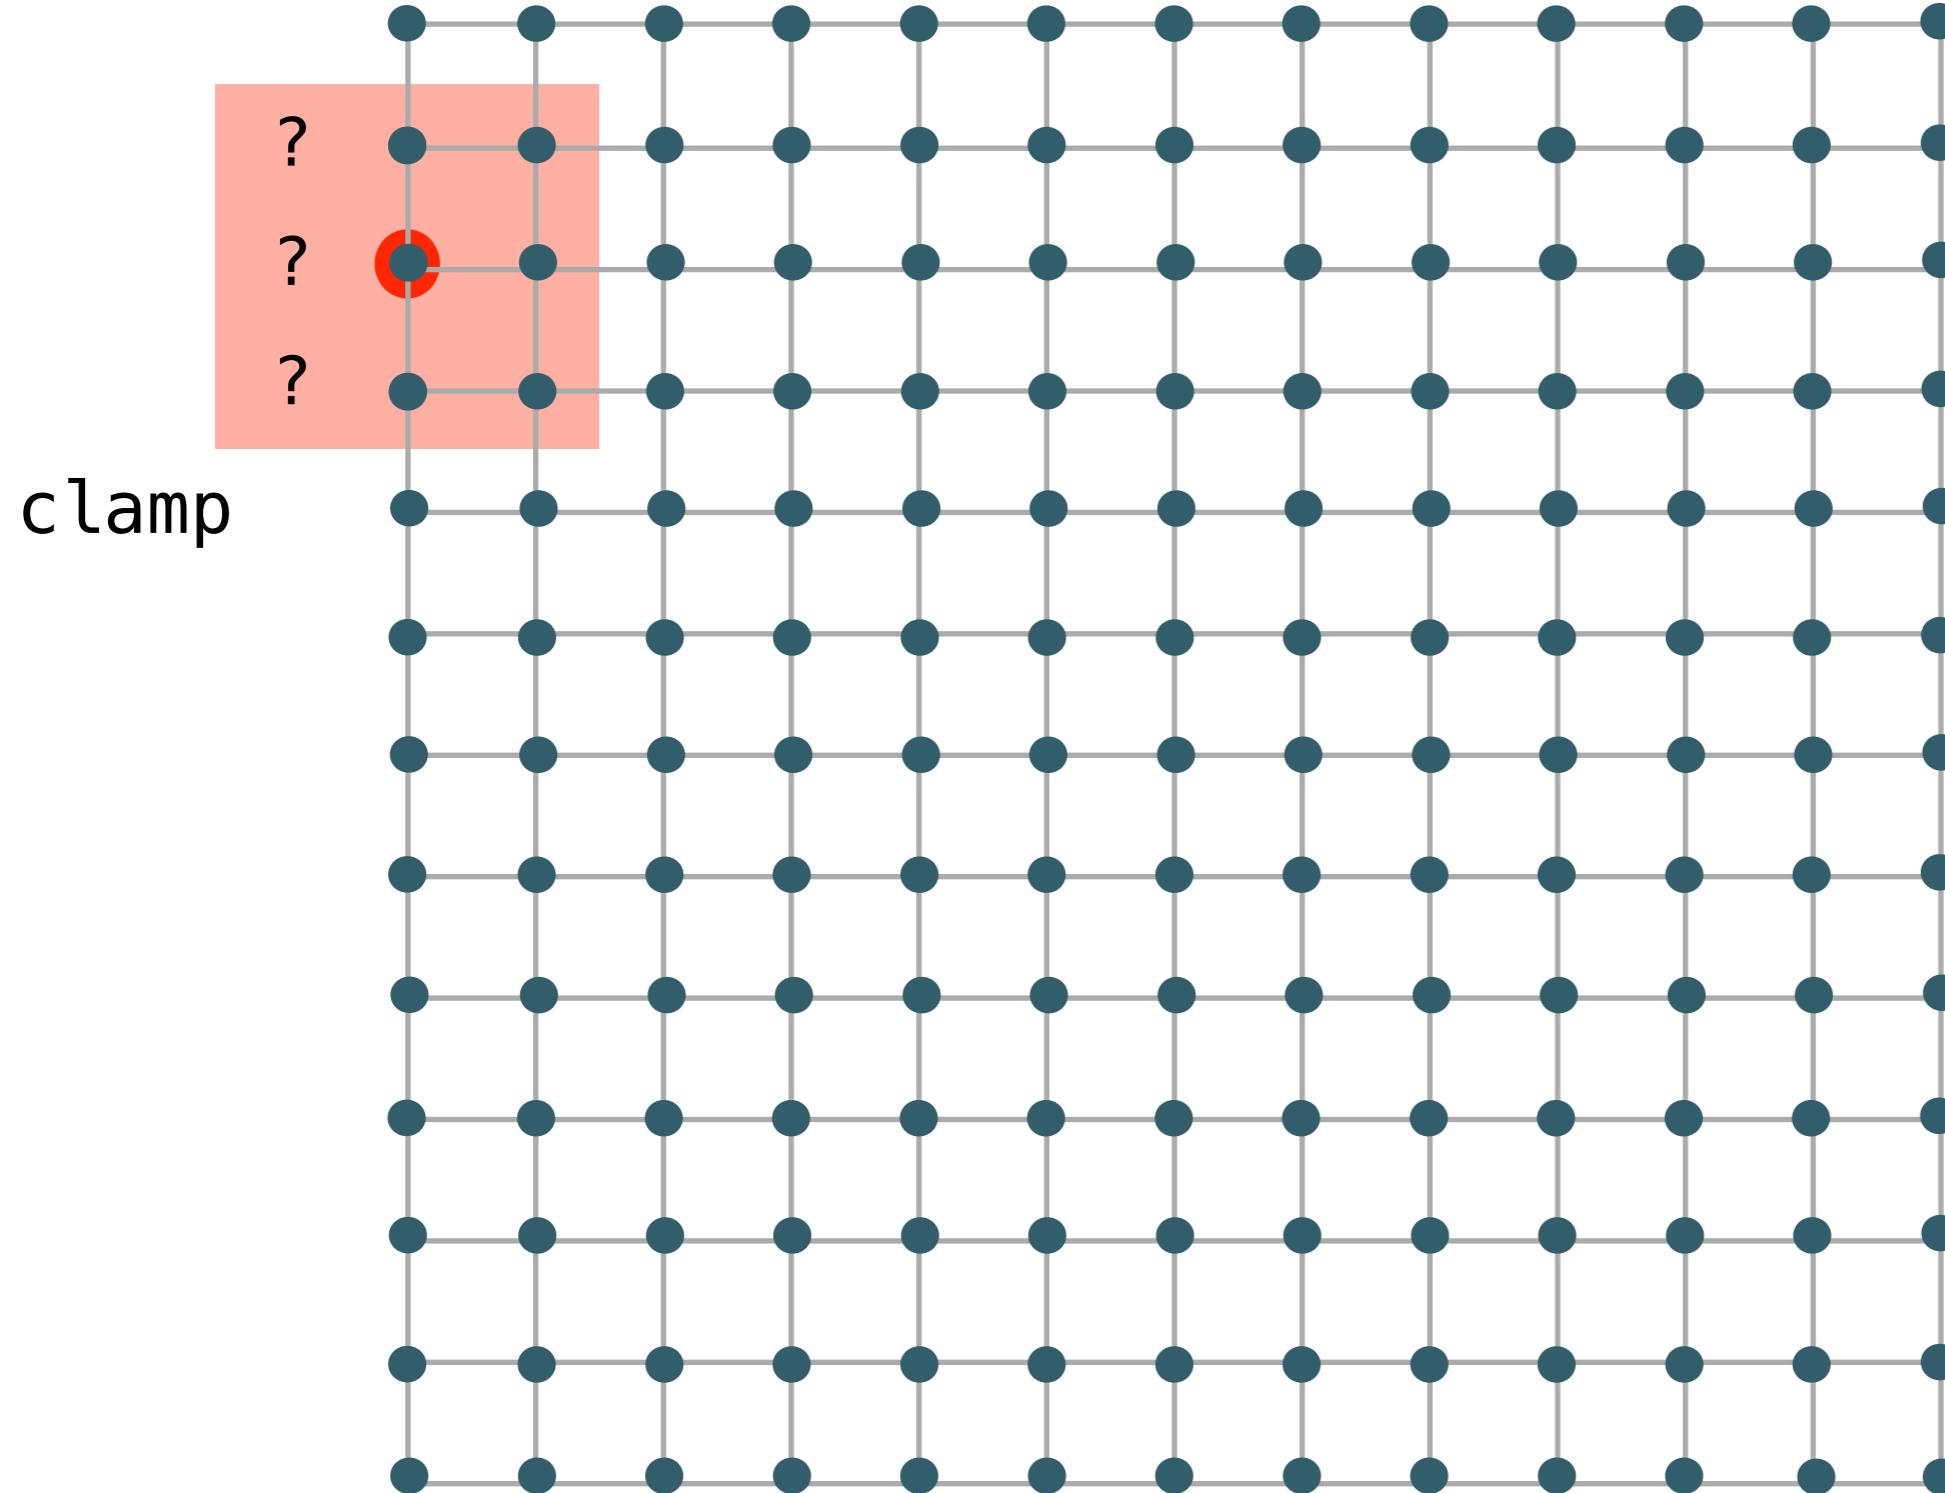


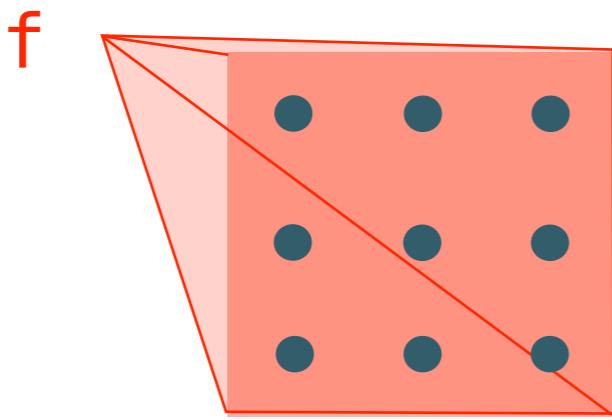
Stencil (convolution matrix) computations

```
stencil :: (Stencil sh a stencil, Elt b)
  => (stencil -> Exp b) -> Boundary (Array sh a) -> Acc (Array sh a)
  -> Acc (Array sh b)
```



Stencil computations - boundaries





$$= \begin{matrix} & + \\ \bullet & + \bullet & + \bullet \\ & + \\ & \bullet \end{matrix}$$

OpenCL

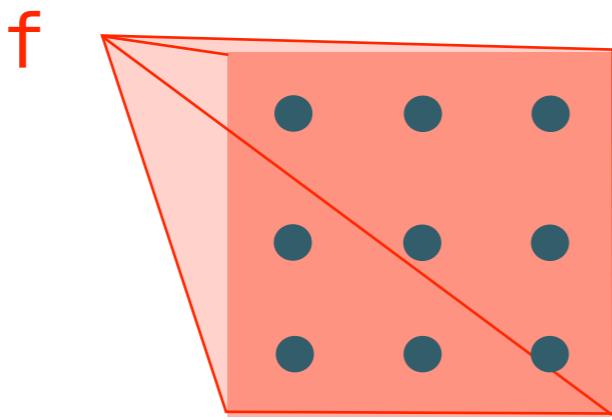
```
__kernel void simulate (__global float* arr
                      ,__global float* new_arr)
{
    const size_t cur = get_global_id(0);
    const size_t row = (size_t)cur/(size_t)Width;
    const size_t col = (size_t)cur%(size_t)Height;

    if (row > 0 && row < height-1
        && col > 0 && col < width-1) {
        new_arr[curr] =
            arr[cur] + arr[row * Width + col-1] ... ;
    } else if (row == 0 && col < width-1) {
        ...
    } else if ...
```

```
simulate :: Stencil3x3 Float-> Exp Float
simulate ((_, top, _),
          (left, curr, right),
          (_, bot, _))
= top + left + curr + right + bot

new_matrix
= stencil simulate clamp matrix
```

Accelerate



$$= \begin{matrix} & \bullet \\ & + \\ \bullet & + \bullet & + \bullet \\ & + \\ & \bullet \end{matrix}$$

OpenCL

```
__kernel void simulate (__global float* arr
                      ,__global float* new_arr)
{
    const size_t cur = get_global_id(0);
    const size_t row = (size_t)cur/(size_t)Width;
    const size_t col = (size_t)cur%(size_t)Height;

    if (row > 0 && row < height-1
        && col > 0 && col < width-1) {
        new_arr[curr] =
            arr[cur] + arr[row * Width + col-1] ... ;
    } else if (row == 0 && col < width-1) {
        ...
    } else if ...
```

Accelerate

```
simulate :: Stencil3x3 Float-> Exp Float
simulate ((_,      top,      _),
          (left,   curr, right),
          (_,      bot,      _))
= top + left + curr + right + bot

new_matrix
= stencil simulate clamp matrix
```

Executing an Accelerate Program

```
run :: Arrays a => Acc a -> a
```

```
import Data.Array.Accelerate
import Data.Array.Accelerate.LLVM.Native - CPU
```

```
vec1, vec2 :: Acc (Array DIM1 Float)
dotp xs ys = fold (+) 0 (zipWith (*) xs ys)
```

```
main =
  putStrLn $ show $ run (dotp vec1 vec2)
```

```
runN f :: (Arrays a, Arrays b, ... Arrays c)
=> a -> b -> ... -> c
```

The Elt class

- Members of the Elt class contain admissible **surface types** for array elements:
 - ()
 - Int, Int32, Int64, Word, Word32, Word64 ...
 - Float, Double
 - Char
 - Bool
 - Array indices formed from Z and (:) ..)
 - Tuples of all of these, e.g. (Bool, Int, (Float, Float))
 - To meet hardware restrictions, there are **no nested arrays** in Accelerate

Lifting values into the embedded language

```
0 :: (Num a, Elt a) => Exp a
```

```
fold (+) 0 (zipWith (*) xs ys)
```

- we declare **Num a => Exp a** to be an instance of **Num**
- **Num** values get automatically lifted into the language
- not possible for non-overloaded types and type constructors

- **Bool** (Exp a, Exp b)
- **(0, 4) ::(Num a, Num b) => (a, b)** Exp (a, b)

```
lift :: Lift e => e -> Expr (Plain e)
```

- where **Plain** is an associated type synonym which strips all the **Expr** constructors

```
Plain (Exp Float, Exp Float) ~ Plain (Exp (Float, Exp Float))
```

Lifting and unlifting values

```
swap :: (a, b) -> (b, a)
swap (x, y) = (y, x)
```

can't use pattern matching directly!

```
swapExp :: Exp (a, b) -> Exp (b, a)
swapExp xy = lift (sndExp xy, fstExp xy)
```

can we have an `unlift` function?

... => `Exp a -> a`

in general, this would mean evaluating the AST (`run`)!

Lifting and unlifting values

But it does work if the type has only one constructor:

$$\begin{array}{ll} \text{Exp ()} & () \\ \text{Exp (a, b)} & (\text{Exp a}, \text{Exp b}) \end{array}$$

unlift :: Unlift e =>

swapExp (unlift -> (x,
= lift (y, x))

swapExp (unlift -> (x :
= lift (y, x))

- Couldn't match type 'Plain a' with 'Plain a0'
Expected type: c (Plain b, Plain a) -> c1 (Plain a, Plain b)
Actual type: c (Plain (b0, a0)) -> c1 (Plain (a0, b0))
NB: 'Plain' is a non-injective type family
The type variable 'a0' is ambiguous
- When checking that the inferred type

```
swapExp :: forall (c1 :: * -> *) b a (c2 :: * -> *).  
          (Unlift c1 (b, a), Lift c2 (a, b)) =>  
          c1 (Plain (b, a)) -> c2 (Plain (a, b))
```

is as general as its inferred signature

```
swapExp :: forall (c1 :: * -> *) b a (c2 :: * -> *).  
          (Unlift c1 (b, a), Lift c2 (a, b)) =>  
          c1 (Plain b, Plain a) -> c2 (Plain a, Plain b)
```
- Could not deduce (Unlift c (b0, a0))
from the context: (Unlift c (b, a), Lift c1 (a, b))
bound by the inferred type for 'swapExp':

```
forall (c :: * -> *) b a (c1 :: * -> *).  
          (Unlift c (b, a), Lift c1 (a, b)) =>  
          c (Plain b, Plain a) -> c1 (Plain a, Plain b)
```

at ... Accelerate Examples/Main.hs:32:1-42
The type variables 'b0', 'a0' are ambiguous
- When checking that the inferred type

```
swapExp :: forall (c1 :: * -> *) b a (c2 :: * -> *).  
          (Unlift c1 (b, a), Lift c2 (a, b)) =>  
          c1 (Plain (b, a)) -> c2 (Plain (a, b))
```

Lifting and unlifting

Haskell's **bidirectional pattern synonyms** can do the job for us:

```
data Calc  = BinOpApp BinOp Calc Calc
           | ...
data BinOp = Plus | Times | ...

pattern Add e1 e2 = BinOpApp Plus e1 e2

swapArgs :: Calc -> Calc
swapArgs (Add e1 e2) = Add e2 e1
```

Accelerate provides pattern synonyms **T2, T3, T4,...** for pairs, triples,...

```
swapExp :: Exp (a, b) -> Exp (b, a)
swapExp (T2 x y) = T2 y x
```

User defined data types - the Elt class

- What about user defined data types?
 - tuples support built in

```
type Point    = (Float, Float)
type Complex = (Float, Float)
```
 - we want to support product and sum types, recursive types
- Two problems need to be solved:
 - how to map them efficiently to target hardware
 - how to represent them in the AST
 - both should happen without much overhead for the user

User defined data types - the Elt class

- We use Haskell Generics to map arbitrary algebraic data types to uniform representation
 - we use that both for user defined and built-in types
 - for representation in the AST and for mapping to hardware
 - for product types, mapping is clear
 - $(\dots, \dots, \dots, \dots) \quad (((), \dots), \dots), \dots)$
 - sum types are much more complicated:
 - $\text{Exp} (\text{Either } a \ b) \quad \text{Either}(\text{Exp } a) \ (\text{Exp } b)$

User defined data types - the Elt class

```
data Point = Point Float Float
  deriving (Show, Generic, Elt, IsProduct Elt)

pattern Point_ :: Exp Float → Exp Float → Exp Point
pattern Point_ x y = Pattern (x,y)

foo :: Exp Point → Exp Point
foo (Point_ x y) = (Point_ y x)
```

User defined data types - the Elt class

```
data Option a
  = None
  | Some a deriving (Generic, Elt, IsTuple)
```

```
mkPattern ''Option
```

```
simple :: Exp (Option Int) → Exp Int
simple None_      = 0
simple (Some_ x) = x
```

**But now, let's have another look at
how to implement an algorithm which is
has nested irregular parallelism**

Implementing Quicksort in Accelerate

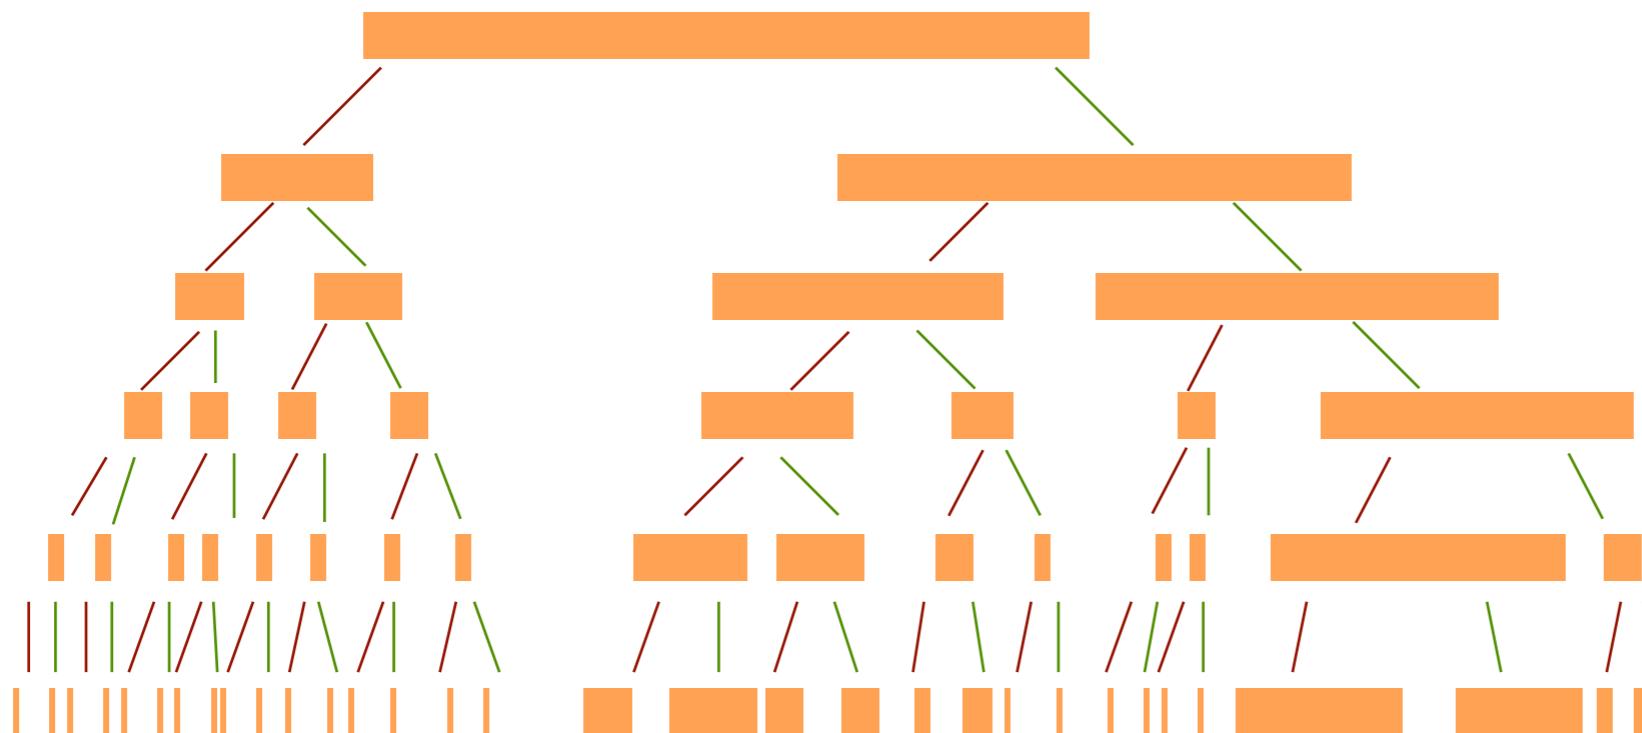
```
qsort []      = []
qsort (p : xs) = smaller' ++ [p] ++ greater'
```

where

```
smaller = filter (< p) xs
```

```
greater = filter (> p) xs
```

```
[smaller', greater'] = map qsort [smaller, greater]
```



Implementing Quicksort in Accelerate

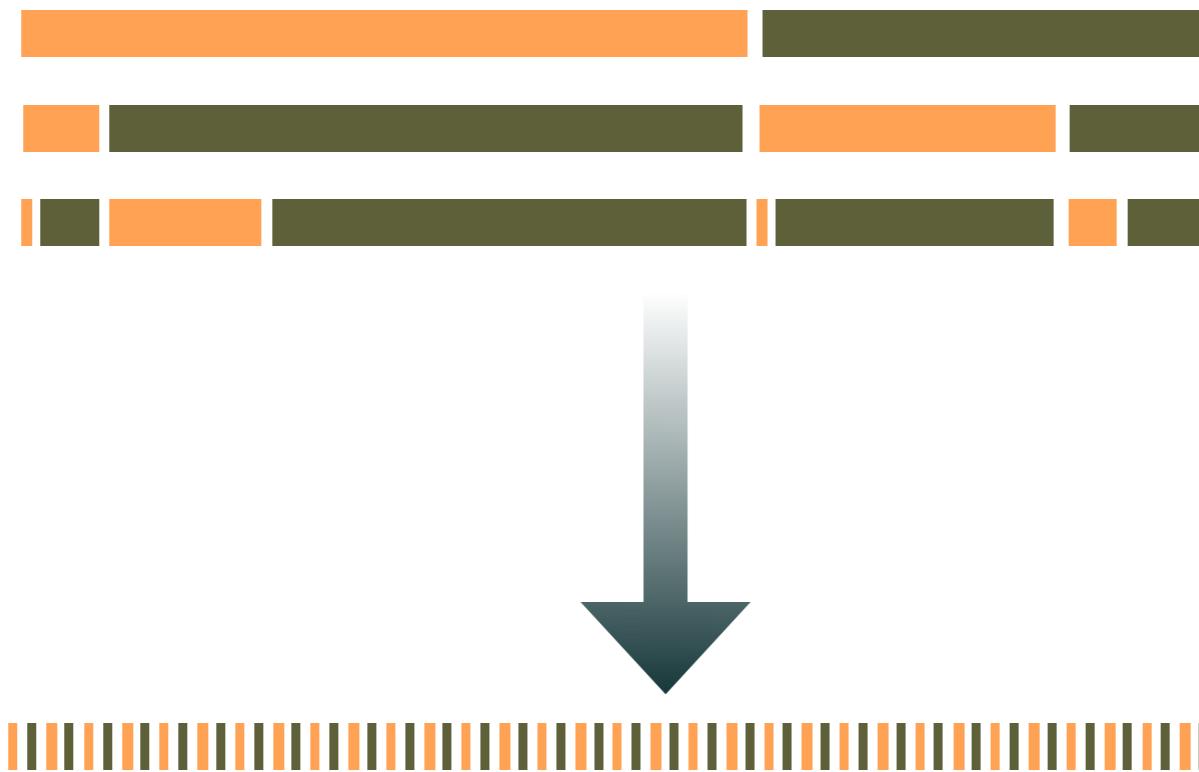
- NESL would generate
 - a vectorised version of quicksort :: [[Float]] -> [[Float]]
 - represent data and shape information separately

```
qsort []      = []
qsort (p : xs) = smaller' ++ [p] ++ greater'
where
  smaller = filter (< p) xs
  greater = filter (> p) xs
  [smaller', greater'] = map qsort [smaller, greater]
```

```
quicksort' :: ([Float], [Int]) -> ([Float], [Int])
quicksort' xs seg = ...
```

```
where
  nonEmptySegs = filter (/= 0) segs
  segStarts    = preScan 0 (+) nonEmptySegs
  ps           = backpermute xs segStarts
  smallerFlags = zipWith (<) (repl' ps seg) xs
  smallerSegs  = fold' (+) ... smallerFlags
  ...
  sortedSegs   =
    quicksort' ((smaller, smallerSegs) ++' (greater, greaterSegs))
```

Implementing Quicksort in Accelerate



- we have segmented ops in Accelerate , so we could implement the manually flattened quick sort
 - not trivial to get right
 - would suffer from the similar performance problems

Implementing Quicksort in Accelerate

- We can still take the inspiration from the NESL approach, but avoid inefficiencies
 - shape information can be represented using segment size:

`[[1,2],[],[3,4],[5,6,7,8]]`

`[1,2, 3,4,5,6,7,8]`

`[2, 0,2, 4]`

- ratio between segment descriptor size and data size can become arbitrarily big (many empty segments)
 - problem for quick sort

Implementing Quicksort in Accelerate

- **flag vectors** can be used to mark segment boundaries

`[[1,2],[],[3,4],[5,6,7,8]]`

`[1,2,3,4,5,6,7,8]
[T,F,T,F,T,F,F]`

- same size as data vector, can't represent empty segments
 - doesn't matter for quicksort - segment information will be discarded when sorted

Implementing Quicksort in Accelerate

- We can implement quicksort using scans, map/zipWiths and a permutation
 - idea: we count the number of elements greater and smaller than the pivot, in parallel for each segment
 - using this information, we calculate the index of each element in the partitioned array



Implementing Quicksort in Accelerate

```
postscanl op v [] = []
postscanl op v (x:xs)
= v' : postscanl op v' xs
where v' = v `op` x
```

```
postscanl (+) 10 [1,2,3]
=
[11, 13, 16]
```

```
postScanSeg op vals flags
= postscanl
  (\x -> \(y, flag) -> if flag
                           then y
                           else x `op` y) undefined
```

[3, 4, 2, 1, 5, 8, 9, 6, 7] vals

[T, F, F, F, T, T, F, F, F] flags

postScanSeg (+) vals flags = [3, 7, 9, 10, 5, 8, 17, 23, 30]

postScanSeg const vals flags = [3, 3, 3, 3, 5, 8, 8, 8, 8]

postScanSeg const [0,1..] flags = [0, 0, 0, 0, 4, 5, 5, 5, 5]

[3, 4, 2, 1, 5, 8, 9, 6, 7]	vals
[T, F, F, F, T, T, F, F, F]	flags
[0, 0, 0, 0, 4, 5, 5, 5, 5]	segStarts
[3, 3, 3, 3, 5, 8, 8, 8, 8]	pivots

zipWith ($<$)

[3, 4, 2, 1, 5, 8, 9, 6, 7]	vals
[T, F, F, F, T, T, F, F, F]	flags
[0, 0, 0, 0, 4, 5, 5, 5, 5]	segStarts
[3, 3, 3, 3, 5, 8, 8, 8, 8]	pivots
[T, T, F, F, T, T, T, F, F]	lessOrEq

[3, 4, 2, 1, 5, 8, 9, 6, 7] vals

[T, F, F, F, T, T, F, F, F] flags

[0, 0, 0, 0, 4, 5, 5, 5, 5] segStarts

[3, 3, 3, 3, 5, 8, 8, 8, 8] pivots

[T, T, F, F, T, T, T, F, F] lessOrEq

```
[3, 4, 2, 1, 5, 8, 9, 6, 7]    vals
[T, F, F, F, T, T, F, F, F]    flags

[0, 0, 0, 0, 4, 5, 5, 5, 5]    segStarts
[T, T, F, F, T, T, T, F, F]    lessOrEq
map (\f -> if f
      then 1
      else 0) [1, 1, 0, 0, 1, 1, 1, 0, 0]
segmented (+)
postscan,
starting at -1
[0, 1, 1, 1, 0, 0, 1, 1, 1]    indexLargerEq
```

```
[3, 4, 2, 1, 5, 8, 9, 6, 7]    vals
[T, F, F, F, T, T, F, F, F]    flags

[0, 0, 0, 0, 4, 5, 5, 5, 5]    segStarts
[T, T, F, F, T, T, T, F, F]    lessOrEq
map (\f -> if f
      then 0
      else 1) [1, 1, 0, 0, 1, 1, 1, 0, 0]
segmented (+)
postscan,
starting at -1
[-1,-1, 0, 1,-1,-1,-1, 0, 1]    indexSmaller
```

```
[3, 4, 2, 1, 5, 8, 9, 6, 7]    vals
[T, F, F, F, T, T, F, F, F]    flags

[0, 0, 0, 0, 4, 5, 5, 5, 5]    segStarts
[T, T, F, F, T, T, T, F, F]    lessOrEq

[0, 1, 1, 1, 0, 0, 1, 1, 1]    indexLargerEq
[-1,-1, 0, 1,-1,-1,-1, 0, 1]    indexSmaller
```

[3, 4, 2, 1, 5, 8, 9, 6, 7] vals

[T, F, F, F, T, T, F, F, F] flags

[0, 0, 0, 0, 4, 5, 5, 5, 5] segStarts

[T, T, F, F, T, T, T, F, F] lessOrEq

[0, 1, 1, 1, 0, 0, 1, 1, 1] indexLargerEq

[-1,-1, 0, 1,-1,-1,-1, 0, 1]
| +1 | +1 | +1 |
[2, 2, 2, 2, 0, 2, 2, 2, 2] indexSmaller

countSmaller

<code>[3, 4, 2, 1, 5, 8, 9, 6, 7]</code>	<code>vals</code>
<code>[T, F, F, F, T, T, F, F, F]</code>	<code>flags</code>
<code>[0, 0, 0, 0, 4, 5, 5, 5, 5]</code>	<code>segStarts</code>
<code>[T, T, F, F, T, T, T, F, F]</code>	<code>lessOrEq</code>
<code>[0, 1, 1, 1, 0, 0, 1, 1, 1]</code>	<code>indexLargerEq</code>
<code>[-1, -1, 0, 1, -1, -1, -1, 0, 1]</code>	<code>indexSmaller</code>
<code>[2, 2, 2, 2, 0, 2, 2, 2, 2]</code>	<code>countSmaller</code>

+ + + +

| | | |

`[, , 0, 1, , , , 5, 6]`

[3, 4, 2, 1, 5, 8, 9, 6, 7] vals

[T, F, F, F, T, T, F, F, F] flags

[0, 0, 0, 0, 4, 5, 5, 5, 5] segStarts

[T, T, F, F, T, T, T, F, F] lessOrEq

[0, 1, 1, 1, 0, 0, 1, 1, 1] indexLargerEq

[2, 2, 2, 2, 0, 2, 2, 2, 2] countSmaller

+ + + + +

+ + + + +

| | | | |
[2, 3, 0, 1, 4, 7, 8, 5, 6]

permute

[3, 4, 2, 1, 5, 8, 9, 6, 7] vals

[T, F, F, F, T, T, F, F, F] flags

[2, 1, 3, 4, 5, 6, 7, 8, 9] vals

[T, F, T, F, T, T, F, T, F] flags

[2, 3, 0, 1, 4, 7, 8, 5, 6]

Implementing Quicksort in Accelerate

- algorithm has depth $(\log n)$ on average
- implementation not Accelerate specific
 - would benefit from destructive updates in the permute (we're working on it!)
- changes and optimisations necessary to get from flattened NDP algorithm to *efficient* manually flattened are non-trivial
 - not likely to be derivable by compiler
 - more promising to identify and support common irregular patterns?

Current research directions

- Support for a limited form of irregular nested computations
 - aimed at detecting regular subcomputations
- Automatic derivation of uniqueness information
 - enables more destructive updates
- Support for more flexible schedules
- Improving fusion optimisation
 - ‘horizontal fusion/tupling’
- Other backends (AMD, Metal)
- Automatic differentiation in Accelerate