


Parallel Functional Programming in Java 8

Peter Sestoft
IT University of Copenhagen

Parallel Functional Programming
Chalmers, Thursday 2023-05-04

The speaker

- MSc 1988 computer science and mathematics and PhD 1991, DIKU, Copenhagen University
- KU, DTU, KVL and ITU; and Glasgow U, AT&T Bell Labs, Microsoft Research UK, Harvard University
- Functional, object-oriented, and parallel software



1993




2002, 2005, 2016




2004, 2012



2007



2012, 2017



2014

Plan

- Java 8 functional programming
 - Package `java.util.function`
 - Lambda expressions, method reference expressions
 - Functional interfaces, targeted function type
- Java 8 streams for bulk data
 - Package `java.util.stream`
- High-level parallel programming
 - Streams: primes, top-down parsing, n-queens, ...
 - Array parallel prefix operations
 - Class `java.util.Arrays` static methods

Materials

- Sestoft: *Java Precisely* 3rd ed., MIT Press 2016
 - § 11.13: Lambda expressions
 - § 11.14: Method reference expressions
 - § 23: Functional interfaces
 - § 24: Streams for bulk data
 - § 25: Class Optional<T>
- Book examples are called Example154.java etc
 - Get them from the book homepage
<http://www.itu.dk/people/sestoft/javaprecisely/>

Functional and stream programming in Java 8 (2014)

- Lambda expressions
`(String s) -> s.length`
- Method reference expressions
`String::length`
- Functional interfaces
`Function<String,Integer>`
- Streams for bulk data processing
`Stream<String> ss = ...`
`Stream<Integer> is = ss.map(String::length)`
- Parallel stream processing
`is = ss.parallel().map(String::length)`
- Parallel array operations
`Arrays.parallelSetAll(arr, i -> sin(i/PI/100.0))`
`Arrays.parallelPrefix(arr, (x, y) -> x+y)`

Functional programming in Java

- *Immutable data* instead of objects with state
- *Recursion* instead of loops
- *Higher-order functions* that may
 - take functions as argument
 - return functions as result

```
class FunList<T> {  
    final Node<T> first;  
    protected static class Node<U> {  
        public final U item;  
        public final Node<U> next;  
        public Node(U item, Node<U> next) { ... }  
    }  
    ...  
}
```

Immutable
list of T


Example154.java

Immutable data

- FunList<T>, linked lists of nodes

```
class FunList<T> {  
    final Node<T> first;  
    protected static class Node<U> {  
        public final U item;  
        public final Node<U> next;  
        public Node(U item, Node<U> next) { ... }  
    }  
    static <T> FunList<T> cons(T item, FunList<T> list) {  
        return new FunList<T>(new Node<T>(item, list.first));  
    }  
}
```

Example154.java



FunList<Integer>


Immutability:

No changes to existing data

```
FunList<Integer> empty = new FunList<>(null),  
list1 = cons(9, cons(13, cons(0, empty))),  
list2 = cons(7, list1),  
list3 = cons(8, list1),  
list4 = list1.insert(1, 12),  
list5 = list2.removeAt(3);
```

Insert 12 before
element number 1

Example154.java



No changes
to existing
data

Recursion in insert

```
public FunList<T> insert(int i, T item) {  
    return new FunList<T>(insert(i, item, this.first));  
}  
  
static <T> Node<T> insert(int i, T item, Node<T> xs) {  
    return i == 0 ? new Node<T>(item, xs)  
        : new Node<T>(xs.item, insert(i-1, item, xs.next));  
}
```

Example154.java

- “If *i* is zero, put *item* in a new node, and let its tail be the old list *xs*”
- “Otherwise, put the first element of *xs* in a new node, and let its tail be the result of inserting *item* in position *i-1* of the tail of *xs*”

Functional programming and immutable data: Pros and cons

- Immutability leads to more data allocation
 - Takes time and space
 - But modern allocators, garbage collectors are fast
- Immutable data can be safely shared
 - May actually reduce amount of allocation
- Immutable data are automatically threadsafe
 - No (other) thread can destructively update it; very good for parallel programming
 - And also due to visibility effects of `final` modifier

Subtle Java point;
Java Precisely § 20.5.2

Mutable and/or shared memory

- Mutable: Data can be updated, reassigned
- Shared: Data can be accessed from two threads

	Shared	Unshared
Mutable	Imperative: C, Java, ...; threads, locks, semaphore	Message passing: Erlang, Scala Akka
Immutable	Functional: Haskell, F#, Java parallel func. streams	

Languages,
techniques

	Shared	Unshared
Mutable	Difficult, race conditions	Easy
Immutable	Easy	Easy

Conceptual
difficulty

	Shared	Unshared
Mutable	Slow (cache states M \leftrightarrow I)	Fast (cache state M)
Immutable	Fast (cache state S)	Fast (cache state E)

MESI cache
hardware

Java lambda expressions 1

Example64.java

- One argument lambda expressions:

```
Function<String,Integer>  
fsi1 = s -> Integer.parseInt(s);
```

```
... fsi1.apply("004711") ...
```

Calling the function

Function that takes a string s and parses it as an integer

```
Function<String,Integer>  
fsi2 = s -> { return Integer.parseInt(s); },  
fsi3 = (String s) -> Integer.parseInt(s);
```

Same, written in other ways

- Two-argument lambda expressions:

```
BiFunction<String,Integer,String>  
fsis1 = (s, i) -> s.substring(i, Math.min(i+3, s.length()));
```

Java method reference expressions

```
BiFunction<String,Integer,Character> charat  
= String::charAt;
```

Same as (s,i) -> s.charAt(i)

```
System.out.println(charat.apply("ABCDEF", 1));
```

Example67.java

```
Function<String,Integer> parseInt = Integer::parseInt;
```

Same as fsi1, fsi2 and fsi3

```
Function<Integer,Character> hex1  
= "0123456789ABCDEF"::charAt;
```

Conversion to hex digit

Class and array constructors

```
Function<Integer,C> makeC = C::new;  
Function<Integer,Double[]> make1DArray = Double[]::new;
```

Targeted function type (TFT)

- A lambda expression or method reference expression *does not have a type in itself*
- Therefore must have a *targeted function type*
- Lambda or method reference must appear as
 - Assignment right hand side:
 - `Function<String,Integer> f = Integer::parseInt;`
 - Argument to call:
 - `stringList.map(Integer::parseInt)`
 - In a cast:
 - `(Function<String,Integer>) Integer::parseInt`
 - Argument to **return** statement:
 - `return Integer::parseInt;`

TFT

map's argument type is TFT

TFT

Enclosing method's
return type is TFT

Functions as arguments: map

```
public <U> FunList<U> map(Function<T,U> f) {  
    return new FunList<U>(map(f, first));  
}  
static <T,U> Node<U> map(Function<T,U> f, Node<T> xs) {  
    return xs == null ? null  
        : new Node<U>(f.apply(xs.item), map(f, xs.next));  
}
```

Example154.java

- Function **map** encodes general behavior
 - Transform each list element to make a new list
 - Argument **f** expresses the specific transformation
- Just as in Haskell, Scala, F#, Scheme, ...
- Same effect as OO “template method pattern”

Calling map

7 9 13

```
FunList<Double> list8 = list5.map(i -> 2.5 * i);
```

17.5 22.5 32.5

```
FunList<Boolean> list9 = list5.map(i -> i < 10);
```

true true false

Java 8 functional interfaces

- A *functional interface* has exactly one abstract method

```
interface Function<T,R> {  
    R apply(T x);  
}
```

Type of functions
from T to R

C#: Func<T,R>

F#: T -> R

```
interface Consumer<T> {  
    void accept(T x);  
}
```

Type of functions
from T to void

C#: Action<T>

F#: T -> unit

(Too) many Java functional interfaces

Interface	Sec.	Function Type	Single Abstract Method Signature
One-Argument Functions and Predicates			
Function<T,R>	23.5	T -> R	R apply(T)
UnaryOperator<T>	23.6	T -> T	T apply(T)
Predicate<T>	23.7	T -> boolean	boolean test(T)
Consumer<T>	23.8	T -> void	void accept(T)
Supplier<T>	23.9	void -> T	T get()
Runnable		void -> void	void run()
Two-Argument Functions and Predicates			
BiFunction<T,U,R>	23.10	T * U -> R	R apply(T, U)
BinaryOperator<T>	23.11	T * T -> T	T apply(T, T)
BiPredicate<T,U>	23.7	T * U -> boolean	boolean test(T, U)
BiConsumer<T,U>	23.8	T * U -> void	void accept(T, U)
Primitive-Type Specialized Versions of the Generic Functional Interfaces			
DoubleToIntFunction	23.5	double -> int	int applyAsInt(double)
DoubleToLongFunction	23.5	double -> long	long applyAsLong(double)
IntToDoubleFunction	23.5	int -> double	double applyAsDouble(int)
IntToLongFunction	23.5	int -> long	long applyAsLong(int)
LongToDoubleFunction	23.5	long -> double	double applyAsDouble(long)
LongToIntFunction	23.5	long -> int	int applyAsInt(long)
DoubleFunction<R>	23.5	double -> R	R apply(double)
IntFunction<R>	23.5	int -> R	R apply(int)
LongFunction<R>	23.5	long -> R	R apply(long)
ToDoubleFunction<T>	23.5	T -> double	double applyAsDouble(T)
ToIntFunction<T>	23.5	T -> int	int applyAsInt(T)
ToLongFunction<T>	23.5	T -> long	long applyAsLong(T)
ToDoubleBiFunction<T,U>	23.10	T * U -> double	double applyAsDouble(T, U)
ToIntBiFunction<T,U>	23.10	T * U -> int	int applyAsInt(T, U)
ToLongBiFunction<T,U>	23.10	T * U -> long	long applyAsLong(T, U)
DoubleUnaryOperator	23.6	double -> double	double applyAsDouble(double)
IntUnaryOperator	23.6	int -> int	int applyAsInt(int)
LongUnaryOperator	23.6	long -> long	long applyAsLong(long)
DoubleBinaryOperator	23.11	double * double -> double	double applyAsDouble(double, double)
IntBinaryOperator	23.11	int * int -> int	int applyAsInt(int, int)
LongBinaryOperator	23.11	long * long -> long	long applyAsLong(long, long)
DoublePredicate	23.7	double -> boolean	boolean test(double)
IntPredicate	23.7	int -> boolean	boolean test(int)
LongPredicate	23.7	long -> boolean	boolean test(long)
DoubleConsumer	23.8	double -> void	void accept(double)
IntConsumer	23.8	int -> void	void accept(int)
LongConsumer	23.8	long -> void	void accept(long)
ObjDoubleConsumer<T>	23.8	T * double -> void	void accept(T, double)
ObjIntConsumer<T>	23.8	T * int -> void	void accept(T, int)
ObjLongConsumer<T>	23.8	T * long -> void	void accept(T, long)
BooleanSupplier	23.9	void -> boolean	boolean getAsBoolean()
DoubleSupplier	23.9	void -> double	double getAsDouble()
IntSupplier	23.9	void -> int	int getAsInt()
LongSupplier	23.9	void -> long	long getAsLong()

```
interface IntFunction<R> {
    R apply(int x);
}
```

Use instead of
Function<Integer,R>
to avoid (un)boxing

Primitive-type
specialized
interfaces

Primitive-type specialized interfaces for int, double, and long

```
interface Function<T,R> {  
    R apply(T x);  
}
```

```
interface IntFunction<R> {  
    R apply(int x);  
}
```

Why both?

What difference?

```
Function<Integer,String> f1 = i -> "#" + i;  
IntFunction<String> f2 = i -> "#" + i;
```

- Calling `f1.apply(42)` will *box* 42 as Integer
 - Allocating object in heap, takes time and memory
- Calling `f2.apply(42)` avoids boxing, is faster
- Purely for performance

Java streams for bulk data

- `Stream<T>` is a finite or infinite sequence of `T`
 - Possibly lazily generated
 - Possibly parallel
- Stream methods
 - `map`, `flatMap`, `reduce`, `filter`, ...
 - These take functions as arguments
 - Can be combined into pipelines
 - Java optimizes (and parallelizes) the pipelines well
- Similar to
 - Java Iterators, but very different implementation
 - The Enumerable extension methods underlying C#/.NET Language Integrated Query (Linq)

Some stream operations

- `Stream<Integer> s = Stream.of(2, 3, 5)`
- `s.filter(p) = those x where p.test(x) holds`
`s.filter(x -> x%2==0) gives 2`
- `s.map(f) = results of f.apply(x) for x in s`
`s.map(x -> 3*x) gives 6, 9, 15`
- `s.flatMap(f) = a flattening of the streams created by f.apply(x) for x in s`
`s.flatMap(x -> Stream.of(x,x+1)) gives 2,3,3,4,5,6`
- `s.findAny() = some element of s, if any, or else the absent Option<T> value`
`s.findAny() gives 2 or 3 or 5`
- `s.reduce(x0, op) = x0 ⋄ s0 ⋄ ... ⋄ sn` if we write `op.apply(x,y)` as `x ⋄ y`
`s.reduce(1, (x,y)->x*y) gives 1*2*3*5 = 30`

Similar functions are everywhere

- Java stream **map** is called
 - **map** in Haskell, Scala, F#, Clojure
 - **Select** in C# Linq
- Java stream **flatMap** is called
 - **concatMap** in Haskell
 - **flatMap** in Scala
 - **collect** in F#
 - **SelectMany** in C# Linq
 - **mapcat** in Clojure
- Java **reduce** is a special (assoc. op.) case of
 - **fold1** in Haskell
 - **foldLeft** in Scala
 - **fold** in F#
 - **Aggregate** in C# Linq
 - **reduce** in Clojure

Counting primes on Java 8 streams

- A standard Java for loop:

```
int count = 0;
for (int i=0; i<range; i++)
    if (isPrime(i))
        count++;
```

Classical efficient
imperative loop

- Sequential Java 8 stream:

```
IntStream.range(0, range)
    .filter(i -> isPrime(i))
    .count()
```

Pure functional
programming ...

- Parallel Java 8 stream:

```
IntStream.range(0, range)
    .parallel()
    .filter(i -> isPrime(i))
    .count()
```

... and thus
parallelizable and
thread-safe

Performance results (!!)

- Counting the primes in 0 ...99,999

Method	Intel i7 (ms)	AMD Opteron (ms)
Sequential for-loop	9.9	40.5
Sequential stream	9.9	40.8
Parallel stream	2.8	1.7
Best thread-parallel	3.0	4.9
Best task-parallel	2.6	1.9

- Functional streams give the simplest solution
- Nearly as fast as tasks and threads, or faster:
 - Intel i7 (4 cores) speed-up: 3.6 x
 - AMD Opteron (32 cores) speed-up: 24.2 x
 - ARM Cortex-A7 (RP 2B) (4 cores) speed-up: 3.5 x
- The future is parallel – and functional 😊

Side-effect freedom: functional!

- From the `java.util.stream` package docs:

Side-effects

function-type

Side-effects in behavioral parameters to stream operations are, in general, discouraged, as they can often lead to unwitting violations of the statelessness requirement, as well as other thread-safety hazards.

Should say
"catastrophic"

- Java compiler and type system cannot enforce side-effect freedom – unlike Haskell
- Java runtime cannot detect violations

Creating Java streams 1

- Explicitly or from array, collection or map:

```
IntStream is = IntStream.of(2, 3, 5, 7, 11, 13);
```

```
String[] a = { "Hoover", "Roosevelt", ... };  
Stream<String> presidents = Arrays.stream(a);
```

```
Collection<String> coll = ...;  
Stream<String> countries = coll.stream();
```

```
Map<String,Integer> phoneNumbers = ...;  
Stream<Map.Entry<String,Integer>> phones  
    = phoneNumbers.entrySet().stream();
```

Example164.java

- Finite, ordered, sequential, lazily generated

Creating Java streams 2

- Useful special-case streams:
 - `IntStream.range(0, 100_000)`
 - `random.ints(5_000)`
 - `bufferedReader.lines()`
 - `regex.matcher(text).results()`
 - `bitset.stream()`
- Functional iterators for infinite streams
- Imperative generators for infinite streams
- `StreamBuilder<T>`: eager, only finite streams

Example164.java

Creating Java streams 3: generators

- Generating 0, 1, 2, 3, ...

Functional

```
IntStream nats1 = IntStream.iterate(0, x -> x+1);
```

Most efficient (!!),
and parallelizable

Object-oriented
imperative

Example165.java

```
IntStream nats2 = IntStream.generate(new IntSupplier() {  
    private int next = 0;  
    public int getAsInt() { return next++; }  
});
```

Imperative, using final
array for mutable state

```
final int[] next = { 0 };  
IntStream nats3 = IntStream.generate(() -> next[0]++);
```

Streams & floating-point sum

- Eg. compute series sum: $\frac{1}{1} + \frac{1}{2} + \frac{1}{3} + \dots + \frac{1}{N}$
for $N=999,999,999$
- Could make a `DoubleStream`, and use `.sum()`

```
IntStream.range(1, N)
    .mapToDouble(i -> 1.0/i)
    .sum();
```

21.300481501347942

- Or *parallel* `DoubleStream` and `.sum()`

```
IntStream.range(1, N)
    .parallel()
    .mapToDouble(i -> 1.0/i)
    .sum();
```

21.300481501347942

Precise (not exact) is
21.300481501347944

TestStreamSums.java

- What about a good old-fashioned for loop?

Summation with good old for-loops

- Compute series sum: $\frac{1}{1} + \frac{1}{2} + \frac{1}{3} + \dots + \frac{1}{N}$
for $N=999,999,999$

- For-loop, forwards summation

```
double sum = 0.0;  
for (int i=1; i<N; i++)  
    sum += 1.0/i;
```

21.300481501348550

- For-loop, backwards summation

```
double sum = 0.0;  
for (int i=1; i<N; i++)  
    sum += 1.0/(N-i);
```

21.300481501346148

Different!

TestStreamSums.java

- Floating-point is sensitive to summation order
- Java DoubleStream .sum() uses *Kahan summation* to avoid these problems

Streams for top-down parsing 1

- Three forms of grammar Rule:
 - a Literal: parse a
 - r1 r2 ... rn Sequence: parse r1, then r2, ...
 - r1 | r2 Alternative: parse r1 or r2

```
static abstract class Rule {  
    final Rule first;  
  
    public abstract IntStream match(String s, int p);  
  
    public static boolean match(String s) {  
        return first.match(s, 0).anyMatch(pos -> pos == s.length());  
    }  
}  
  
static class Lit extends Rule { ... }  
static class Seq extends Rule { ... }  
static class Alt extends Rule { ... }
```

Stream of integers q
such that the rule can
match s[p..q-1]

Streams for top-down parsing 2

```
static class Lit extends Rule {
    final String val;
    public IntStream match(String s, int pos) {
        return s.startsWith(val, pos) ? IntStream.of(pos+val.length())
                                         : IntStream.empty();
    }
}

static class Seq extends Rule {
    final Rule[] rules;
    public IntStream match(String s, int pos0) {
        IntStream result = IntStream.of(pos0);
        for (int rule : rules)
            result = result.flatMap(pos -> rule.match(s, pos));
        return result;
    }
}

static class Alt extends Rule {
    final Seq alt1, alt2;
    public IntStream match(String s, int pos0) {
        return IntStream.concat(alt1.match(s, pos0), alt2.match(s, pos0));
    }
}
```


Streams for backtracking

- Eg. generate all n-permutations of 0, 1, ..., n-1
n=3: [2,1,0], [1,2,0], [2,0,1], [0,2,1], [0,1,2], [1,0,2]

Set of numbers not
yet used in tail

An incomplete
permutation

```
public static Stream<IntList> perms(BitSet todo, IntList tail) {  
    if (todo.isEmpty())  
        return Stream.of(tail);  
    else  
        return todo.stream().boxed()  
            .flatMap(r -> perms(minus(todo, r), new IntList(r, tail)));  
}
```

Example175.java

```
public static Stream<IntList> perms(int n) {  
    BitSet todo = new BitSet(n); todo.flip(0, n);  
    return perms(todo, null);  
}
```

{ 0, ..., n-1 }

Empty
permutation []

A closer look at generation for n=3

todo

tail

({0,1,2}, [])

r=0

({1,2}, [0])

({2}, [1,0])

({}, [2,1,0])

To result stream

({1}, [2,0])

({}, [1,2,0])

To result stream

r=1

({0,2}, [1])

({2}, [0,1])

({}, [2,0,1])

To result stream

({0}, [2,1])

({}, [0,2,1])


To result stream

r=2


({0,1}, [2])

...


A permutation is a safe rook (tårn) placement on a chessboard




[2, 1, 0]




[1, 2, 0]




[2, 0, 1]



[0, 2, 1]



[0, 1, 2]



[1, 0, 2]

Solutions to the n-queens problem

- For queens, just take diagonals into account:
 - consider only r that are safe for the partial solution

```
public static Stream<IntList> queens(BitSet todo, IntList tail) {  
    if (todo.isEmpty())  
        return Stream.of(tail);  
    else  
        return todo.stream()  
            .filter(r -> safe(r, tail)).boxed()  
            .flatMap(r -> queens(minus(todo, r), new IntList(r, tail)));  
}
```

Diagonal check

.parallel()

Example176.java

```
public static boolean safe(int mid, IntList tail) {  
    return safe(mid+1, mid-1, tail);  
}  
public static boolean safe(int d1, int d2, IntList tail) {  
    return tail==null || d1!=tail.item && d2!=tail.item && safe(d1+1, d2-1, tail.next);  
}
```

- Simple, and parallelizable for free, 3.5 x faster
- Solve or generate sudokus: much the same

Versatility of streams

- Many uses of a stream of solutions

- Print the number of solutions

```
System.out.println(queens(8).count());
```

- Print all solutions

```
queens(8).forEach(System.out::println);
```

- Print an arbitrary solution (if there is one)

```
System.out.println(queens(8).findAny());
```

- Print the 20 first solutions

```
queens(8).limit(20).forEach(System.out::println);
```

- Much harder in an imperative version
- Separation of concerns (Dijkstra): *production* of solutions versus *consumption* of solutions

Example174.java

Streams for quasi-infinite sequences

- van der Corput numbers
 - $1/2, 1/4, 3/4, 1/8, 5/8, 3/8, 7/8, 1/16, \dots$
 - Dense and uniform in interval $[0, 1]$
 - For simulation and finance, Black-Scholes options
- Trick: vd Corput numbers as base-2 fractions
 $0.1, 0.01, 0.11, 0.001, 0.101, 0.011, 0.111, \dots$
are *bit-reversals* of $1, 2, 3, 4, 5, 6, 7, \dots$ in binary

```
public static DoubleStream vanDerCorput() {  
    return IntStream.range(1, 31).asDoubleStream()  
        .flatMap(b -> bitReversedRange((int)b));  
}  
  
private static DoubleStream bitReversedRange(int b) {  
    final long bp = Math.round(Math.pow(2, b));  
    return LongStream.range(bp/2, bp)  
        .mapToDouble(i -> (double)(bitReverse((int)i) >>> (32-b)) / bp);  
}
```

Example183.java

Java 8 stream properties

- Some stream dimensions
 - Finite versus infinite
 - Lazily generated (by `iterate`, `generate`, ...) versus eagerly generated (stream builders)
 - Ordered (`map`, `filter`, `limit` ... preserve element order) versus unordered
 - Sequential (all elements processed on one thread) versus parallel
- Java streams
 - can be lazily generated, like Haskell lists
 - but are *use-once*, unlike Haskell lists
 - reduces risk of space leaks
 - limits expressiveness, harder to compute average ...

How are Java streams implemented?

- Spliterators

```
interface Spliterator<T> {  
    long estimateSize();  
    void forEachRemaining(Consumer<T> action);  
    boolean tryAdvance(Consumer<T> action);  
    void Spliterator<T> trySplit();  
}
```

- Many method calls, well inlined/fused by the JIT

- Parallelization

- Divide stream into chunks using `trySplit`
- Process each chunk in a Java task (Haskell “spark”)
- Run on thread pool using work-stealing queues
- ... thus similar to Haskell `parBuffer/parListChunk`

Parallel (functional) array operations

- Simulating random motion on a line
 - Take n random steps of length at most $[-1, +1]$:

```
double[] a = new Random().doubles(n, -1.0, +1.0)
               .toArray();
```

- Compute the positions at end of each step:

$a[0], a[0]+a[1], a[0]+a[1]+a[2], \dots$

```
Arrays.parallelPrefix(a, (x,y) -> x+y);
```

NB: Updates
array a

- Find the maximal absolute distance from start:

```
double maxDist = Arrays.stream(a).map(Math::abs)
                       .max().getAsDouble();
```

- A lot done, fast, without loops or assignments
 - Just arrays and streams and functions

Example25.java

Array and streams and parallel ...

- Associative array aggregation

```
Arrays.parallelPrefix(a, (x,y) -> x+y) ;
```

- Such operations can be parallelized well
 - So-called *prefix scans* (Blelloch 1990)
- Streams and arrays complement each other:
- Streams: lazy, possibly infinite, non-materialized, use-once, parallel pipelines
- Arrays: eager, always finite, materialized, use-many-times, parallel prefix scans

Some problems with Java streams

- Streams are use-once & have other restrictions
 - Probably to permit easy parallelization
- Hard to create lazy finite streams
 - Probably to allow simple high-performance implementation
- Difficult to control resource consumption
- A single side-effect may mess all up completely
- Sometimes `.parallel()` hurts performance a lot
 - See exercise
 - And strange behavior, in parallel + limit in Sudoku generator
- Laziness in Java easily goes wrong, try-with-resource:

```
static Stream<String> getPageAsStream(String url) throws IOException {  
    try (BufferedReader in  
        = new BufferedReader(new InputStreamReader(  
                                new URL(url).openStream())) {  
        return in.lines();  
    }  
}
```

Example216.java

Closes the Reader eagerly, so any use of the `Stream<String>` causes `IOException: Stream closed`

Useless

A multicore performance mystery

- K-means clustering 2P: Assign – Update – Assign – Update ... till convergence

Pseudocode

```
while (!converged) {  
  let taskCount parallel tasks do {  
    final int from = ..., to = ...;  
    for (int pi=from; pi<to; pi++)  
      myCluster[pi] = closest(points[pi], clusters);  
  }  
  let taskCount parallel tasks do {  
    final int from = ..., to = ...;  
    for (int pi=from; pi<to; pi++)  
      myCluster[pi].addToMean(points[pi]);  
  }  
  ...  
}
```

Assign

Update

TestKMeansSolution.java

Imperative

- Assign: writes a point to `myCluster[pi]`
- Update: calls `addToMean` on `myCluster[pi]`

A multicore performance mystery

- “Improved” version 2Q:
 - call `addToMean` directly on point
 - instead of first writing it to `myCluster` array

```
while (!converged) {  
    let taskCount parallel tasks do {  
        final int from = ..., to = ...;  
        for (int pi=from; pi<to; pi++)  
            closest(points[pi], clusters).addToMean(points[pi]);  
    }  
    ...  
}
```

Performance of k-means clustering

- Sequential: as you would expect, 5% speedup
- Parallel: surprisingly bad!

"Improved"

	2P	2Q	Ratio 2Q/2P
Sequential	4.240	4.019	0.95
4-core parallel i7	1.310	2.234	1.70
24-core parallel Xeon	0.852	6.587	7.70

Bad

Very bad

Time in seconds for 200,000 points, 81 clusters, 1/8/48 tasks, 108 iterations

- Q: WHY is the "improved" code slower?
- A: Cache invalidation and false sharing

The Point and Cluster classes

```
class Point {  
    public final double x, y;  
}
```

```
static class Cluster extends ClusterBase {  
    private volatile Point mean;  
    private double sumx, sumy;  
    private int count;  
    public synchronized void addToMean(Point p) {  
        sumx += p.x;  
        sumy += p.y;  
        count++;  
    }  
    ...  
}
```



Cluster object
layout (maybe)

Parallel streams to the rescue, 3P

3P

```
while (!converged) {  
    final Cluster[] clustersLocal = clusters;  
    Map<Cluster, List<Point>> groups =  
        Arrays.stream(points).parallel()  
            .collect(Collectors.groupingBy(p -> closest(p, clustersLocal)));  
    clusters = groups.entrySet().stream().parallel()  
        .map(kv -> new Cluster(kv.getKey().getMean(), kv.getValue()))  
        .toArray(Cluster[]::new);  
    Cluster[] newClusters =  
        Arrays.stream(clusters).parallel()  
            .map(Cluster::computeMean).toArray(Cluster[]::new);  
    converged = Arrays.equals(clusters, newClusters);  
    clusters = newClusters;  
}
```

Assign

Update

Functional

	2P	2Q	3P stream
Sequential	4.240	4.019	5.353
4-core parallel i7	1.310	2.234	1.350
24-core parallel Xeon	0.852	6.587	0.553

!!!

Time in seconds for 200,000 points, 81 clusters, 1/8/48 tasks, 108 iterations

Mutable and/or shared memory

- Mutable: Data can be updated, reassigned
- Shared: Data can be accessed from two threads

	Shared	Unshared
Mutable	Imperative: C, Java, ...; threads, locks, semaphore	Message passing: Erlang, Scala Akka
Immutable	Functional: Haskell, F#, Java parallel func. streams	

Languages,
techniques

	Shared	Unshared
Mutable	Difficult, race conditions	Easy
Immutable	Easy	Easy

Conceptual
difficulty

	Shared	Unshared
Mutable	Slow (cache states M \leftrightarrow I)	Fast (cache state M)
Immutable	Fast (cache state S)	Fast (cache state E)

MESI cache
hardware

Materials

- Reading

- Java Precisely 3rd ed. § 11.13, 11.14, 23, 24, 25
- Optional:
 - <http://www.itu.dk/people/sestoft/papers/benchmarking.pdf>
 - <http://www.itu.dk/people/sestoft/papers/cpucache-20170319.pdf>



"A multicore performance mystery"

- Exercises

- Extend immutable list class with functional programming; use parallel array operations; use streams of words and streams of numbers
- Alternatively: Make a faster and more scalable k-means clustering implementation, if possible, in any language