



# Composable Memory Transactions in Haskell

Parallel Functional Programming

John Hughes

# Recall Race Conditions...



**1**

**1**



# Transaction

“a group of related actions that need to be performed as a single action”



- **Single global lock**  
Inefficient!
- **Multiple locks**  
Inefficient!  
Error prone!

# Database solution: optimistic concurrency

- Transactions are often *independent*
- Run them in parallel, and *detect conflicts*
  - A conflicted transaction is automatically retried until it succeeds

# Software Transactional Memory

- Optimistic concurrency for updates to memory

# Composable Memory Transactions

Tim Harris

Simon Marlow

Simon Peyton Jones

Maurice Herlihy

Microsoft Research

7 J J Thomson Avenue, Cambridge, UK, CB3 0FB

{tharris,simonmar,simonpj,t-maherl}@microsoft.com

## ABSTRACT

Writing concurrent programs is notoriously difficult, and is of increasing practical importance. A particular source of concern is that even correctly-implemented concurrency abstractions cannot be composed together to form larger abstractions. In this paper we present a new concurrency model, based on *transactional memory*, that offers far richer composition. All the usual benefits of transactional memory are present (e.g. freedom from deadlock), but in addition we describe new modular forms of *blocking* and *choice* that have been inaccessible in earlier work.

**Categories and Subject Descriptors:** D.1.3 [Programming Techniques]: Concurrent Programming – Parallel programming

**General Terms:** Algorithms, Languages

**Keywords:** Non-blocking algorithms, locks, transactions

## 1. INTRODUCTION

Concurrent programming is notoriously tricky. Current lock-based abstractions are difficult to use and make it hard to design computer systems that are reliable and scalable. Furthermore, systems built using locks are difficult to compose without knowing about their internals.

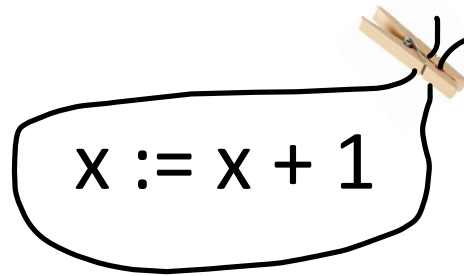
- We re-express the ideas of transactional memory in the setting of Concurrent Haskell (Section 3). This is much more than a routine “port” into a new setting. As we show, STM can be expressed particularly elegantly in a declarative language, and we are able to use Haskell’s type system to give far stronger guarantees than are conventionally possible. Furthermore transactions are compositional: small transactions can be glued together to form larger transactions.

- We present a new, modular form of blocking, which appears to the programmer as a simple function called **retry** (Section 3.2). Unlike most existing approaches, the programmer does not have to identify the condition under which the transaction will complete: **retry** can occur at any time, blocking it until the condition becomes possible.

- The **retry** function can be composed in a way that provides **orElse**, which allows a programmer to specify *alternatives*, so that the second alternative is tried if the first fails (Section 3.4). This ability allows threads to wait for many things at once, like the Unix **select** system call – except that **orElse** composes well, whereas **select** does not. It turns out that **orElse** requires the order

1112  
citations

# STM in Haskell



`x := x + 1`

**IO Monad**

`x :: IORef Int`

```
do n <- readIORef x
   writeIORef x (n+1)
```

**STM Monad**

`x :: TVar Int`

```
do n <- readTVar x
   writeTVar x (n+1)
```



# STM Types

- `newIORef :: a -> IO (IORef a)`
- `readIORef :: IORef a -> IO a`
- `writeIORef :: IORef a -> a -> IO ()`




*put together  
into IO actions*

- `forkIO :: IO () -> IO ThreadId`

- `newTVar :: a -> STM (TVar a)`
- `readTVar :: TVar a -> STM a`
- `writeTVar :: TVar a -> a -> STM ()`

*put together  
into STM actions*

- `atomically :: STM a -> IO a`





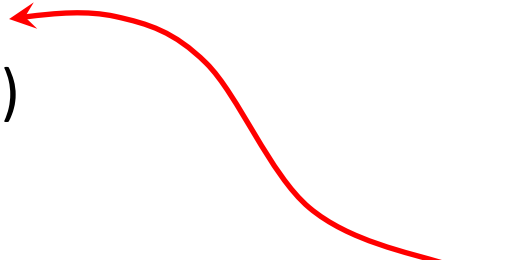
# Example: transferring money

**newtype** Account = Account (TVar Int)

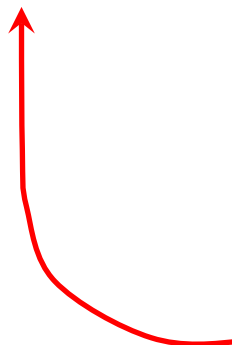
transfer n (Account a) (Account b) = **do**

  modifyTVar' b (+n)

  modifyTVar' a (+(-n))



*add n to  
contents of b*



*evaluate the contents  
of the reference*

transfer :: Int -> Account -> Account -> **STM** ()

# Example: exchanging money

atomically \$ **do**


transfer 1000 aSEK bSEK  
transfer 97 bUSD aUSD

} STM } IO

## Implementation

aSEK	<del>3000</del>
bSEK	<del>20000</del>
aUSD	<del>90</del>
bUSD	<del>2000</del>

Log
bSEK = 20000
bSEK := 21000
aSEK = 5000
aSEK := 4000
bUSD = 2000
bUSD := 1903
aUSD = 0
aUSD := 97



# Example: exchanging money

atomically \$ **do**

transfer 1000 aSEK bSEK  
transfer 97 bUSD aUSD


} **STM** } **IO**

## Implementation

aSEK	<b>6000</b>
bSEK	20000
aUSD	0
bUSD	2000



Log
bSEK = 20000
bSEK := 21000
aSEK = 5000
aSEK := 4000
bUSD = 2000
bUSD := 1903
aUSD = 0
aUSD := 97



# Blocking transactions

```
newtype Account = Account (TVar Int)
```


```
transfer n (Account a) (Account b) = do
```

```
  m <- readTVar a
```

```
  if m < n then retry else do
```

```
    modifyTVar' b (+n)
```

```
    writeTVar a (m-n)
```



*Discard this execution & **block**  
the transaction to be retried  
later...*

*...once one of the variables  
read in the log changes*

# Example: when would this block?

atomically \$ **do**

transfer 1000 aSEK bSEK

transfer 97 bUSD aUSD

**If bUSD contains too little, *when is the transaction retried?***

# Example: paying from two accounts

atomically \$ **do**

transfer 500    aSEK   bSEK

transfer 500    cSEK   bSEK

transfer 97     bUSD   aUSD

# Example: paying from one of two accounts

atomically \$ **do**

transfer 1000 aSEK bSEK `orElse` transfer 1000 cSEK bSEK  
transfer 97 bUSD aUSD



*Do the first, but if it retries, then*

- *Undo its writes*
- *Try the second*
- *If that retries too, wake up when **any** Tvar read changes*

# orElse generalises select

```
wait for A...  
`orElse`  
wait for B...  
`orElse`  
wait for C...  
...
```



*Wait for **any**  
combination of events,  
not just a file descriptor*



# Algebra of ``orElse``

- (``orElse``) is associative, but *not* commutative  
    `return 1 `orElse` return 2 /=`  
    `return 2 `orElse` return 1`
- `retry `orElse` m`  
    `= m`  
    `= m `orElse` retry`

# Algebra of **retry**

- **retry**  $\gg=$  **f** = **retry**

Nothing after a **retry** matters

- **m**  $\gg=$  (**\x -> retry**) = **retry**

Nothing *before* a **retry** matters either

(except to determine when a transaction needs rescheduling)

# STM can simulate other primitives

- E.g. one-place buffer

*new type, same  
representation*



```
newtype Buffer a = Buffer (TVar (Maybe a))
```

```
newBuffer = Buffer <$> newTVar Nothing
```

# STM can simulate other primitives

- E.g. one-place buffer

```
send (Buffer t) a = do
  m <- readTVar t
  case m of
    Nothing -> do
      writeTVar t (Just a)
      return a
    Just _ ->
      retry
```

# STM can simulate other primitives

- E.g. one-place buffer

```
receive (Buffer t) = do
  m <- readTVar t
  case m of
    Nothing ->
      retry
    Just a -> do
      writeTVar t Nothing
      return a
```

# What does this do?

do

```
buf1 <- atomically newBuffer
buf2 <- atomically newBuffer
atomically (send buf1 "hello")
atomically (send buf2 "goodbye")
x <- atomically (send buf1 "hohoho"
                  `orElse` receive buf2)
y <- atomically (receive buf1
                  `orElse` receive buf2)
return (x,y)
```

# What does this do?

```
atomically $ do
  transfer 25 wallet cokeMachine
  Just <$> receive cokeMachineSlot
`orElse`
  return Nothing
```

# Example: Futures

```
newtype Future a = Future (TVar (Maybe a))
```

```
future io = do start a task whose result will be available in the future
  v <- atomically $ newTVar Nothing
  forkIO $ do a <- io
              atomically $ writeTVar v (Just a)
  return (Future v)
```

```
await (Future t) = the future is now
  atomically $ do m <- readTVar t
                  case m of
                    Nothing -> retry
                    Just a   -> return a
```



# Example: Snapshotting a state

```
snapshot :: Eq f => STM s -> (s -> f) -> IO (TVar s)
snapshot getState fingerprint = do
  s <- atomically $ getState
  v <- atomically $ newTVar s
  forkIO $ loop v (fingerprint s)
  return v
where loop v print = do
  print' <- atomically $ do
    s' <- getState
    let print' = fingerprint s'
    if print'==print then retry else do
      writeTVar v s'
      return print'
  loop v print'
```

# Benchmark: binary trees

*strict—to  
simplify  
benchmarking*



```
data Tree k v = Leaf
              | Branch !(Tree k v) k v !(Tree k v)
  deriving (Eq, Show, Generic)
```

```
find k Leaf = Nothing
find k (Branch l k' v' r)
  | k < k' = find k l
  | k ==k' = Just v'
  | k > k' = find k r
```

```
insert k v Leaf = Branch Leaf k v Leaf
insert k v (Branch l k' v' r)
  | k < k' = Branch (insert k v l) k' v' r
  | k ==k' = Branch l k v r
  | k > k' = Branch l k' v' (insert k v r)
```

# Benchmark: binary trees

```
delete k Leaf = Leaf
delete k (Branch l k' v' r)
  | k < k' = Branch (delete k l) k' v' r
  | k == k' = merge l r
  | k > k' = Branch l k' v' (delete k r)
```

```
merge Leaf t = t
merge t Leaf = t
```

```
merge (Branch l k v r) (Branch l' k' v' r') =
  Branch l k v (Branch (merge r l') k' v' r')
```

# Simplest imperative version

```
newTree = atomically $ newTVar Leaf
```

```
insertI t k v = atomically $  
    modifyTVar' t (insert k v)
```

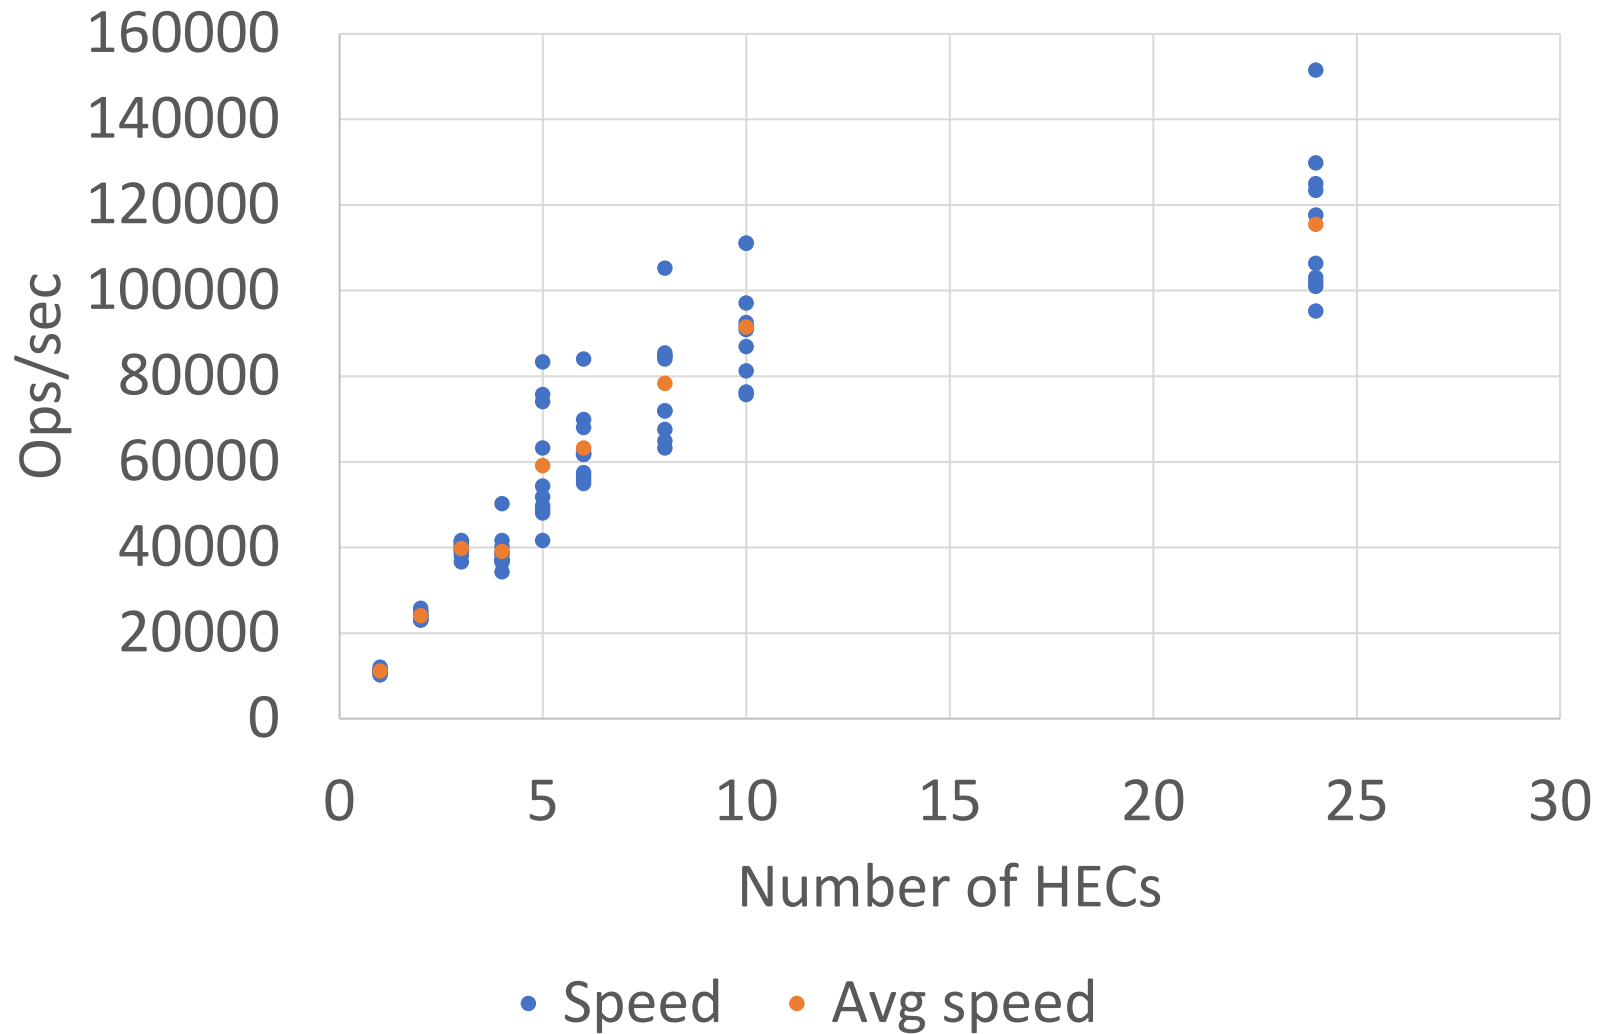
```
deleteI t k    = atomically $  
    modifyTVar' t (delete k)
```

```
findI    t k    = find k <$> atomically $  
    readTVar t
```


# Benchmark

- 500,000 random operations  
*(all performed in parallel)*
- find::insert::delete is 1::1::1
- Keys selected uniformly from the range 1..1000  
*(each key appears ~500 times)*
- Total Haskell run-time is the measure
- Average of 10 runs with each number-of-cores
- Hardware: core i9 with 16 cores (8+8)/24 threads

# Operations per second



# Speedup over 1 HEC



# A transactional version

```
type TTree k v = TVar (TNode k v)
```

```
data TNode k v = TLeaf  
                | TBranch (TTree k v) k v (TTree k v)
```

```
findT t k = do  
  n <- readTVar t  
  case n of  
    TLeaf ->  
      return Nothing  
    TBranch l k' v' r  
      | k < k' -> findT l k  
      | k == k' -> return (Just v')  
      | k > k' -> findT r k
```



```
insertT t k v = do
  n <- readTVar t
  case n of
    TLeaf -> do
      l <- newTVar TLeaf
      r <- newTVar TLeaf
      writeTVar t (TBranch l k v r)
    TBranch l k' v' r
      | k < k' -> insertT l k v
      | k ==k' -> writeTVar t (TBranch l k v r)
      | k > k' -> insertT r k v
```

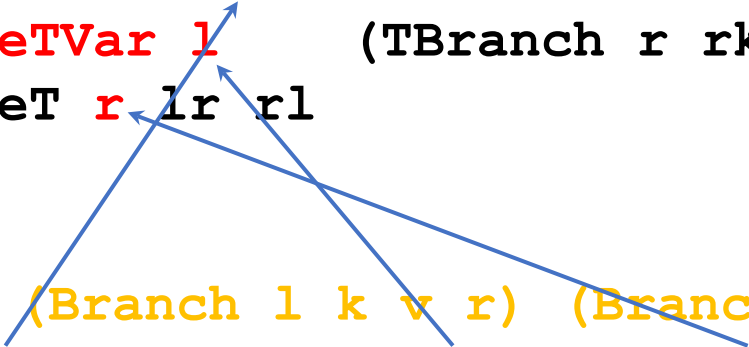
```
deleteT t k = do
  n <- readTVar t
  case n of
    TLeaf ->
      return ()
    TBranch l k' v' r
      | k < k' -> deleteT l k
      | k ==k' -> mergeT t l r
      | k > k' -> deleteT r k
```

```

mergeT dest l r = do
  ln <- readTVar l
  rn <- readTVar r
  case ln of
    TLeaf ->
      writeTVar dest rn
    TBranch ll lk lv lr ->
      case rn of
        TLeaf ->
          writeTVar dest ln
        TBranch rl rk rv rr -> do
          writeTVar dest (TBranch ll lk lv l)
          writeTVar l      (TBranch r rk rv rr)
          mergeT r lr rl

```

$\text{merge } (\text{Branch } l \ k \ v \ r) \ (\text{Branch } l' \ k' \ v' \ r') =$   
 $\text{Branch } l \ k \ v \ (\text{Branch } (\text{merge } r \ l') \ k' \ v' \ r')$



# Does this work?

*List of 'actions'*

```
prop_TTree acts = ioProperty $ do
  t <- atomically $ newTVar TLeaf
  agrees t Leaf acts
```

*corresponding  
pure Tree*

```
where agrees t t' [] = do
  t_ <- atomically $ treeOf t
  return (t_ == t')
```

*TVar tree*

```
  agrees t t' (Find k:acts) = do
    v <- atomically $ findT t k
    if v==find k t'
    then agrees t t' acts
    else return (v==find k t')
```

*convert a TVar tree  
into a pure one*

```
  agrees t t' (Insert k v:acts) = do
    atomically $ insertT t k v
    agrees t (insert k v t') acts
```

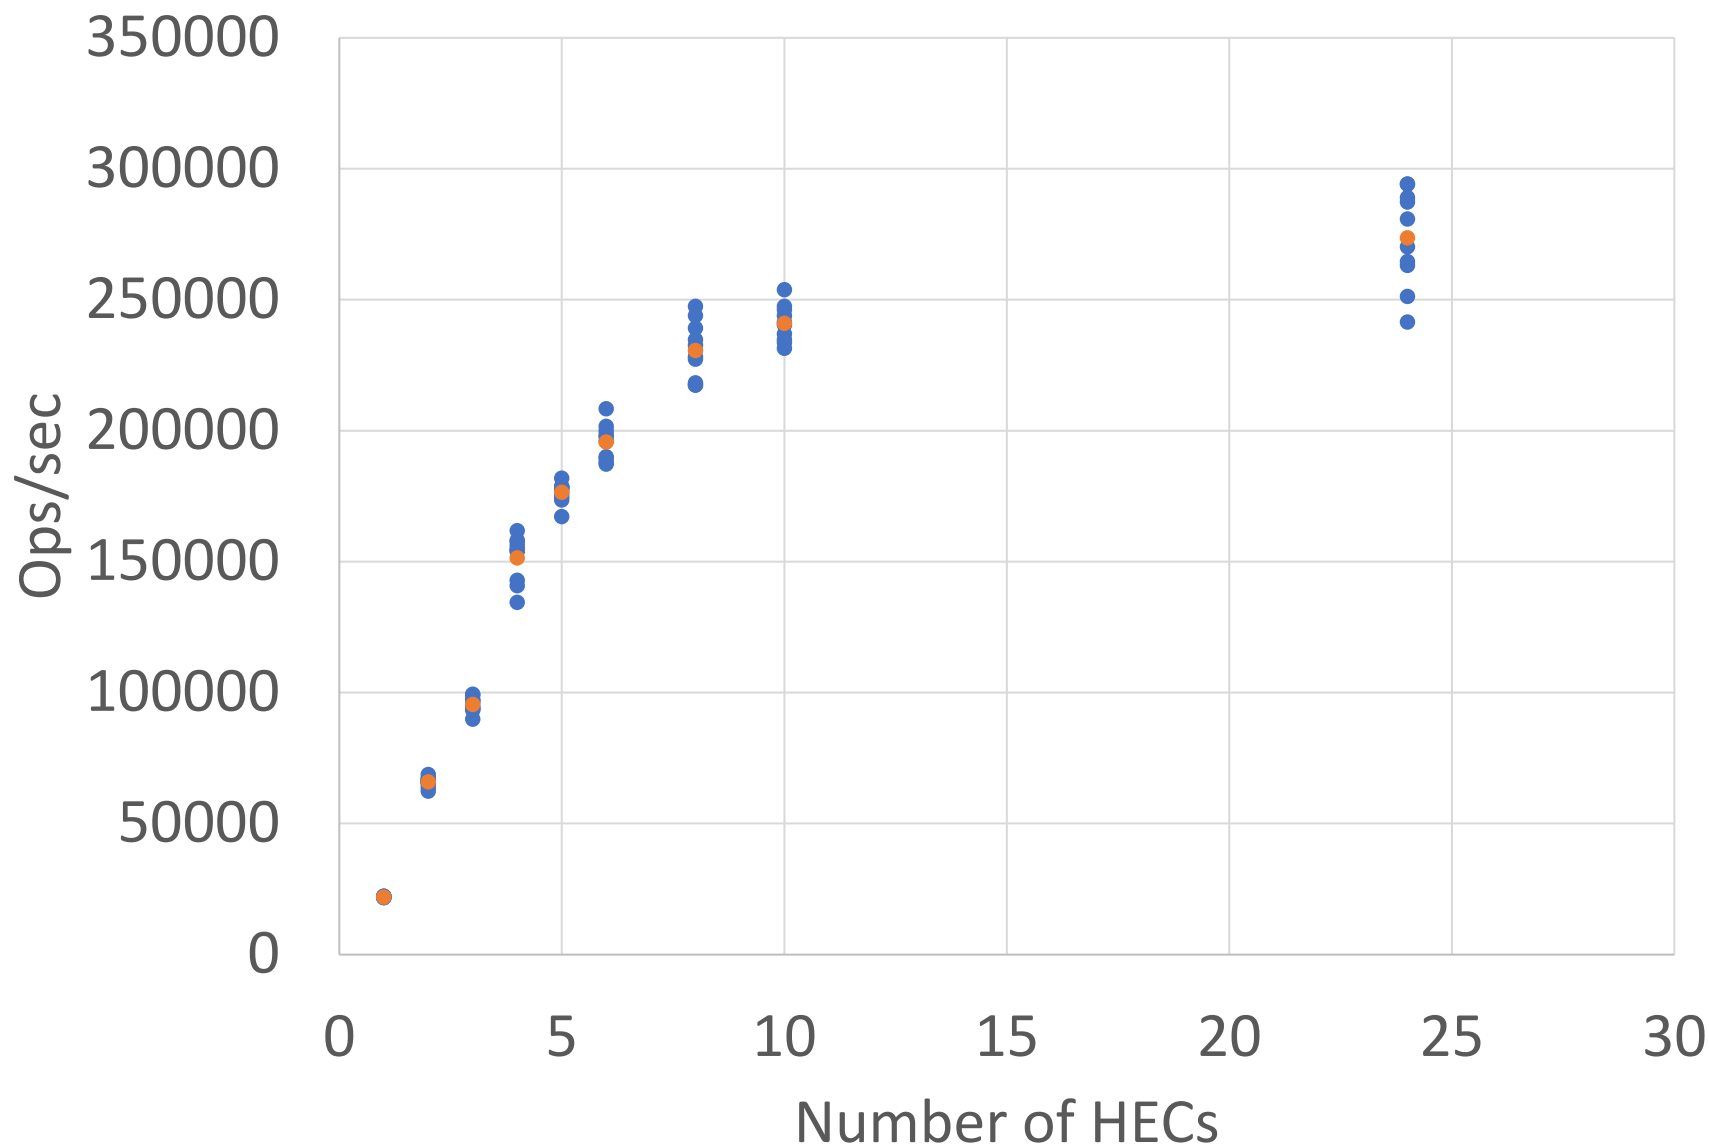
```
  agrees t t' (Delete k:acts) = do
    atomically $ deleteT t k
    agrees t (delete k t') acts
```

# One bug...

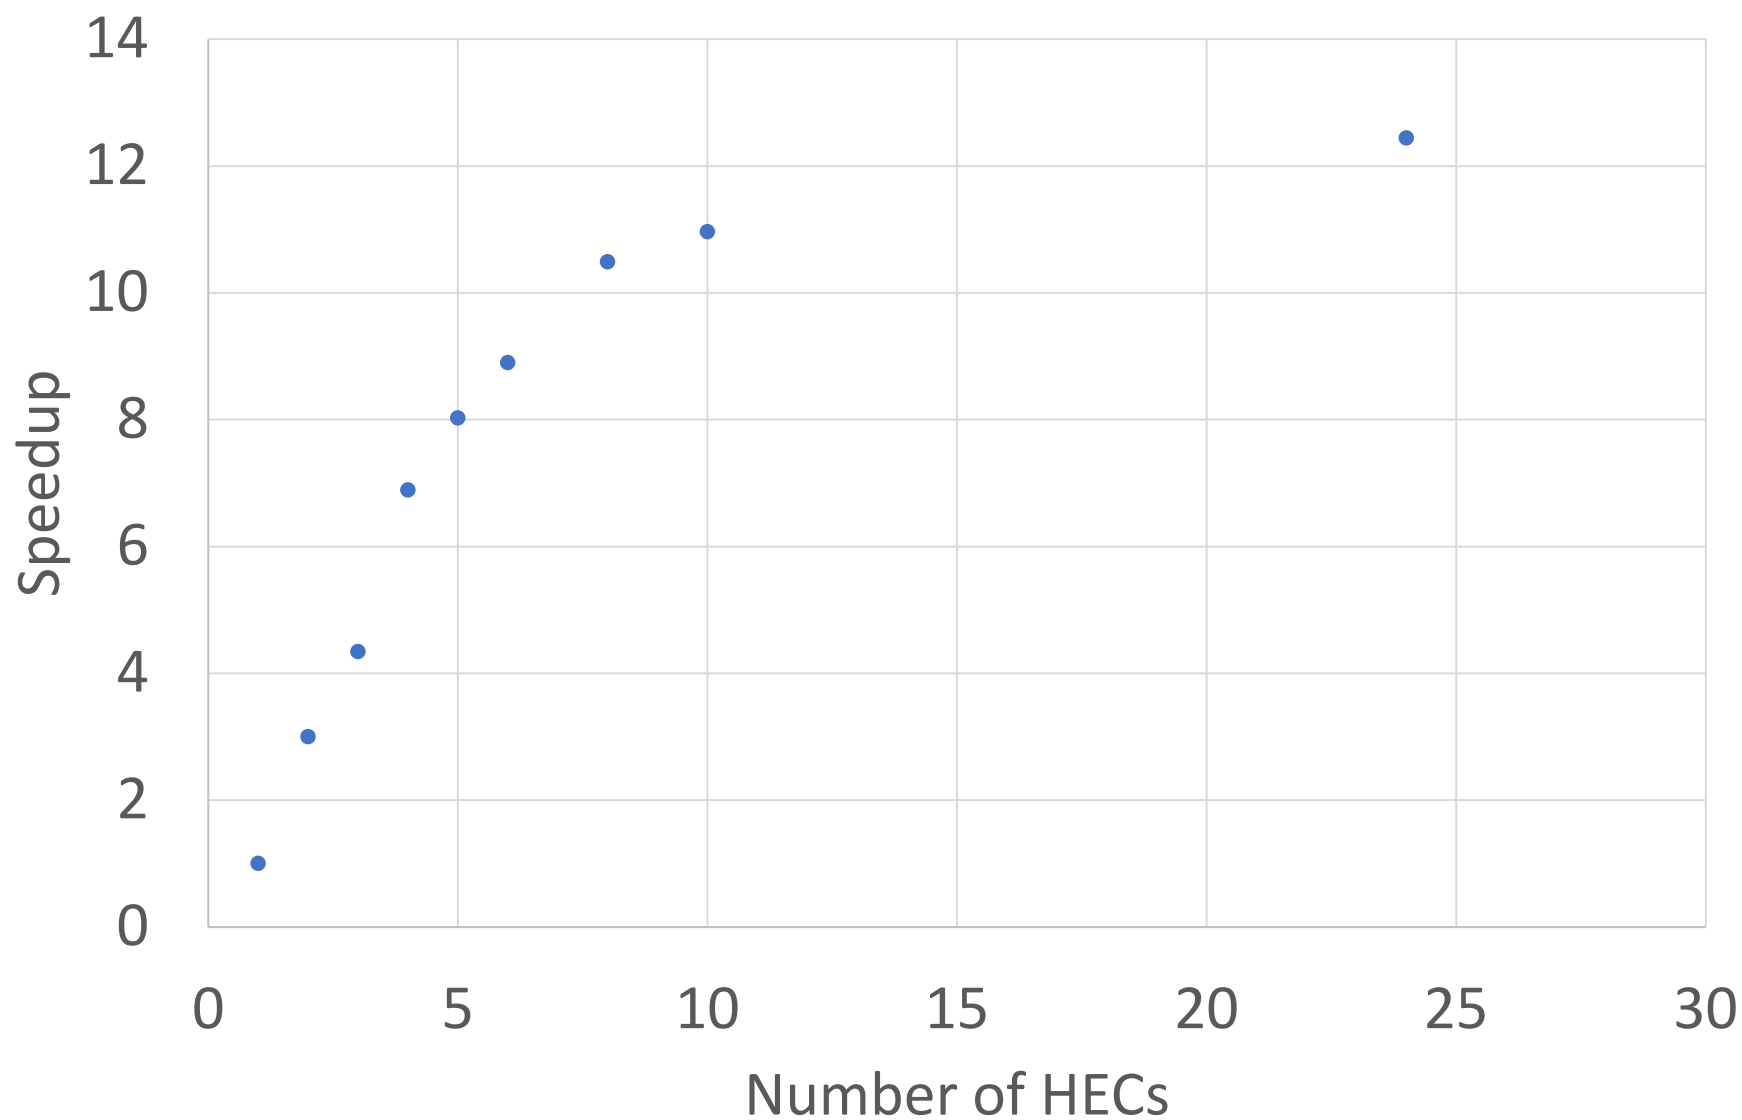
```
*Main> quickCheck prop_TTree
*** Failed! Falsified (after 22 tests and 11 shrinks):
[Insert 0 0,Insert 1 0,Delete 1]
Branch Leaf 0 0 (Branch Leaf 1 0 Leaf)
  /= Branch Leaf 0 0 Leaf
```

```
deleteT t k = do
  n <- readTVar t
  case n of
    TLeaf ->
      return ()
    TBranch l k' v' r
      | k < k' -> deleteT l k'
      | k ==k' -> mergeT t l r
      | k > k' -> deleteT r k'
```

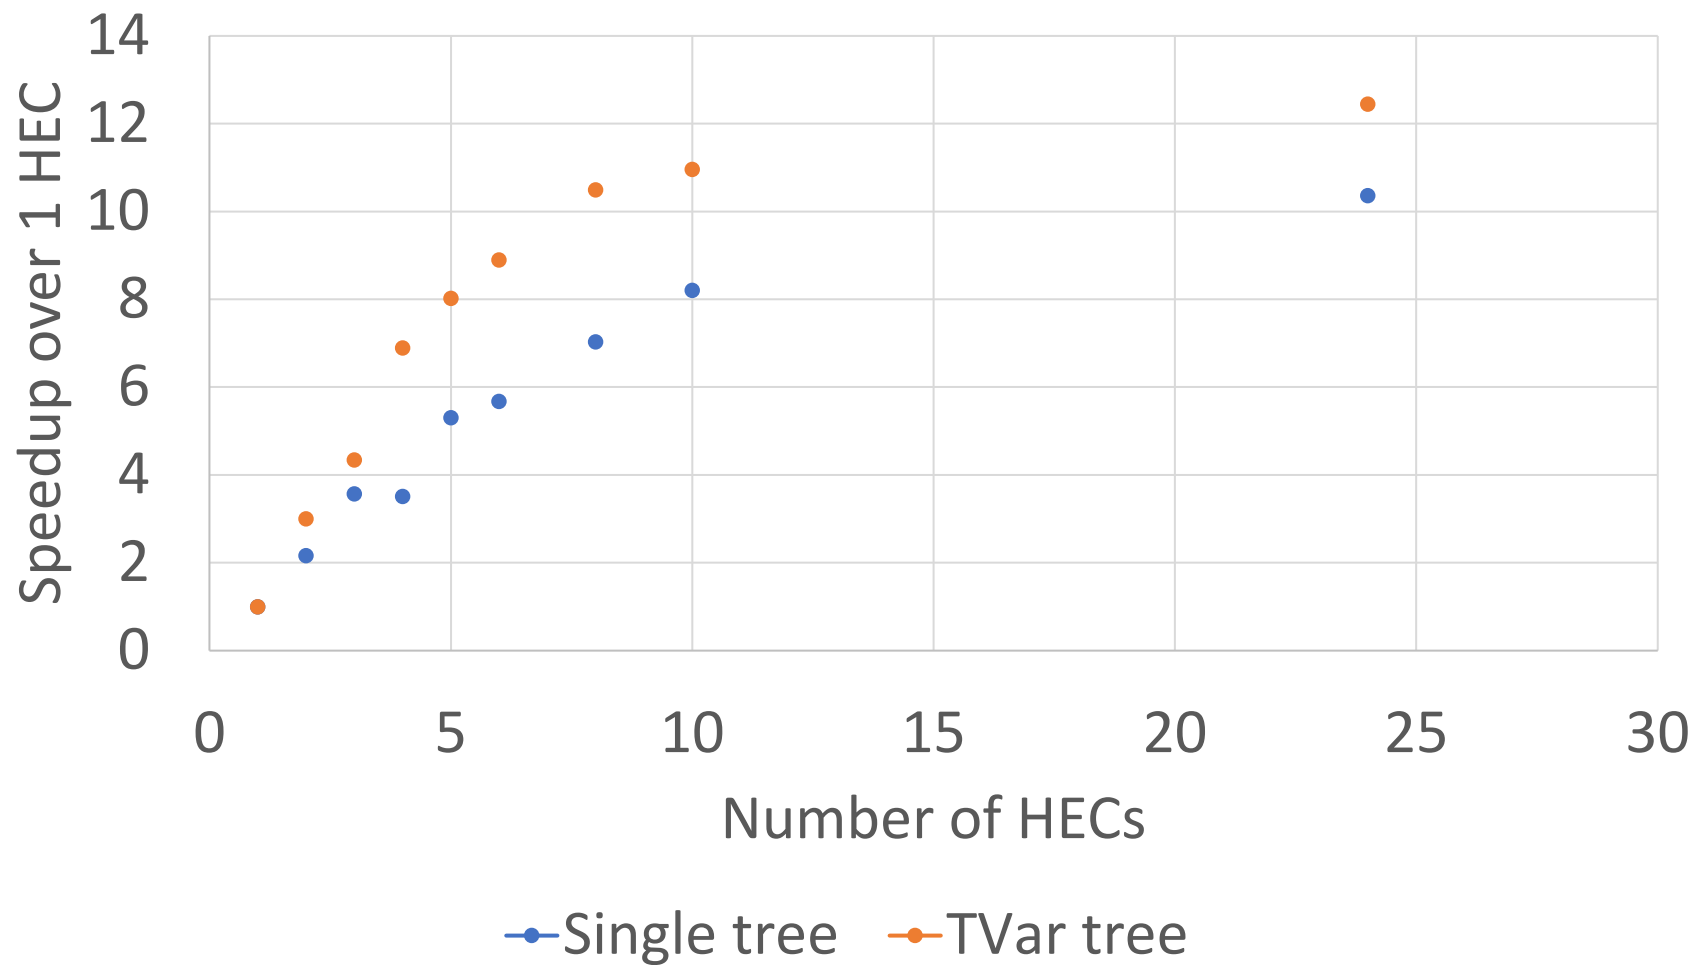
# Operations per second



# Speedup over 1 HEC




# Comparison of pure trees & Tvar trees

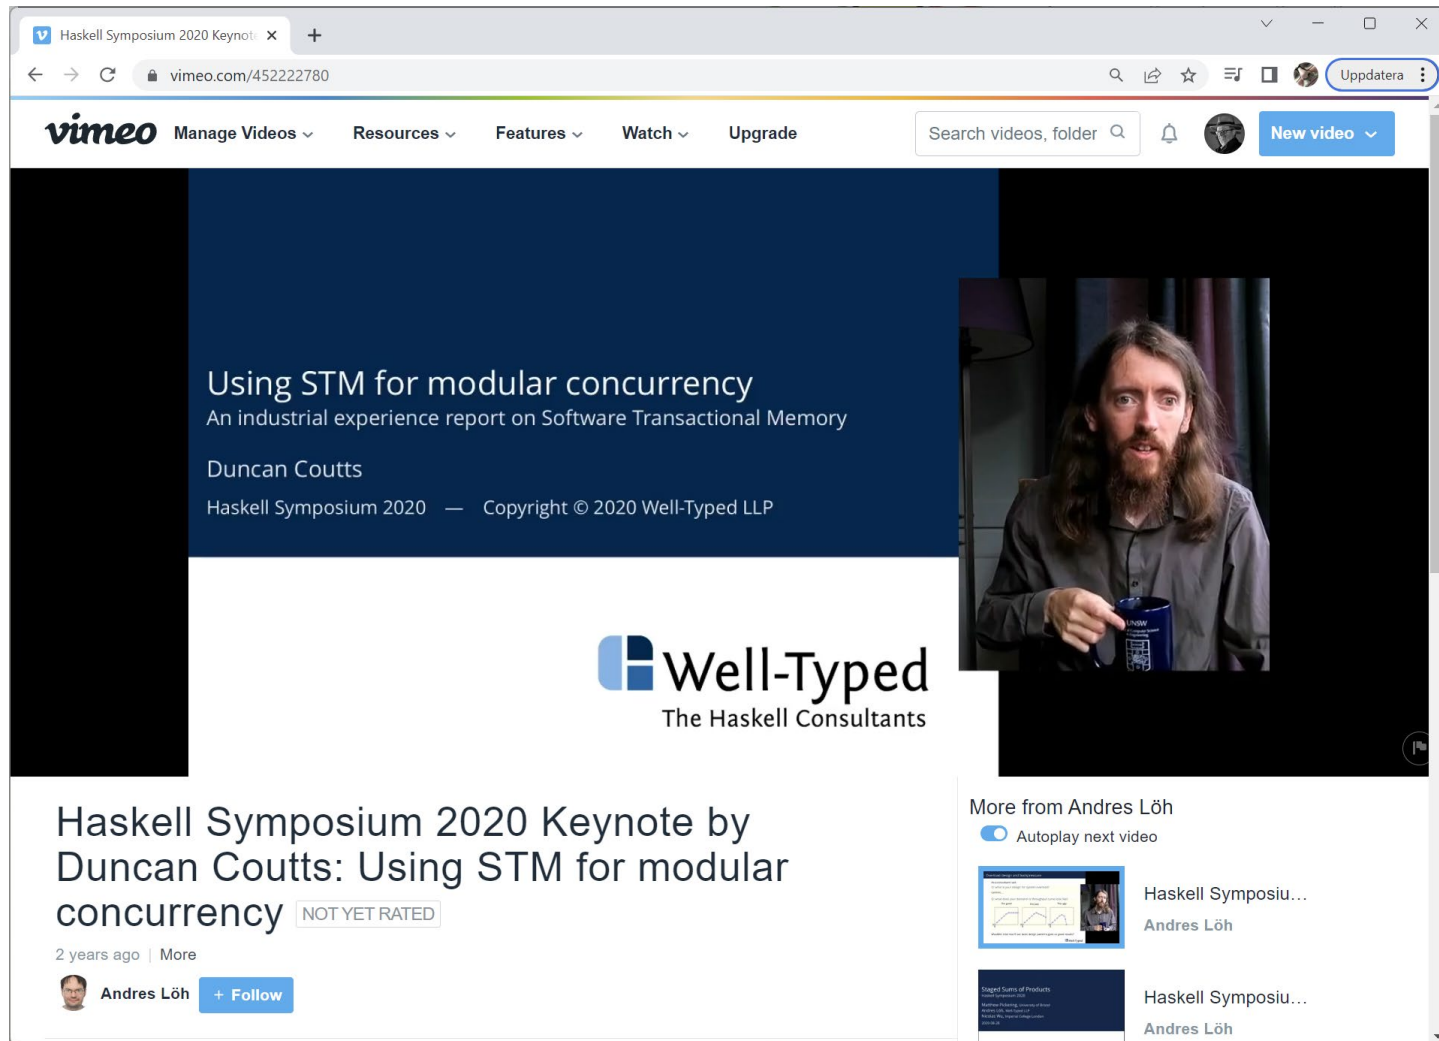




# Performance of pure trees and Tvar trees



# System Design with STM



The screenshot shows a web browser window with a single tab titled "Haskell Symposium 2020 Keynote". The address bar displays "vimeo.com/452222780". The Vimeo navigation bar includes links for "Manage Videos", "Resources", "Features", "Watch", and "Upgrade", along with a search bar and a "New video" button. The video player area features a dark blue background with the title "Using STM for modular concurrency" and subtitle "An industrial experience report on Software Transactional Memory" by Duncan Coutts. The video thumbnail shows a man with a beard and long hair, identified as Andres Löh, holding a blue mug. Below the video, the text "Haskell Symposium 2020 Keynote by Duncan Coutts: Using STM for modular concurrency" is displayed, along with a "NOT YET RATED" badge and a "2 years ago | More" link. A "Follow" button for Andres Löh is also present. The right sidebar shows a section titled "More from Andres Löh" with a toggle for "Autoplay next video" and two video thumbnails, both titled "Haskell Symposiu..." by Andres Löh.

Haskell Symposium 2020 Keynote by Duncan Coutts: Using STM for modular concurrency

NOT YET RATED

2 years ago | More

Andres Löh + Follow

More from Andres Löh

Autoplay next video

Haskell Symposiu... Andres Löh

Haskell Symposiu... Andres Löh

# Overload design and backpressure


As a consultant I ask:

Q: what is your design for system overload?


Ummm...

Q: what does your demand vs throughput curve look like?


The good



The bad



The ugly



Wouldn't it be nice if our basic design patterns gave us good results?



## Project context



### Commercial context

- ▶ a blockchain and a crypto-currency
- ▶ a top 10 crypto-currency (by market capitalisation)

### Technical context

- ▶ a **from-scratch** blockchain implementation in Haskell
- ▶ interacting networked nodes, **lots of concurrency**
- ▶ design assumption that **'they really are out to get you'**

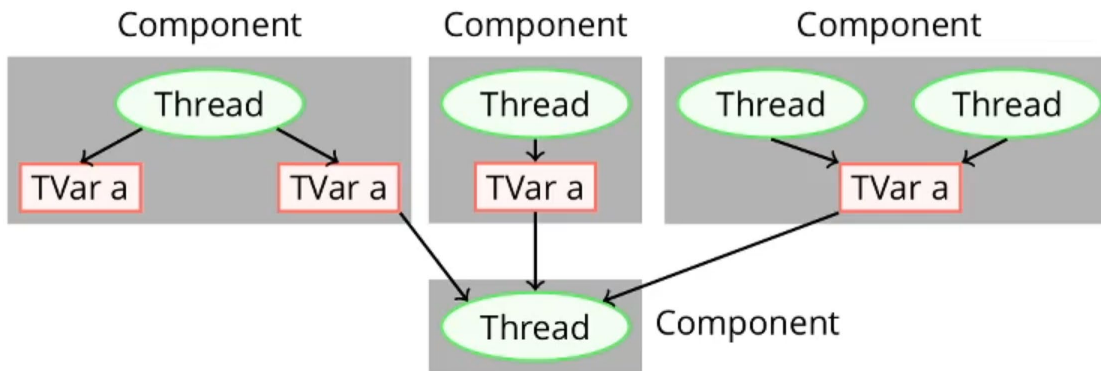


# Background ideas

Ideas from previous projects working with networking experts

- ▶ Queues often make things worse in overload situations and are a source of timing **variability**
- ▶ **Pull**-based designs are often better than **push**-based
- ▶ Aim for designs that do not become **less efficient** under load

# Design thought process

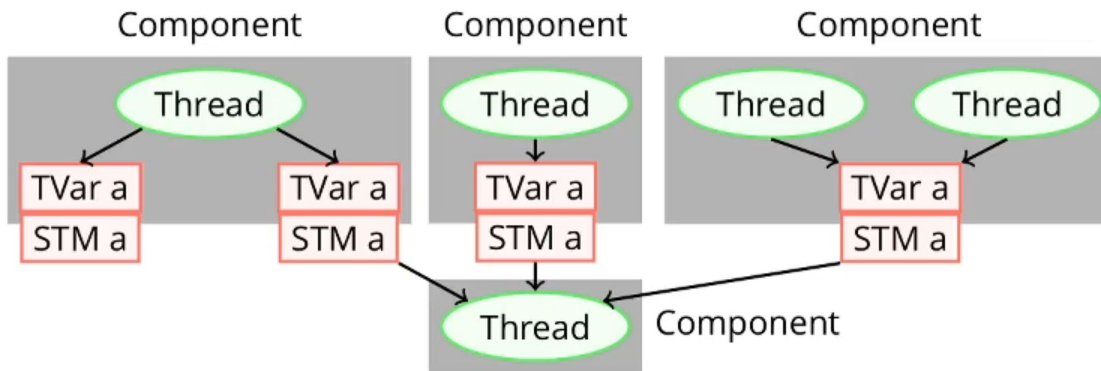


- ▶ Unidirectional data flow for each TVar
- ▶ Associate TVars with the components that write to them





# Design thought process



- ▶ Unidirectional data flow for each TVar
- ▶ Associate TVars with the components that write to them
- ▶ Expose TVar reads as opaque STM queries  
Think of such STM queries as time-varying observations



## Acting on the current state

We observe the **current** state, not all intermediate changes.

This encourages a pattern where we act based on the current state.

- ▶ Irrespective of how many changes there have been
- ▶ Can miss intermediate states if there are frequent changes
- ▶ Can become **more efficient** as we get more overloaded





## Use of STM within Cardano

The use of STM within Cardano has been a **clear success**

- ▶ Allowed a modular design by appropriate use of concurrency
- ▶ Used with explicit (pull-based) protocols for distributed concurrency
- ▶ Handles overload well: slows down asking for more work
- ▶ Concurrency testing found **lots** of bugs, very few found in production
- ▶ Did not hit any STM weak spots
  - no long-running STM transactions
  - no fairness problems
  - no low level performance problems



## Contrast with message-passing

### Message passing

- ▶ push-based
- ▶ act on individual change events
- ▶ implicit queues
- ▶ resource control is implicit (size of queues)
- ▶ no natural backpressure

### State observation

- ▶ pull-based
- ▶ act on changed state eventually
- ▶ no queues
- ▶ resource control is explicit (content of state variables)
- ▶ natural backpressure by slowdown



Haskell + STM =  ?

- STM is available for
  - Scala, Rust, .NET, the JVM, Go, C, C++...
- Haskell uses types to *guarantee* transactions only use TVars.
- Most Haskell memory accesses are *not* to Tvars, and need not be `transactionalized`.

# Summary

- STM is an approach to *concurrent programming* that reduces pain
- STM is *compositional*; transactions can be combined in arbitrary ways
- STM is *expressive*; many other concurrency primitives can be expressed in terms of STM
- STM *performs well*; it can help reduce contention
- STM enables 'state observation design', with built-in ability to withstand overload.