

Parallel Functional Programming

Lecture 3

Mary Sheeran

(with thanks to Simon Marlow for use of slides)

<https://chalmers.instructure.com/courses/23443>

Remember nfib

```
nfib :: Integer -> Integer
nfib n | n < 2 = 1
nfib n = nfib (n-1) + nfib (n-2) + 1
```

- A trivial function that returns the number of calls made—and makes a very large number!

n	nfib n
10	177
20	21891
25	242785
30	2692537

Parallelism

par x y

- “Spark” x in parallel with computing y
 - (and return y)
- The run-time system *may* convert a spark into a parallel task—or it may not
- Starting a task is cheap, but not free

Parallelism

$x \text{ `par` } y$

Sequencing

pseq x y

- Evaluate x *before* y (and return y)
- Used to *ensure* we get the right evaluation order

Sequencing

$x \text{ `pseq` } y$

- Binds more tightly than par

Using par and pseq

```
import Control.Parallel

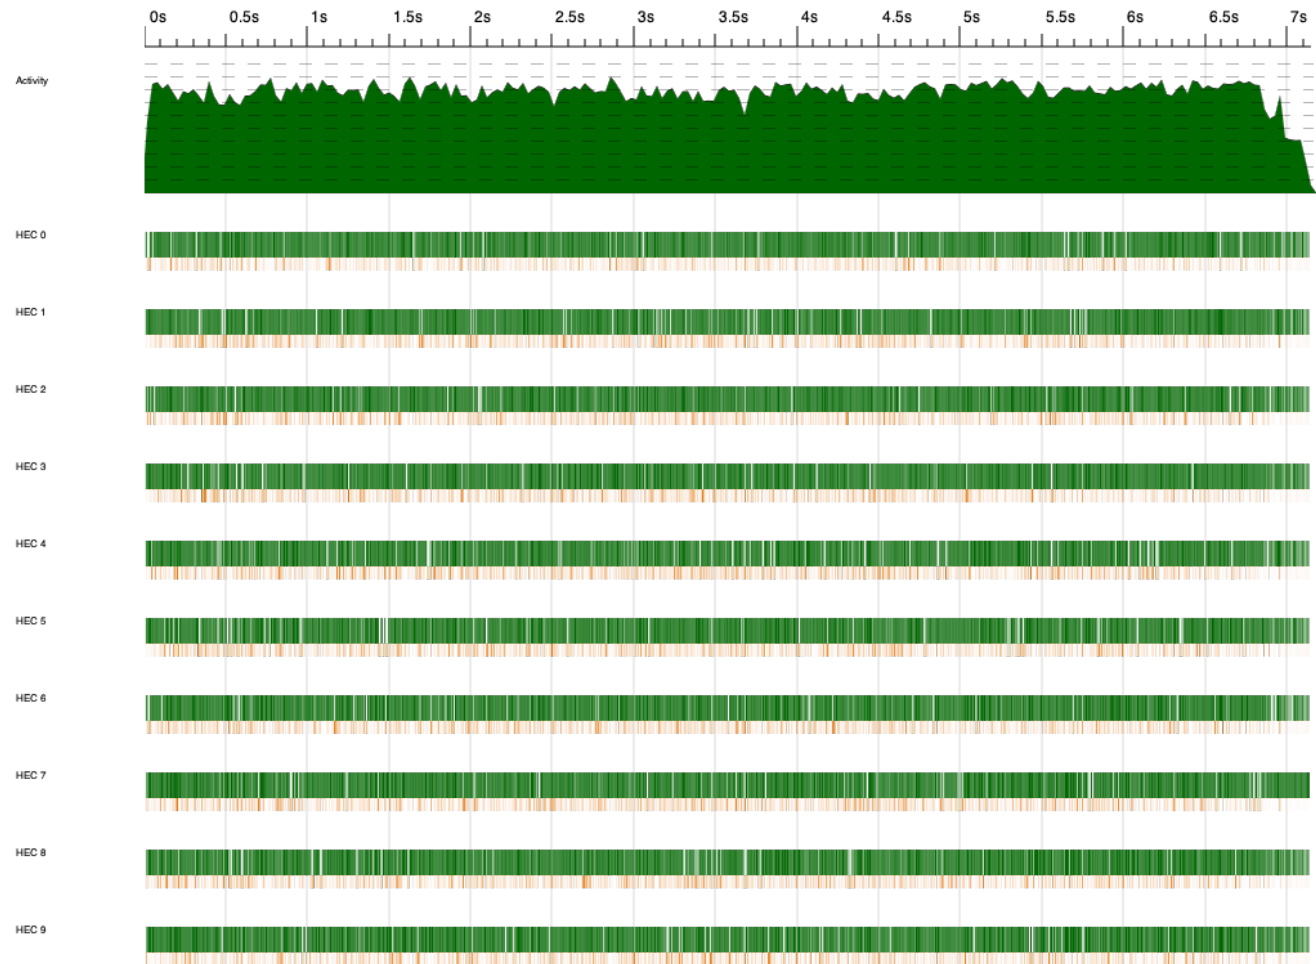
rfib :: Integer -> Integer
rfib n | n < 2 = 1
rfib n = nf1 `par` (nf2 `pseq` nf2 + nf1 + 1)
    where nf1 = rfib (n-1)
          nf2 = rfib (n-2)
```

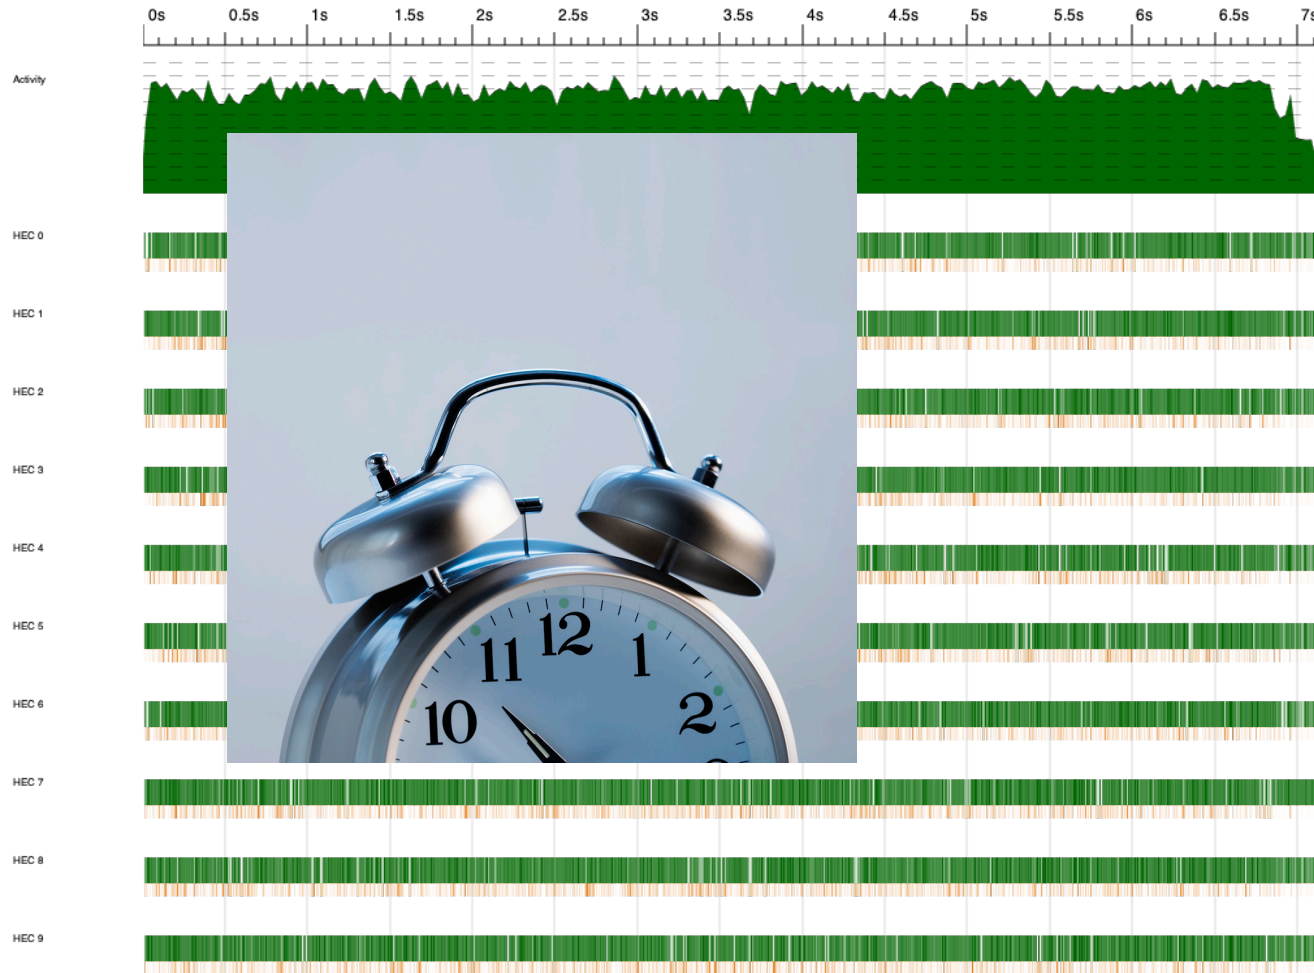
Using par and pseq

```
import Control.Parallel

rfib :: Integer -> Integer
rfib n | n < 2 = 1
rfib n = nf1 `par` (nf2 `pseq` nf2 + nf1 + 1)
    where nf1 = rfib (n-1)
          nf2 = rfib (n-2)
```

- Evaluate *nf1 in parallel with* (Evaluate *nf2 before ...*)





What's happening?

```
$stack run -- +RTS -N10 -s
```



-s to get stats

Hah

331160281

...

SPARKS: 165597162 (7874 converted, 94153001 overflowed,
0 dud, 7257490 GC'd, 37481069 fizzled)

INIT	time	0.000s	(0.007s elapsed)
MUT	time	17.660s	(2.256s elapsed)
GC	time	0.099s	(0.097s elapsed)
EXIT	time	0.000s	(0.010s elapsed)
Total	time	17.760s	(2.371s elapsed)

Hah

331160281

...

SPARKS: 165597162 (7874 converted, 94153001 overflowed,
0 dud, 7257490 GC'd, 37481069 fizzled)

INIT	time	0.000s
MUT	time	17.660s
GC	time	0.099s
EXIT	time	0.000s
Total	time	17.760s (2.371s elapsed)

converted = turned into
useful parallelism

Controlling Granularity

- Let's use a threshold for going sequential, **t**

```
tfib :: Integer -> Integer -> Integer
tfib t n | n < t = sfib n
tfib t n = nf1 `par` nf2 `pseq` nf1 + nf2 + 1
  where nf1 = tfib t (n-1)
        nf2 = tfib t (n-2)
```

Better

tfib 32 40 gives

SPARKS: 88 (63 converted, 0 overflowed, 0 dud,
0 GC'd, 25 fizzled)

INIT	time	0.000s	(0.006s	elapsed)
MUT	time	5.044s	(0.648s	elapsed)
GC	time	0.024s	(0.029s	elapsed)
EXIT	time	0.000s	(0.002s	elapsed)
Total	time	5.068s	(0.685s	elapsed)


Better

tfib 32 40 gives


SPARKS: 88 (63 converted, 0 overflowed, 0 dud,
0 GC'd, 25 fizzled)

INIT	time	0.000s	(0.000s elapsed)
MUT	time	5.068s	(5.068s elapsed)
GC	time	0.000s	(0.000s elapsed)
EXIT	time	0.000s	(0.000s elapsed)
Total	time	5.068s	(0.685s elapsed)

Was 1.7s on my previous
(4 core) Mac 😊



8 cores slightly better (efficiency cores perhaps not so useful)



30413 events, 0.726s

(next step is to move to Criterion, then no eventlog, get slightly better speed)

What are we controlling?

The division of the work into possible parallel tasks (**par**) including choosing size of tasks

GHC runtime takes care of choosing which sparks to actually evaluate in parallel and of distribution (load balancing)

Need also to control order of evaluation (**pseq**) and degree of evaluation

Dynamic behaviour is the term used for how a pure function gets partitioned, distributed and run

Remember, this is deterministic parallelism. The answer is always the same!

positive so far (par and pseq)

Don't need to

express communication

express synchronisation

deal with threads explicitly

BUT

par and pseq are difficult to use 😞

BUT

par and pseq are difficult to use ☹️

MUST


Pass an unevaluated computation to par

It must be somewhat expensive

Make sure the result is not needed for a bit


Make sure the result is shared by the rest of the program

Even if you get it right

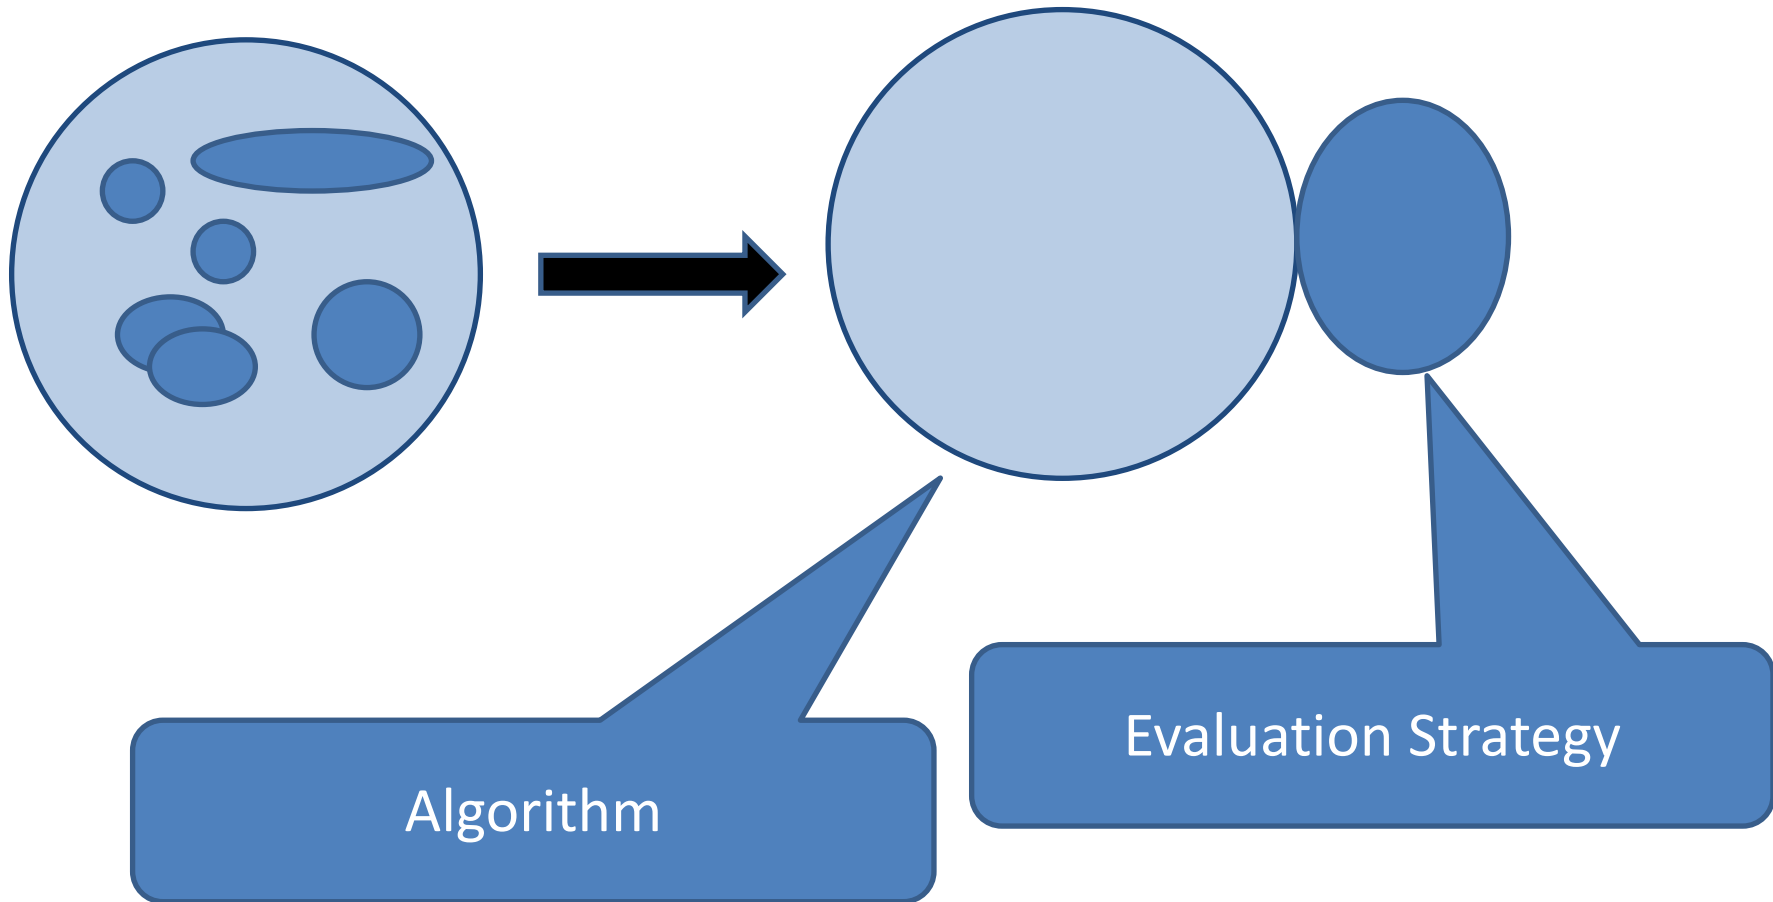


Original code + par + pseq + rnf etc.
can be opaque

Separate concerns



Separate concerns



Evaluation Strategies

express dynamic behaviour independent of the algorithm

provide abstractions above par and pseq

are modular and compositional
(they are ordinary higher order functions)

can capture patterns of parallelism

Papers (last lecture)

Haskell on a Shared-Memory Multiprocessor

Tim Harris Simon Marlow Simon Peyton Jones
Microsoft Research, Cambridge
{tharris,simonmar,simonpj}@microsoft.com

Haskell'05

Runtime Support for Multicore Haskell

Simon Marlow
Microsoft Research
Cambridge, U.K.
simonmar@microsoft.com

Simon Peyton Jones
Microsoft Research
Cambridge, U.K.
simonpj@microsoft.com

Satnam Singh
Microsoft Research
Cambridge, U.K.
satnams@microsoft.com

ICFP'09

Papers (last lecture)

Haskell on a Shared-Memory Multiprocessor

Tim Harris Simon Marlow Simon Peyton Jones
Microsoft Research, Cambridge
{tharris,simonmar,simonpj}@microsoft.com

Haskell'05

85

Runtime Support for Multicore Haskell

Simon Marlow
Microsoft Research
Cambridge, U.K.
simonmar@microsoft.com

Simon Peyton Jones
Microsoft Research
Cambridge, U.K.
simonpj@microsoft.com

Satnam Singh
Microsoft Research
Cambridge, U.K.
satnams@microsoft.com

ICFP'09

Papers (last lecture)

Haskell on a Shared-Memory Multiprocessor

Tim Harris Simon Marlow Simon Peyton Jones
Microsoft Research, Cambridge
{tharris,simonmar,simonpj}@microsoft.com

Haskell'05

85

Runtime Support for Multicore Haskell

Simon Marlow
Microsoft Research
Cambridge, U.K.
simonmar@microsoft.com

Simon Peyton Jones
Microsoft Research
Cambridge, U.K.
simonpj@microsoft.com

Satnam Singh
Microsoft Research
Cambridge, U.K.
satnams@microsoft.com

ICFP'09

183

Latter won ICFP 10 year most
influential paper
=> a retrospective blog post

<https://blog.sigplan.org/2019/12/16/runtime-support-for-multicore-haskell-a-retrospective/>

Papers

Algorithm + Strategy = Parallelism

P.W. TRINDER

Department of Computing Science, University of Glasgow, Glasgow, UK

K. HAMMOND

Division of Computing Science, University of St Andrews, St Andrews, UK

H.-W. LOIDL AND S.L. PEYTON JONES [†]

Department of Computing Science, University of Glasgow, Glasgow, UK

JFP 1998

Seq no more: Better Strategies for Parallel Haskell

Simon Marlow

Microsoft Research, Cambridge, UK
simonmar@microsoft.com

Patrick Maier

Heriot-Watt University, Edinburgh, UK
P.Maier@hw.ac.uk

Hans-Wolfgang Loidl

Heriot-Watt University, Edinburgh, UK
H.W.Loidl@hw.ac.uk

Mustafa K. Aswad

Heriot-Watt University, Edinburgh, UK
mka19@hw.ac.uk

Phil Trinder

Heriot-Watt University, Edinburgh, UK
P.W.Trinder@hw.ac.uk

Haskell'10

Papers

Algorithm + Strategy = Parallelism

P.W. TRINDER

Department of Computing Science, University of Glasgow, Glasgow, UK

K. HAMMOND

Division of Computing Science, University of St Andrews, St Andrews, UK

H.-W. LOIDL AND S.L. PEYTON JONES [†]

Department of Computing Science, University of Glasgow, Glasgow, UK

JFP 1998

389

Seq no more: Better Strategies for Parallel Haskell

Simon Marlow

Microsoft Research, Cambridge, UK
simonmar@microsoft.com

Patrick Maier

Heriot-Watt University, Edinburgh, UK
P.Maier@hw.ac.uk

Hans-Wolfgang Loidl

Heriot-Watt University, Edinburgh, UK
H.W.Loidl@hw.ac.uk

Mustafa K. Aswad

Heriot-Watt University, Edinburgh, UK
mka19@hw.ac.uk

Phil Trinder

Heriot-Watt University, Edinburgh, UK
P.W.Trinder@hw.ac.uk

Haskell'10

Papers

Algorithm + Strategy = Parallelism

P.W. TRINDER

Department of Computing Science, University of Glasgow, Glasgow, UK

K. HAMMOND

Division of Computing Science, University of St Andrews, St Andrews, UK

H.-W. LOIDL AND S.L. PEYTON JONES [†]

Department of Computing Science, University of Glasgow, Glasgow, UK

JFP 1998

389

Seq no more: Better Strategies for Parallel Haskell

Simon Marlow

Microsoft Research, Cambridge, UK
simonmar@microsoft.com

Patrick Maier

Heriot-Watt University, Edinburgh, UK
P.Maier@hw.ac.uk

Hans-Wolfgang Loidl

Heriot-Watt University, Edinburgh, UK
H.W.Loidl@hw.ac.uk

Mustafa K. Aswad

Heriot-Watt University, Edinburgh, UK
mka19@hw.ac.uk

Phil Trinder

Heriot-Watt University, Edinburgh, UK
P.W.Trinder@hw.ac.uk

Haskell'10

109

Papers

Algorithm + Strategy = Parallelism

P.W. TRINDER

Department of Computing Science, University of Glasgow

K. HAMMOND

Division of Computing Science, University of St Andrews

H.-W. LOIDL AND S.L. PEYTON JONES

Department of Computing Science, University of Glasgow

Redesigns strategies

richer set of parallelism combinators

Better specs (evaluation order)

Allows new forms of coordination

generic regular strategies over data structures

speculative parallelism

based on a simple monad

Presentation is about New Strategies

Seq no more: Better Parallelism

Simon Marlow

Microsoft Research, Cambridge, UK
simonmar@microsoft.com

Patrick Maier

Heriot-Watt University, Edinburgh, UK
P.Maier@hw.ac.uk

Heriot-Watt University, Edinburgh, UK
H.W.Loid@hw.ac.uk

HaskeR 10

Mustafa K. Kowad

Heriot-Watt University, Edinburgh, UK
mka19@hw.ac.uk

Phil Trinder

Heriot-Watt University, Edinburgh, UK
P.W.Trinder@hw.ac.uk

The Eval monad

```
import Control.Parallel.Strategies

data Eval a
instance Monad Eval

runEval :: Eval a -> a

rpar :: a -> Eval a
rseq :: a -> Eval a
```

- Eval is pure
- Just for expressing sequencing between rpar/rseq – nothing more
- Compositional – larger Eval sequences can be built by composing smaller ones using monad combinators
- Internal workings of Eval are very simple (see Haskell Symposium 2010 paper)

Expressing evaluation order

```
qfib :: Integer -> Integer
qfib n | n < 2 = 1
qfib n = runEval $ do
    nf1 <- rpar (qfib (n-1))
    nf2 <- rseq (qfib (n-2))
    return (nf1 + nf2 + 1)
```

Expressing evaluation order

```
qfib :: Integer -> Integer
qfib n | n < 2 = 1
qfib n = runEval $ do
    nf1 <- rpar (qfib (n-1))
    nf2 <- rseq (qfib (n-2))
    return (nf1 + nf2 + 1)
```

do this
spark qfib (n-1)

"My argument could be evaluated in parallel"

Expressing evaluation order

```
qfib :: Integer -> Integer
qfib n | n < 2 = 1
qfib n = runEval $ do
    nf1 <- rpar (qfib (n-1))
    nf2 <- rseq (qfib (n-2))
    return (nf1 + nf2 + 1)
```

do this
spark qfib (n-1)

"My argument could be evaluated in parallel"

Remember that the argument should be a thunk!

Expressing evaluation order

```
qfib :: Integer -> Integer
qfib n | n < 2 = 1
qfib n = runEval $ do
    nf1 <- rpar (qfib (n-1))
    nf2 <- rseq (qfib (n-2))
    return (nf1 + nf2 + 1)
```

and **then** this
Evaluate qfib(n-2)
and wait for
result

"Evaluate my argument and wait for the result."

Expressing evaluation order

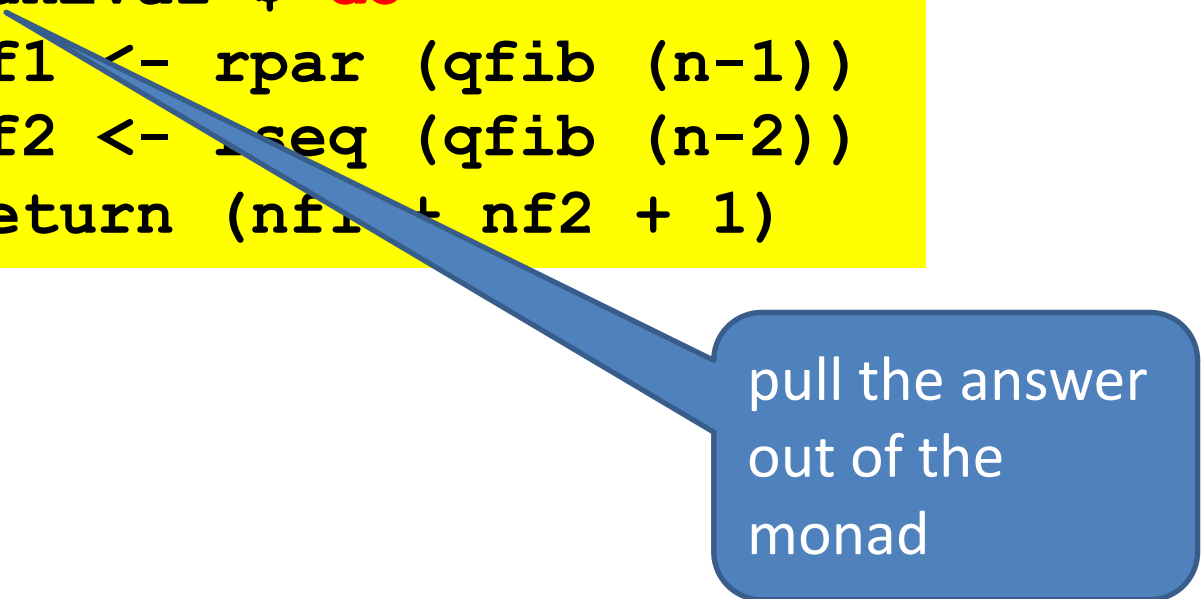
```
qfib :: Integer -> Integer
qfib n | n < 2 = 1
qfib n = runEval $ do
    nf1 <- rpar (qfib (n-1))
    nf2 <- rseq (qfib (n-2))
    return (nf1 + nf2 + 1)
```



the result

Expressing evaluation order

```
qfib :: Integer -> Integer
qfib n | n < 2 = 1
qfib n = runEval $ do
    nf1 <- rpar (qfib (n-1))
    nf2 <- rseq (qfib (n-2))
    return (nf1 + nf2 + 1)
```




pull the answer
out of the
monad


(should add granularity control too!)

What does *rpar* actually do?

```
x <- rpar e
```

- *rpar* creates a *spark* by writing an entry in the *spark pool*
 - *rpar* is very cheap! (not a thread)
- the spark pool is a circular buffer
- when a processor has nothing to do, it tries to remove an entry from its own spark pool, or steal an entry from another spark pool (*work stealing*)
- when a spark is found, it is evaluated
- The spark pool can be full – watch out for spark overflow!





Read Chapters 2 and 3

What do we have?

The Eval monad raises the level of abstraction for pseq and par; it makes fragments of evaluation order first class, and lets us compose them together. We should think of the Eval monad as an Embedded Domain-Specific Language (EDSL) for expressing evaluation order, embedding a little evaluation-order constrained language inside Haskell, which does not have a strongly-defined evaluation order.

(from Haskell 10 paper)

a possible parallel map

```
pMap :: (a -> b) -> [a] -> Eval [b]
pMap f [] = return []
pMap f (a:as) = do
    b  <- rpar (f a)
    bs <- pMap f as
    return (b:bs)
```

a possible parallel map

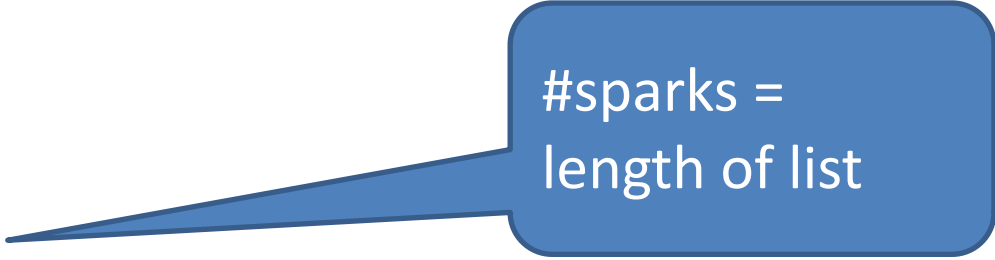
```
import Control.Parallel.Strategies

foo :: Integer -> Integer
foo a = sum [1 .. a]

main
  = print $ sum $ runEval $
    pMap foo (reverse [1..10000])
```

SPARKS: 10000 (8195 converted, 1805 overflowed, 0 dud, 0
GC'd, 0 fizzled)

INIT	time	0.003s (0.009s elapsed)
MUT	time	1.346s (0.410s elapsed)
GC	time	0.010s (0.003s elapsed)
EXIT	time	0.001s (0.000s elapsed)
Total	time	1.361s (0.423s elapsed)




#sparks =
length of list

SPARKS: 10000 (8195 converted, 1805 overflowed, 0 dud, 0
GC'd, 0 fizzled)

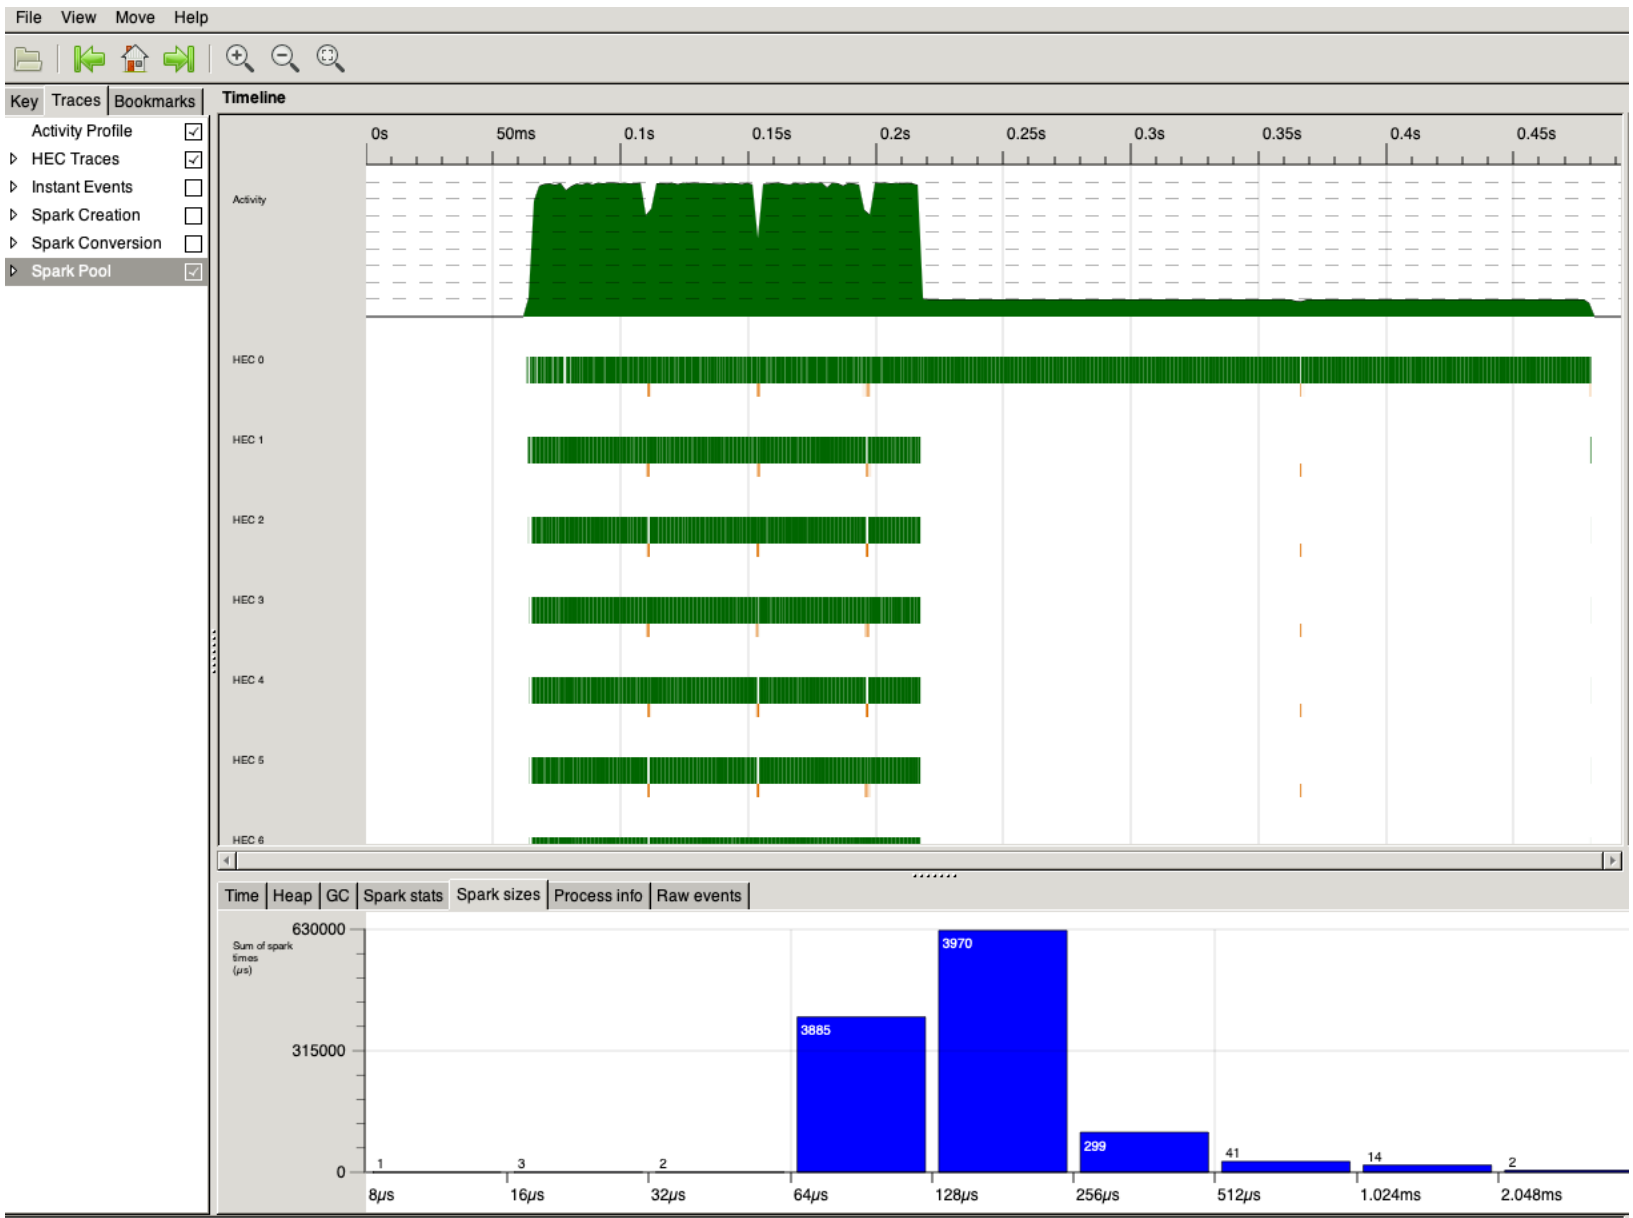
INIT	time	0.003s (0.009s elapsed)
MUT	time	1.346s (0.410s elapsed)
GC	time	0.010s (0.003s elapsed)
EXIT	time	0.001s (0.000s elapsed)
Total	time	1.361s (0.423s elapsed)

Timeline



Time	Heap	GC	Spark stats	Spark sizes	Process info	Raw events
HEC	Total	Converted	Overflowed	Dud	GC'd	Fizzled
Total	10000	8204	1796	0	0	0
HEC 0	0	2721	0	0	0	0
HEC 1	10000	267	1796	0	0	0
HEC 2	0	2613	0	0	0	0
HEC 3	0	2603	0	0	0	0

Figure 1/2/2-exe eventlog (21005 events, 0.445s)



converted

real parallelism at runtime

overflowed

no room in spark pool

dud

first arg of rpar already eval'd

GC'd

sparked expression unused
(removed from spark pool)

fizzled

uneval'd when sparked, later
eval'd independently => removed

our parallel map

```
pMap :: (a -> b) -> [a] -> Eval [b]
pMap f [] = return []
pMap f (a:as) = do
    b  <- rpar (f a)
    bs <- pMap f as
    return (b:bs)
```

parallel map

- + Captures a pattern of parallelism
- + good to do this for standard higher order function like map
- + can easily do this for other standard sequential patterns

return (b.bs)

BUT

- had to write a new version of map
- mixes algorithm and dynamic behaviour



return (b.bs)

Evaluation Strategies



Raise level of abstraction

Encapsulate parallel programming idioms as reusable components that can be composed

Strategy (as of 2010)

```
type Strategy a = a -> Eval a
```

function

evaluates its input to some degree

traverses its argument and uses rpar and rseq to express dynamic behaviour / sparking

returns an equivalent value in the Eval monad

using

```
using :: a -> Strategy a -> a  
x `using` strat = runEval (strat x)
```

Program typically applies the strategy to a structure and then uses the returned value, discarding the original one (which is why the value had better be equivalent)

An almost identity function that does some evaluation and expresses how that can be parallelised



“The key idea in our reformulation is that a strategy returns a new version of its argument, in which the sparked computations have been embedded.”

Basic strategies

```
r0 :: Strategy a  
r0 x = return x
```

```
rpar :: Strategy a  
rpar x = x `par` return x
```

```
rseq :: Strategy a  
rseq x = x `pseq` return x
```

```
rdeepseq :: NFData a => Strategy a  
rdeepseq x = rnf x `pseq` return x
```

Basic strategies

```
r0 :: Strategy a  
r0 x = return x
```

NO evaluation

```
rpar :: Strategy a  
rpar x = x `par` return x
```

```
rseq :: Strategy a  
rseq x = x `pseq` return x
```

```
rdeepseq :: NFData a => Strategy a  
rdeepseq x = rnf x `pseq` return x
```

Basic strategies

```
r0 :: Strategy a  
r0 x = return x
```

```
rpar :: Strategy a  
rpar x = x `par` return x
```



spark x

```
rseq :: Strategy a  
rseq x = x `pseq` return x
```

```
rdeepseq :: NFData a => Strategy a  
rdeepseq x = rnf x `pseq` return x
```

Basic strategies

```
r0 :: Strategy a  
r0 x = return x
```

```
rpar :: Strategy a  
rpar x = x `par` return x
```

```
rseq :: Strategy a  
rseq x = x `pseq` return x
```

```
rdeepseq :: NFData a => Strategy a  
rdeepseq x = rnf x `pseq` return x
```



evaluate x
to WHNF

Basic strategies

```
r0 :: Strategy a  
r0 x = return x
```

```
rpar :: Strategy a  
rpar x = x `par` return x
```

```
rseq :: Strategy a  
rseq x = x `pseq` return x
```

```
rdeepseq :: NFData a => Strategy a  
rdeepseq x = rnf x `pseq` return x
```



fully evaluate x

evalList

```
evalList :: Strategy a -> Strategy [a]
evalList s []      = return []
evalList s (x:xs) = do x' <- s x
                      xs' <- evalList s xs
                      return (x':xs')
```


evalList

```
evalList :: Strategy a -> Strategy [a]
evalList s [] = return []
evalList s (x:xs) = do x' <- s x
                      xs' <- evalList s xs
```

Takes a Strategy on a and returns a Strategy on lists of a
Building strategies from smaller ones

parList

```
evalList :: Strategy a -> Strategy [a]
evalList s [] = return []
evalList s (x:xs) = do x' <- s x
                      xs' <- evalList s xs
                      return (x':xs')
```

```
parList :: Strategy a -> Strategy [a]
parList s = evalList (rpar `dot` s)
```

Composing strategies

```
dot :: Strategy a -> Strategy a -> Strategy a
strat2 `dot` strat1 = strat2 . runEval . strat1
```

In reality

```
evalList :: Strategy a -> Strategy [a]  
evalList = evalTraversable
```

```
parList :: Strategy a -> Strategy [a]  
parList = parTraversable
```

In reality

```
evalList :: Strategy a -> Strategy [a]  
evalList = evalTraversable
```

```
parList :: Strategy a -> Strategy [a]  
parList = parTraversable
```

The equivalent of evalList and of parList are available for many data structures (Traversable). So defining parX for many X is really easy

=> generic strategies for data-oriented parallelism



How do we *use* a Strategy?

```
type Strategy a = a -> Eval a
```

- We could just use runEval
- But this is better:

```
x `using` s = runEval (s x)
```

- e.g.

```
myList `using` parList rdeepseq
```

- Why better? Because we have a “law”:
 - $x \text{ `using` } s \approx x$
 - We can insert or delete “`using` s” without changing the semantics of the program

Is that really true?

- Well, not entirely.
 1. It relies on Strategies returning “the same value” (*identity-safety*)
 - Strategies from the library obey this property
 - Be careful when writing your own Strategies
 2. `x`using`s` might do more evaluation than just `x`.
 - So the program with `x`using`s` might be `_|_`, but the program with just `x` might have a value
- if identity-safety holds, adding `using` cannot make the program produce a different result (other than `_|_`)

```
parListChunk :: Int -> Strategy a -> Strategy [a]
parListChunk n strat xs
  | n <= 1      = parList strat xs
  | otherwise
    = concat `fmap` parList (evalList strat) (chunk n xs)
```




```
parListChunk :: Int -> Strategy a -> Strategy [a]
parListChunk n strat xs
  | n <= 1      = parList strat xs
  | otherwise
    = concat `fmap` parList (evalList strat) (chunk n xs)
```

```
chunk :: Int -> [a] -> [[a]]
chunk _ [] = []
chunk n xs = as : chunk n bs
  where
    (as,bs) = splitAt n xs
```

```
parListChunk :: Int -> Strategy a -> Strategy [a]
```

```
parListChunk n strat
```



```
parListChunk :: Int -> Strategy a -> Strategy [a]
```

Before

```
print $ sum $ runEval $ pMap foo (reverse [1..10000])
```

Now

```
print $ sum $  
(map foo (reverse [1..10000])) `using` parListChunk 50 rdeepseq )
```

Question: how many sparks? Why?

```
parListChunk :: Int -> Strategy a -> Strategy [a]
```

Before

```
print $ sum $ runEval $ pMap foo (reverse [1..10000])
```

Now

```
print $ sum $  
(map foo (reverse [1..10000])) `using` parListChunk 50 rdeepseq )
```

SPARKS: 200 (200 converted, 0 overflowed, 0 dud, 0 GC'd, 0 fizzled)

```
parListChunk :: Int -> Strategy a -> Strategy [a]
```

Before

```
print $ sum $
```

Now

```
print $ sum $  
(map foo (reverse [1..1000000])  
  using `parListChunk 50 rdeepseq` )
```

Remember not to be a control freak, though.
Generating plenty of sparks gives the
runtime the freedom it needs to make good
choices (=> Dynamic partitioning for free)

SPARKS: 200 (200 converted, 0 overflowed, 0 dud, 0 GC'd, 0 fizzled)

Research questions (surprisingly unexplored)

What is a good set of combinators for expressing and controlling parallelism, including granularity?

Can recent(ish) developments in the Haskell type system (such as type level natural numbers) help??

Divide and conquer

Capturing patterns of parallel computation is a major strong point of strategies

D&C is a typical example (see also parBuffer, parallel pipelines etc)

These are covered in Lab A

Skeletons


- encode fixed set of common coordination patterns and provide efficient parallel implementations (Cole, 1989)
- Popular in both functional and non-functional languages. See particularly Eden (Loogen et al, 2005)

A difference: one can / should roll ones own strategies


Strategies: summary

- + elegant redesign by Marlow et al (Haskell 10)
- + better separation of concerns
- + Laziness is essential for modularity
- + generic strategies for (Traversable) data structures
- + Marlow's book contains a nice **kmeans** example. Read it!
- Having to think so much about evaluation order is worrying!
Laziness is not only good here. **(Cue the Par Monad Lecture!)**


Strategies: summary




Better visualisation





Better visualisation



Better visualisation





Landscape for parallel Haskell (Marlow)

Parallel

par / pseq (1) (monads (2))

Strategies (3)

Par Monad. (4)

DPH, Repa , Accelerate (Keller guest lec.)

Landscape for parallel Haskell (Marlow)

Parallel

par / pseq (1) (monads (2))
Strategies (3)
Par Monad. (4)
Repa, DPH, Accelerate (Keller guest lec.)

Concurrent

ForkIO

MVAR

STM



async

Cloud Haskell

parallel FP

Parallel

par / pseq (1) (monads (2))

Strategies (3)

Par Monad. (4)

Repa, DPH, Accelerate (Keller guest lec.)

Parallelising QuickCheck in Haskell

Data Parallel Programming (NESL) (5)

Futhark Java

Parallel/ Concurrent

robust parallel programming in Erlang

Map Reduce

noSQL databases etc.

(STM)

