

Parallel Programming in Erlang

John Hughes

What is Erlang?

Erlang 

Haskell

- Types
- Lazyness
- Purity
- + Concurrency
- + Syntax

If you know Haskell, Erlang is easy to learn!

QuickSort again

- Haskell

```
qsort [] = []
qsort (x:xs) = qsort [y | y <- xs, y < x]
                ++ [x]
                ++ qsort [y | y <- xs, y >= x]
```

- Erlang

```
qsort([]) -> [];
qsort([X|Xs]) -> qsort([Y || Y <- Xs, Y < X])
                ++ [X]
                ++ qsort([Y || Y <- Xs, Y >= X]) .
```

qsort [] =

- Haskell

qsort [] = []

**qsort (x:xs) = qsort [y | y <- xs, y < x]
++ [x]**

qsort([]) ->

- Erlang

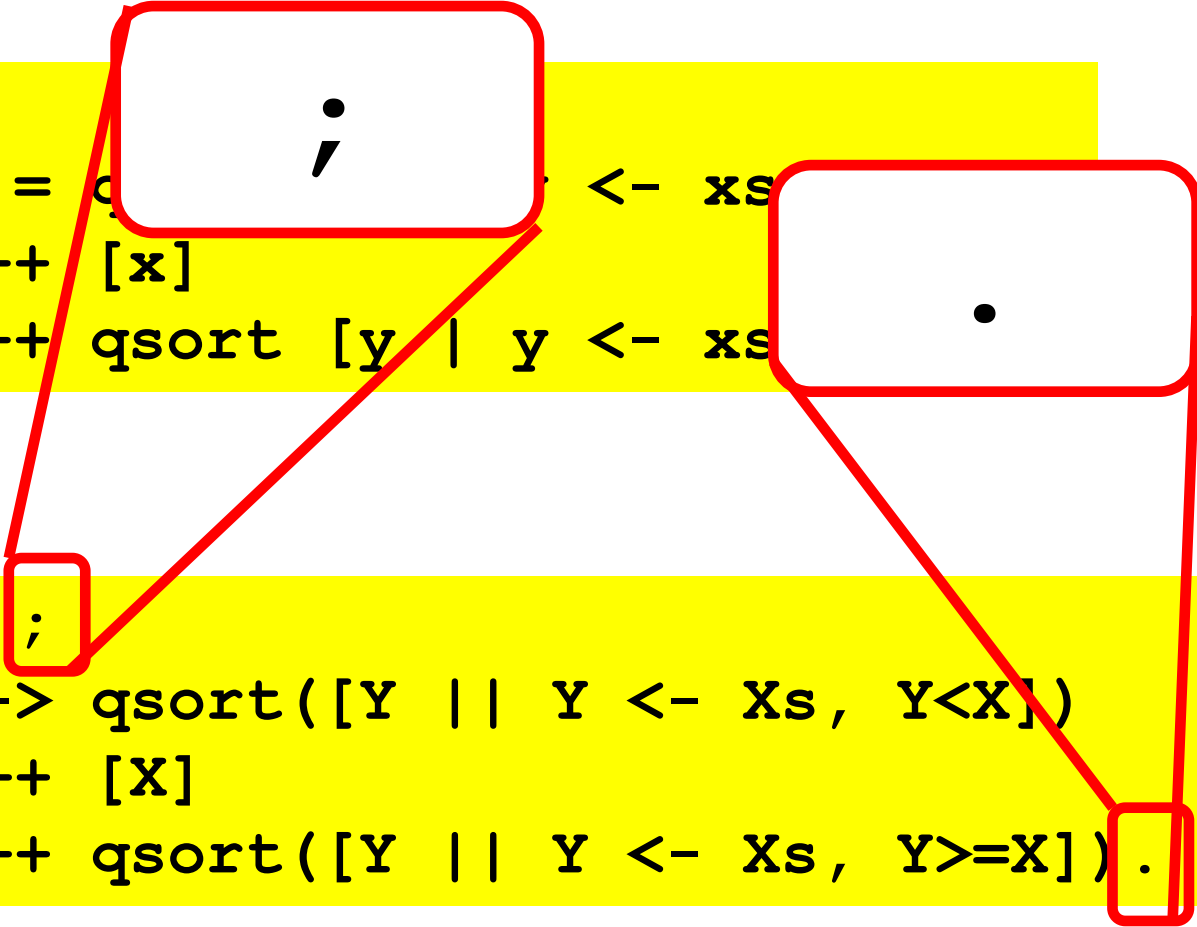
qsort([]) -> [];

**qsort([X|Xs]) -> qsort([Y || Y <- Xs, Y < X])
++ [X]
++ qsort([Y || Y <- Xs, Y >= X]).**

QuickSort again

- Haskell

```
qsort [] = []  
qsort (x:xs) = qsort [y | y <- xs, y < x]  
              ++ [x]  
              ++ qsort [y | y <- xs, y >= x]
```



- Erlang

```
qsort([]) -> [];  
qsort([X|Xs]) -> qsort([Y || Y <- Xs, Y < X])  
                ++ [X]  
                ++ qsort([Y || Y <- Xs, Y >= X]).
```

Qu **$x : xs$** gain

- Haskell

```
qsort [] = []  
qsort (x:xs) = qsort [y | y <- xs, y < x]  
              ++ [x]  
              ++ qsort [y | y <- xs, y >= x]
```

$[X | Xs]$

- Erlang

```
qsort([]) -> [];  
qsort([X|Xs]) -> qsort([Y || Y <- Xs, Y < X])  
                  ++ [X]  
                  ++ qsort([Y || Y <- Xs, Y >= X]).
```

QuickSort again

- Haskell

```
qsort [] = []  
qsort (x:xs) = qsort [y | y <- xs, y < x]  
               ++ [x]  
               ++ qsort [y | y <- xs, y >= x]
```

|

||

- Erlang

```
qsort([]) -> [];  
qsort([X|Xs]) -> qsort([Y | Y <- Xs, Y < X])  
                  ++ [X]  
                  ++ qsort([Y | Y <- Xs, Y >= X]).
```

||

Declare the
module name

foo.erl

Simplest just to
export everything

```
-module(foo) .  
-compile(export_all) .  
  
qsort([]) ->  
    [];  
qsort([X|Xs]) ->  
    qsort([Y || Y <- Xs, Y<X]) ++  
    [X] ++  
    qsort([Y || Y <- Xs, Y>=X]) .
```

```
-compile([export_all,nowarn_export_all]) .
```


werl/erl REPL

```
Erlang
File Edit Options View Help
Erlang/OTP 25 [erts-
hreads:1] [jit:ns]

Eshell V13.1.3 (sh with ^G)
1> c(foo).
{ok,foo}
2> foo:qsort([1,9,2,5,4,3,6,8,7]).
[1,2,3,4,5,6,7,8,9]
3> █
```

Compile foo.erl
"foo" is an *atom*—a constant

Don't forget the "."!

foo:qsort calls qsort from the foo module

- Much like ghci



Test Data

- Create some test data; in foo.erl:

```
random_list(N) ->  
    [rand:uniform(1000000) || _ <- lists:seq(1,N)].
```

Side-
effects!

Instead of
[1..N]

- In the

```
L = foo:random_list(200000) .
```

Timing calls

Module

Function

Arguments

```
4> timer:tc(foo, qsort, [L]).  
{84172,  
 [1, 5, 7, 10, 23, 24, 25, 29,  
 58, 63, 61, 62, 62, 62, 67, 68,  
 71, 81, 94, 96, 102, 105, 106 | ...]}
```

atoms—i.e.
constants

Microseconds

{A,B,C} is a tuple

Benchmarking

Binding a
name... c.f. **let**

Macro: current
module name

```
benchmark (Fun, L) ->  
  NRuns = 500,  
  {T, _} = timer:tc(fun() ->  
    [?MODULE:Fun(L)  
    || _ <- lists:seq(1, NRuns)],  
    ok  
  end),  
  T / (1000*NRuns) .
```

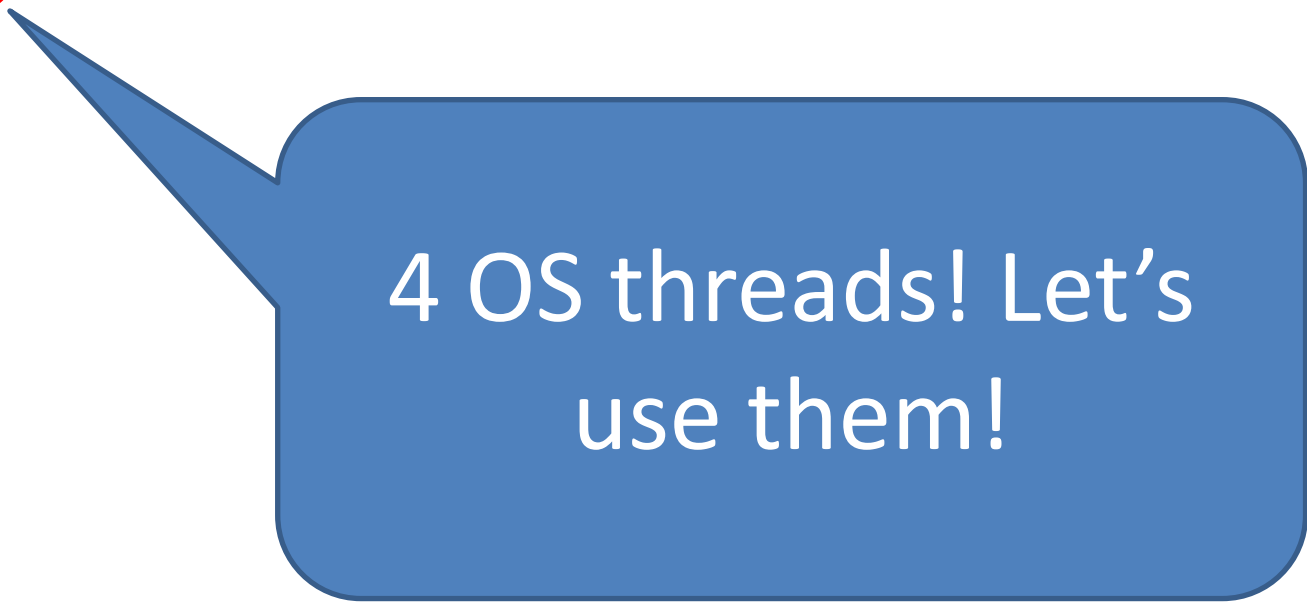
```
4> foo:benchmark(qsort, L) .
```

```
74.653528
```

Parallelism

```
68> erlang:system_info(schedulers_online) .
```

4



4 OS threads! Let's
use them!

Parallelism in Erlang

- Processes are created *explicitly*

```
Pid = spawn_link(fun() -> ...Body... end)
```

- Start a process which executes ...Body...
- `fun() -> Body end` *~* `\() -> Body`
- `Pid` is the *process identifier*

<0.86.0>

Parallel Sorting

```
psort([]) ->  
  [];  
psort([X|Xs]) ->  
  spawn_link(  
    fun() ->  
      psort([Y || Y <- Xs, Y >= X])  
    end) ,  
  psort([Y || Y <- Xs, Y < X]) ++  
    [X] ++  
    ???.
```

Sort second half in
parallel...

But how do we get the
result?

Message Passing

Pid ! Msg

- Send a message to Pid
- *Asynchronous*—do not wait for delivery



Message Receipt

```
receive
```

```
    Msg -> ...
```

```
end
```

- Wait for a message, then bind it to Msg

Parallel Sorting

```
psort([]) ->  
  [];  
psort([X|Xs]) ->  
  Parent = self(),  
  spawn_link(  
    fun() ->  
      Parent !  
      psort([Y || Y <- Xs, Y >= X])  
    end),  
  psort([Y || Y <- Xs, Y < X]) ++  
    [X] ++  
    receive Ys -> Ys end.
```

The Pid of the
executing process

Send the result back
to the parent

Wait for the result *after* sorting the first half

Benchmarks

```
69> foo:benchmark(qsort,L) .  
74.318884  
70> foo:benchmark(psort,L) .  
154.869018
```

- Parallel sort is slower! *Why?*



Controlling Granularity

```
psort2(Xs) -> psort2(5,Xs) .
```

```
psort2(0,Xs) -> qsort(Xs) ;
```

```
psort2(_,[]) -> [];
```

```
psort2(D,[X|Xs]) ->
```

```
    Parent = self(),
```

```
    spawn_link(fun() ->
```

```
        Parent !
```

```
        psort2(D-1,[Y || Y <- Xs, Y >= X])
```

```
    end),
```

```
    psort2(D-1,[Y || Y <- Xs, Y < X]) ++
```

```
    [X] ++
```

```
    receive Ys -> Ys end.
```

Benchmarks

```
69> foo:benchmark(qsort,L) .  
74.318884
```

```
70> foo:benchmark(psort,L) .  
154.869018
```

```
71> foo:benchmark(psort2,L) .  
35.882206
```

- >2x speedup

Correctness

```
91> foo:psort2(L) == foo:qsort(L) .  
false  
92> foo:psort2("hello world") .  
" edh11loorw"
```

Oops!

What's going on?

```
psort2(D,[X|Xs]) ->  
  Parent = self(),  
  spawn_link(fun() ->  
    Parent ! ...  
    end),  
  psort2(D-1,[Y || Y <- Xs, Y < X]) ++  
  [X] ++  
  receive Ys -> Ys end.
```

What's going on?


```
psort2(D, [X|Xs]) ->  
  Parent = self(),  
  spawn link(fun() ->  
    Parent ! ...  
    end),  
  Parent = self(),  
  spawn link(fun() ->  
    Parent ! ...  
    end),  
  psort2(D-2, [Y || Y <- Xs, Y < X]) ++  
  [X] ++  
  receive Ys <-> Ys end ++  
  [X] ++  
  receive Ys <-> Ys end.
```

The diagram illustrates a process spawning two child processes. The parent process sends messages to both children. The children then receive these messages and return them to the parent. The parent process then returns the results to the caller.

Message Passing Guarantees



Message Passing Guarantees





Tagging Messages Uniquely

```
Ref = make_ref()
```

- Create a globally unique reference

```
Parent ! {Ref,Msg}
```

- Send the message tagged with the reference

```
receive {Ref,Msg} -> ... end
```

- Match the reference on receipt... picks the right message from the mailbox

A correct parallel sort

```
psort3(Xs) ->  
    psort3(5,Xs) .
```

```
psort3(0,Xs) ->  
    qsort(Xs) ;
```

```
psort3(_,[]) ->  
    [] ;
```

```
psort3(D,[X|Xs]) ->  
    Parent = self(),  
    Ref = make_ref(),  
    spawn_link(fun() ->  
        Parent ! {Ref,psort3(D-1,[Y || Y <- Xs, Y >= X])}  
    end),  
    psort3(D-1,[Y || Y <- Xs, Y < X]) ++  
    [X] ++  
    receive {Ref,Greater} -> Greater end.
```


Tests

```
69> foo:benchmark(qsort,L) .  
74.318884  
...  
72> foo:benchmark(psort3,L) .  
36.818348  
73> foo:qsort(L) == foo:psort3(L) .  
true
```

- >2x speedup, and now it works 😊

Parallelism in Erlang vs Haskell


- Haskell processes *share memory*






Parallelism in Erlang vs Haskell

- Erlang processes each have their own heap



Pid ! Msg



In Haskell, *forcing to nf* is linear time

- Messages have to be *copied*
- No global garbage collection—each process collects its own heap

What's copied here?

```
psort3(D, [X|Xs]) ->  
  Parent = self(),  
  Ref = make_ref(),  
  spawn_link(fun() ->  
    Parent !  
    {Ref,  
      psort3(D-1, [Y || Y <- Xs, Y >= X]) }  
  end),
```

- Is it sensible to copy *all of* *Xs* to the new process?

Better

A small improvement—but Erlang lets us *reason* about copying

```
psort4(D, [X|Xs]) ->
    Parent = self(),
    Ref = make_ref(),
    Grtr = [Y || Y <- Xs, Y >= X],
    spawn_link(fun() ->
        Parent ! {Ref, psort4(D-1, Grtr)}
    end),
```

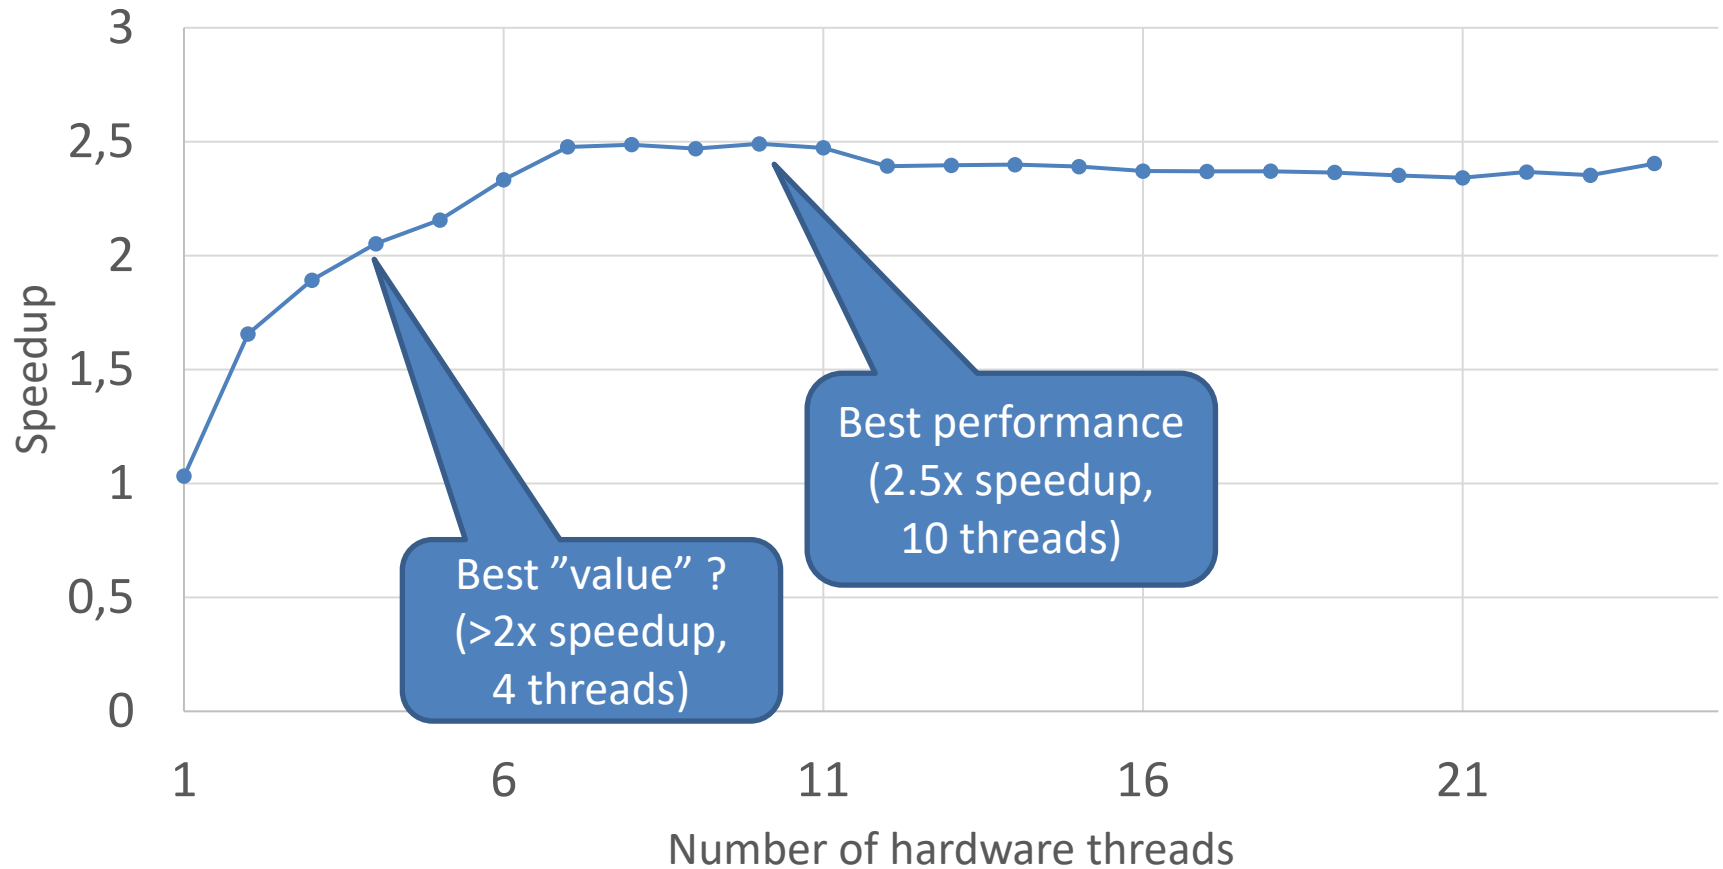
```
72> foo:benchmark(psort3,L) .
36.818348
```

...


```
74> foo:benchmark(psort4,L) .
30.543886
```



Speedups (parallel to depth 8 now)




Speedups



Core i9, 16 cores, L1 cache 1.4MB *in total*, L2 cache 14MB, L3 cache 30MB
(Data size for this problem: 200K x 16B = 3.2MB)

Reducing copying

$$[1,2] ++ [3,4] == [1,2,3,4]$$



Recall qsort

```
qsort([]) -> [];  
qsort([X|Xs]) -> qsort([Y || Y <- Xs, Y<X])  
                ++ [X]  
                ++ qsort([Y || Y <- Xs, Y>=X]) .
```



Reducing copying with an *accumulating parameter*

- Define `asort` s.t. `asort(L,Acc) == qsort(L) ++ Acc`

```
asort(L) -> asort(L, []).
```

```
asort([], Acc) -> Acc;
```

```
asort([X|Xs],Acc) ->
```

```
    asort([Y || Y <- Xs, Y<X],
```

```
        [X | asort([Y || Y <- Xs, Y>=X], Acc)]).
```

```
24> foo:benchmark(qsort,L).
```


```
79.116342
```

```
25> foo:benchmark(asort,L).
```

```
64.204434
```



Three sorting algorithms



Haskell vs Erlang

- Sorting (different) random lists of 200K integers, on 2-core i7

	Haskell	Erlang
Sequential sort	353 ms	312 ms
Depth 5 //el sort	250 ms	153 ms

- *Despite* Erlang running on a VM!

Erlang scales
better

Erlang Distribution

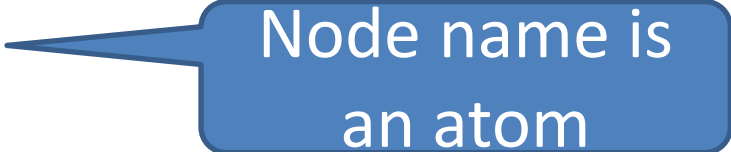
- Erlang processes can run on *different machines* with the same semantics
- No shared memory between processes!
- Just a little slower to communicate...

Named Nodes

```
werl -sname baz
```

- Start a node with a *name*

```
(baz@HALL) 1> node () .  
baz@HALL
```



Node name is
an atom

```
(baz@HALL) 2> nodes () .  
[]
```



List of connected nodes

Connecting to another node

```
net_admin:ping(Node) .
```


```
3> net_admin:ping(foo@HALL) .  
pong
```

Success—pong means
connection failed

```
4> nodes() .  
[foo@HALL,baz@JohnsTablet2014]
```

Now connected to foo and
other nodes foo knows of

Node connections





Gotcha! the Magic Cookie

- All communicating nodes must share the same *magic cookie* (an atom)
- Must be the same on all machines
 - By default, randomly generated on each machine
- Put it in `$HOME/.erlang.cookie`
 - E.g. cookie



A Distributed Sort

```
dsort([]) ->
  [];
dsort([X|Xs]) ->
  Parent = self(),
  Ref = make_ref(),
  Grtr = [Y || Y <- Xs, Y >= X],
  spawn_link(baz@JohnsDesktop,
    fun() ->
      Parent ! {Ref,psort4(Grtr)}
    end),
  psort4([Y || Y <- Xs, Y < X]) ++
  [X] ++
  receive {Ref,Greater} -> Greater
end.
```

Benchmarks

```
5> foo:benchmark (psort4,L) .  
87.23  
6> foo:benchmark (dsort,L) .  
109.27
```

- Distributed sort is *slower*
 - Communicating between nodes is slower
 - Nodes on the same machine are sharing the cores anyway!

OK...

A 2-core laptop... silly
to send it half the work

```
dsort2 ([X|Xs]) ->
```

```
...
```

```
spawn_link (baz@JohnsTablet2014 ,  
            fun () ->
```

```
....
```

```
5> foo:benchmark (psort4,L) .
```

```
87.23
```

```
6> foo:benchmark (dsort,L) .
```

```
109.27
```

```
7> foo:benchmark (dsort2,L) .
```

```
1190.33
```


Distribution Strategy

- Divide the work into 32 chunks on the master node
- Send *one chunk at a time* to each node for sorting
 - Slow nodes will get fewer chunks
- Use the fast parallel sort on each node

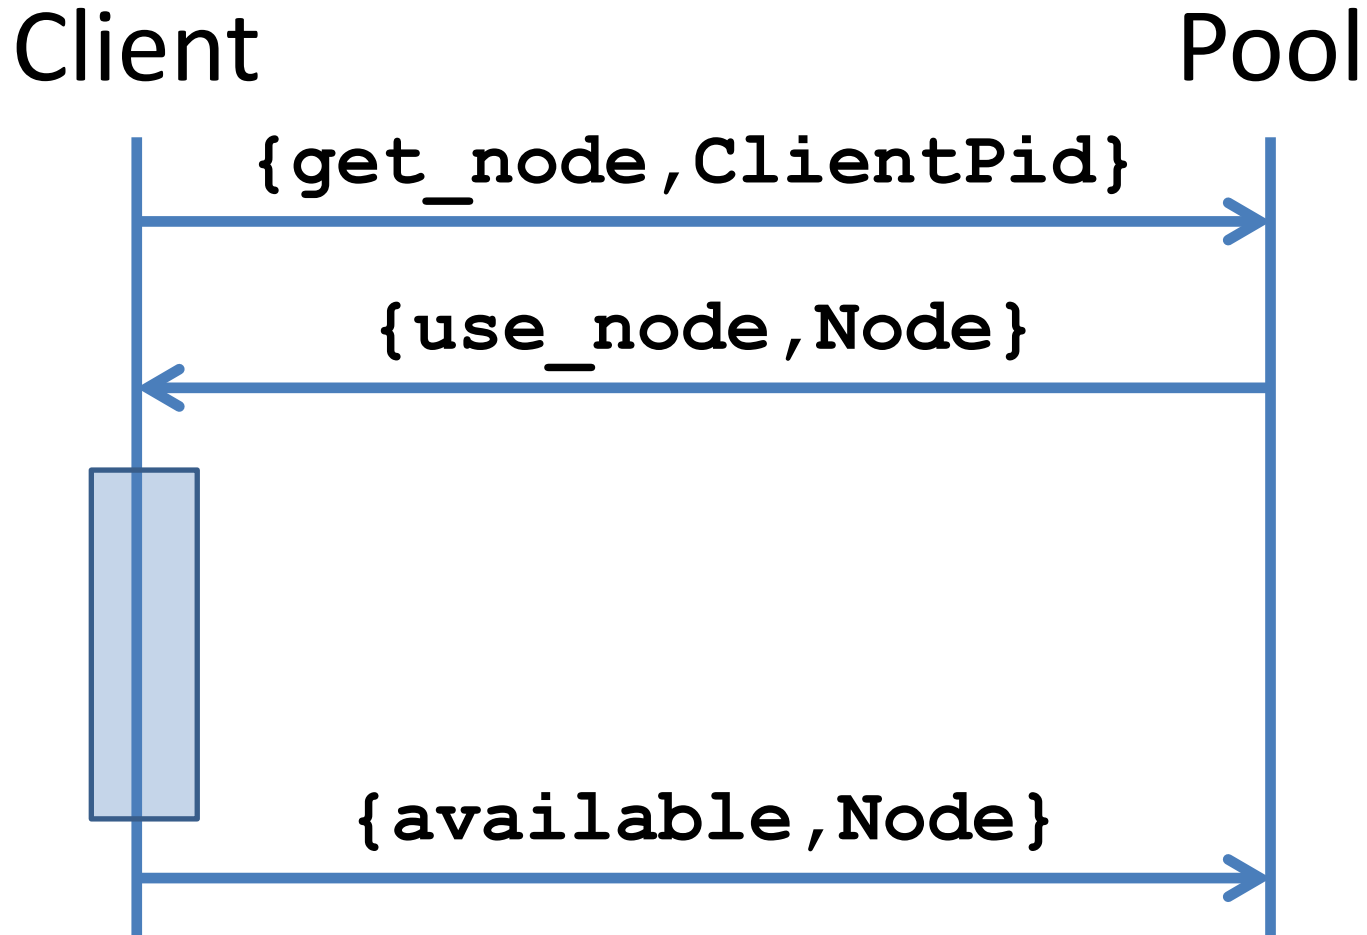
Node Pool

- We need a pool of *available nodes*

```
pool() ->  
  Nodes = [node() | nodes()],  
  spawn_link(fun() ->  
    pool(Nodes)  
  end) .
```

- We create a process to manage the pool, initially containing all the nodes

Node Pool Protocol



Node Pool Behaviour

```
pool([]) ->
  receive
    {available, Node} ->
      pool([Node])
  end;
pool([Node|Nodes]) ->
  receive
    {get_node, Pid} ->
      Pid ! {use_node, Node},
      pool(Nodes)
  end.
```

If the pool is empty, wait for a node to become

If nodes are available, wait for a request and give one out

Selective receive is really useful!

dwsort

Parallel
recursion to
depth 5

```
dwsort(Xs) -> dwsort(pool(), 5, Xs) .

dwsort(_, _, []) -> [];
dwsort(Pool, D, [X|Xs]) when D > 0 ->
    Grtr = [Y || Y <- Xs, Y >= X],
    Ref = make_ref(),
    Parent = self(),
    spawn_link(fun() ->
        Parent ! {Ref, dwsort(Pool, D-1, Grtr)}
    end),
    dwsort(Pool, D-1, [Y || Y <- Xs, Y < X]) ++
    [X] ++
    receive {Ref, Greater} -> Greater end;
```

dwsort

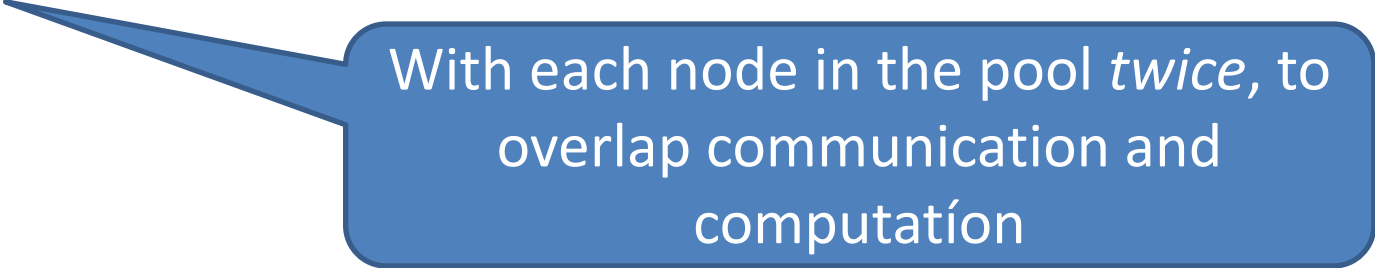
```
dwsort(Pool,0,Xs) ->
  Pool ! {get_node,self()},
  receive
    {use_node,Node} ->
      Ref = make_ref(),
      Parent = self(),
      spawn_link(Node, fun() ->
        Ys = psort4(Xs),
        Pool ! {available,Node},
        Parent ! {Ref,Ys}
      end),
      receive {Ref,Ys} -> Ys end
  end.
```

A further optimisation: if we should use the *current* node, don't spawn a new process

} Messages between nodes

Benchmarks

```
(baz@HALL) 17> foo:benchmark(qsort,L) .  
271.97  
(baz@HALL) 18> foo:benchmark(psort4,L) .  
88.65  
(baz@HALL) 19> foo:benchmark(dsorth2,L) .  
1190.33  
(baz@HALL) 20> nodes() .  
[baz@JohnsTablet2014]  
(baz@HALL) 21> foo:benchmark(dwsorth,L) .  
295.59  
(baz@HALL) 22> foo:benchmark(dwsorth2,L) .  
195.05
```



With each node in the pool *twice*, to overlap communication and computation

Oh well!

- It's quicker to *sort* a list, than to send it to another node and back!

Another Gotcha!

- All the nodes must be running *the same code*
 - Otherwise sending functions to other nodes cannot work
- **n1 (Mod)** loads the module on *all* connected nodes.

Summary

- Erlang parallelism is more explicit than in Haskell
- Processes do not share memory
- All communication is explicit by message passing
- Performance and scalability are strong points
- Distribution is easy
 - (But sorting is cheaper to do than to distribute☹)

References

- *Programming Erlang: Software for a Concurrent World*, Joe Armstrong, Pragmatic Bookshelf, 2007.
- *Learn you some Erlang for Great Good*, Frederic Trottier-Hebert , <http://learnyousomeerlang.com/>

